

Programming Models for Concurrency and Real-time



Jan Vitek

Outline

□: Real-time and embedded systems

1 : Real-time Java with Ovm

2: Memory management with Minuteman

3: Low latency programming with Flexotasks

4: Java in aerospace with the Fiji VM

5: Conclusion



What is a real-time system?

- A real-time system is any information processing system which has to respond to externally generated input stimuli within a finite and specified period
 - ▶ *correctness depends not only on logical result but also time it is delivered*
 - ▶ *failure to respond as bad as a wrong response*

What is an embedded system?

- Computer that is part of some *other* piece of equipment
 - ▶ Usually dedicated software
 - ▶ Often no “real” keyboard or general purpose display
- ... we use 100+ embedded computers daily
- ... embedded hardware growth rate of 14% to reach \$78 billion

<http://www.bccresearch.com/comm/G229R.html>, <http://www.ecpe.vt.edu/news/ar03/embedded.html>

Characteristics of real-time embedded systems

- **Large and complex** — from a few hundred lines of assembly to 20 mio lines of Ada for the Space Station Freedom
- **Concurrent control of separate components** — devices operate in parallel in the real-world; model this by concurrent entities
- **Facilities to interact with special purpose hardware** — need to be able to program devices in a reliable and abstract way
- **Extreme reliability and safe** — embedded systems control their environment; failure can result in loss of life, or economic loss
- **Guaranteed response times** — must predict with confidence the worst case; efficiency important but predictability is essential

A new software crisis?

- Development time, code & certification are increasingly criteria
- For instance in the automotive industry:
 - ▶ 90% of innovation driven by electronics and software — *Volkswagen*
 - ▶ 80% of car electronics in the future will be software-based — *BMW*
 - ▶ 80% of our development time is spent on software— *JPL*
- Worst, software is often the source of missed project deadlines.

A new software crisis?

- **Typical productivity**

- ▶ *5 Line of Code / person / day*
 - ▶ From requirements to testing: *1 kloc / person / year*

- **Typical avionics “box”**

- ▶ *100 kloc \Rightarrow 100 person years of effort*
 - ▶ Costs of modern aircraft is $\sim \$500M$

A new software crisis?

- The important metrics are thus
 - ▶ Reusability
 - ▶ Software quality
 - ▶ Development time
- The challenges are
 - ▶ Sheer number and size of systems
 - ▶ Poor programmer productivity
- The solutions are
 - ▶ Better processes (software engineering)
 - ▶ Better tools (verification, static analysis, program generation)
 - ▶ Better languages and programming models

What programming models?

- The **programming model** for most real-time systems is ‘defined’ as a function of the hardware, operating system, and libraries.
 - ▶ Consequently real-time systems **are not portable** across platforms
- **Good news**
 - ▶ programming languages, such as Java and C#, are wresting control from the lower layers of the stack and impose well-defined semantics (on threads, scheduling, synchronization, memory model)


What programming model?

- “Real-time systems require fine grained control over resources and thus the language of choice is C, C++ or assembly”
- ..entails the software engineering drawbacks of low-level code
- Consider the following list of defects that have to be eradicated
(c.f. “Diagnosing Medical Device Software Defects” Medical DeviceLink, May 2009):
 - ▶ Buffer overflow and underflow (does not occur in a HLL)
 - ▶ Null object dereference (checked exception in a HLL)
 - ▶ Uninitialized variable (does not occur in a HLL)
 - ▶ Inappropriate cast (all casts are checked in a HLL)
 - ▶ Division by zero (checked exception in a HLL)
 - ▶ Memory leaks (garbage collection in a HLL)

What programming models?

- There are many dimensions:
 - ▶ **Imperative** vs. **Functional**
 - ▶ **Shared memory** vs. **Message passing**
 - ▶ **Explicit lock-based synchronization** vs. **Higher-level abstractions**
(data-centric synchronization, transactional memory)
 - ▶ **Time-triggered** vs. **synchronous / logic execution time**
- And multiple languages, systems:
 - ▶ C, C++, Ada, SystemC, Assembler, Erlang, Esterel, Lustre, Giotto ...

Are object oriented technologies the silver bullet for the real-time software crisis?



1



Ovm

The Real-time Java
experience



Java?


- Object-oriented programming helps **software reuse**
- Mature **development environment** and **libraries**
- **Garbage collected & Memory-safe** high-level language
- **Portable**, little implementation-specific behavior
- **Concurrency** built-in, support for SMP, memory model
- **Popular** amongst educators and programmers

Java?


- Predictable?
- Not *really*.

Call sleep(10ms) and get up
20 milis.sec. variability.

Hard real-time often
requires microsecond
accuracy.




Time scale



Java?

- ▶ Predictable?



- ▶ Java Collision Detector running at 20Hz.
 - Bartlett's Mostly Copying Collector. Ovm. Pentium IV 1600 MHz, 512 MB RAM, Linux 2.6.14, GCC 3.4.4
- ▶ GC pauses cause the collision detector to miss up to three deadlines...*this is not a particularly hard should support KHz periods*

The Real-time Specification for Java (RTSJ)

- Java-like programming model:
 - ▶ Shared-memory, lock-based synchronization, first class threads.
- Main real-time additions:
 - ▶ **Physical memory access** (memory mapped I/O, devices, ...)
 - ▶ **Real-time threads** (heap and no-heap)
 - ▶ **Synchronization, Resource sharing** (priority inversion avoidance)
 - ▶ **Memory Management** (region allocation + real-time GC)
 - ▶ **High resolution Time values and Clocks**
 - ▶ **Asynchronous Event Handling and Timers**
 - ▶ **Asynchronous Transfer of Control**

Ovm

- Started on Real-time Java in 2001, in a DARPA funded project.
At the time, no real RTSJ implementation.
- Developed the Ovm virtual machine framework, a clean-room, open source RT Java virtual machine.
- Fall 2005, first flight test with Java on a plane.



Duke's Choice
Award




Case Study: ScanEagle



ScanEagle



ScanEagle



- Flight Software:


- ▶ 953 Java classes, 6616 methods.
- Multiple Priority Processing:
 - High (20Hz) - Communicate with Flight Controls
 - Medium (5 Hz) - Computation of navigation data
 - Low (1 Hz) - Performance Computation
- ▶ Embedded Planet 300 Mhz PPC, 256MB memory, Embedded Linux

- Java performed better than C++



References and acknowledgements

- Team
 - ▶ J. Baker, T. Cunei, C. Flack, D. Holmes, C. Grothoff, K. Palacz, F. Pizlo, M. Prochazka and also J. Thomas, K. Grothoff, E. Pla, H. Yamauchi, P. McGachey, J. Manson, A. Madan, B. Titzer
- Funding: DARPA, NSF, Lockheed Martin, Boeing
- Availability: open source, <http://www.cs.purdue.edu>
- Paper trail
 - A Real-time Java Virtual Machine for Avionics. RTAS, 2006
 - Scoped Types and Aspects for Real-Time Systems. ECOOP, 2006
 - A New Approach to Real-time Checkpointing. VEE, 2006
 - Real-Time Java scoped memory: design patterns, semantics. ISORC, 2004
 - Subtype tests in real time. ECOOP, 2003
 - Engineering a customizable intermediate representation. IVME, 2003



2



Minuteman

Real-time Garbage Collection

Memory management and programming models

- The choice of memory management affects productivity
- Object-oriented languages naturally hide allocation behind abstraction barriers
 - ▶ Taking care of de-allocation manually is more difficult in OO style
- Concurrent algorithms usually emphasize allocation
 - ▶ because freshly allocated data is guaranteed to be thread local
 - ▶ “transactional” algorithms generate a lot of temporary objects
- ... but garbage collection is a global, costly, operation that introduces unpredictability

Alternative I: No Allocation

- If there is no allocation, GC does not run.
 - This approach is used in JavaCard

Alt 2: Allocation in Scoped Memory


- RTSJ provides scratch pad memory regions which can be used for temporary allocation
 - Used in deployed systems, but tricky as they can cause exceptions

```
s = new SizeEstimator();
s.reserve(Decrypt.class, 2);
...
shared = new LTMemory(s.getEstimate());
shared.enter(new Run(){ public void run(){
    ...d1 = new Decrypt() ...
}});
```

Alt 3: Real-time Garbage Collection

- There are three main families of RTGC implementations
- **Work-based**
 - ▶ *Aicas JamaicaVM*
- **Time-triggered, periodic**
 - ▶ *IBM Websphere*
- **Time-triggered, slack**
 - ▶ *SUN Java Real Time System*


Garbage Collection



Phases

- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction


Garbage Collection



Phases

- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction


Garbage Collection



Phases

- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction


Garbage Collection



Phases

- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction


Garbage Collection



Phases

- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction


Garbage Collection



Phases

- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction

Garbage Collection




Phases

- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction



Garbage Collection

Phases

- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction



Incrementalizing marking



Collector marks object




Application updates
reference field




Compiler inserted
write barrier marks object

Incrementalizing compaction


- Forwarding pointers refer to the current version of objects
- Every access must start with a derefence



Time-based GC Scheduling



Slack-based GC Scheduling



Worst case = 114ms


Latency (millis)

100
80
60
40
20


0 50 100 150 200 250 300

Iteration Number

- ▶ GC pauses cause the collision detector to miss deadlines...
and this is not a particularly hard problem should support KHz periods




CD with periodic RTGC



Slack-based GC

References and acknowledgements

- Team
 - ▶ J. Baker, T. Cunei, T. Kalibera, T. Hosking, F. Pizlo, M. Prochazka
- Funding: NSF
- Availability: open source
- Paper trail
 - Accurate Garbage Collection in Uncooperative Environments. CC:P&E, 2009
 - Memory Management for Real-time Java: State of the Art. ISORC, 2008
 - Garbage Collection for Safety Critical Java. JTRES, 2007
 - Hierarchical Real-time Garbage Collection. LCTES, 2007
 - Scoped Types and Aspects for Real-time Java Memory management. RTS, 2007
 - Accurate Garbage Collection in Uncooperative Environments with Lazy Stacks. CC, 2007
 - An Empirical Evaluation of Memory Management Alternatives for Real-time Java. RTSS, 2006
 - Real-Time Java scoped memory: design patterns, semantics. ISORC, 2004



3

Flexotask

Flexible Task Graphs


Goals

- Design a new real-time programming model that allows embedding hard real-time computations in timing-oblivious Java applications
- Principle of Least Surprise
 - ▶ Semantics of non-real-time code unchanged
 - ▶ Semantics of real-time code unsurprising
- Limited set of new abstractions that compose flexibly
- No cheating
 - ▶ Run efficiently in a production environment

Unification of previous work

- **Eventrons** [PLDI'06] (IBM)
- **Reflexes** [VEE'07] (Purdue/EPFL)
 - ▶ *Inspired by RTSJ and Eventrons*
- **Exotasks** [LCTES'07] (IBM)
 - ▶ *Inspired by Giotto, and E-machine*
- **StreamFlex** [OOPSLA'07] (Purdue/EPFL)
 - ▶ *Inspired by Reflexes, StreamIt and dataflow languages*

Design space



Programming model

- **Basic model:**


- ▶ No shared state, (but local state), no low-level data races
- ▶ Components communicate via atomic channels
- ▶ Memory management is either GCed or Region-allocated
- ▶ Time triggered scheduler
 - *Inspiration: Actors, Erlang, ...*

- **Extensions:**

- ▶ Rate driven schedulers
 - *Inspiration: StreamIt, Giotto, ...*
- ▶ Weak isolation for throughput
- ▶ Transactional memory for external interaction


Flexible Task Graphs

- A *FlexoTask Graph* is a set of concurrently executing, isolated, tasks communicating through non-blocking channels
- Semantics of legacy code is unaffected
- Real-time code has restricted semantics, enforced by compile and start-up time static checks




Task Graph

- A *FlexoTask Graph* is a set of concurrently executing, isolated, tasks communicating through channels
- Schedulers control the execution of tasks with user-defined policies (eg. logical execution time, data driven)
 - ▶ atomically update task's in ports
 - ▶ invoke task's **execute()**
 - ▶ update the task's output ports



Memory management

- Either garbage collected with a real-time GC, or a region allocator for sub-millisecond response times.
- Region tasks are split between
 - Stable objects
 - Transient (per invocation) objects
- Region-allocated tasks preempt task RTGC and Java GC




Channels

- Stable channels


- Can refer to any stable object (complex structures)
- Deep copy on read (atomic)

- Transient channels

- Can refer to Capsules (transient objects, arrays)
- Zero-copy (linear reference)




Stable objects,
copy on read



Pinned Transient objects,
allocated on Java heap


Communication with Java

- Every task has an automatically generated proxy-object
- User-defined atomic methods can be called from Java with transactional semantics
- Arguments are reference-immutable pinned objects




Communication with Java

- Atomic Methods:
 - ▶ acquire a lock on guard & pin all reference-immutable arguments
 - ▶ start transaction;
 - ▶ execute method
 - ▶ commit transaction
 - ▶ reclaim transient memory
 - ▶ unpin all arguments & release lock on guard
- If during execution of the method the Task is scheduled, the transaction is immediately aborted.




Static safety

- Safety checks prevent references to transient objects after they have been deallocated and to capsules once they have been sent.
- A simple form of **ownership types** is used where **Stable** is a marker interface for data allocated in the stable heap and **Capsule** for messages. Some polymorphism needed for arrays.
- Checking is done statically, no dynamic tests are need.




Predictability

- ▶ 600K periodic invocations
- ▶ Inter-arrival time bw 57 and 144us
- ▶ 516 aborts of the atomic method




References and acknowledgements

- Team
 - ▶ J. Spring, J. Auerbach, D. Bacon, F. Pizlo, R. Guerraoui, J. Manson
- Funding: NSF & IBM
- Availability: released open source by IBM on sourceforge
- Paper trail
 - A Unified Restricted Thread Programming Model for Java. LCTES, 2008
 - StreamFlex: High-throughput Stream Programming in Java. OOPSLA, 2007
 - Reflexes: Abstractions for Highly Responsive Systems. VEE, 2007
 - Scoped Types and Aspects for Real-time Java Memory management. RTS, 2007
 - Scoped Types and Aspects for Real-Time Systems. ECOOP, 2006
 - Preemptible Atomic Regions for Real-time Java. RTSS, 2005
 - Transactional lock-free data structure for Real Time Java. CSJP, 2004



4



Fiji

Safety Critical Java

SC Java Goal

- A specification for Safety Critical Java capable of being certified under **DO-178B Level A**
 - ▶ Implies small, reduced complexity infrastructure (i.e. JVM)
 - ▶ Emphasis on defining a minimal set of capabilities required by implementations
 - ▶ Based on HIJA – High-Integrity Java Application (EU project)
 - ▶ **Final draft due this year (already 300+ page book)**

Fiji VM technology

- Proprietary ahead-of-time compiler

- ▶ Java bytecode to portable ANSI C
- ▶ high-performance, predictable execution
- ▶ Multi-core ready


- Proprietary real-time garbage collection

- ▶ easy-to-use, fully preemptible, small overhead
- ▶ zero pause times for RT tasks

- Current platforms

- ▶ OS X, Linux, RTEMS
- ▶ x86 and x64, SPARC, LEON2/3, ERC32, and PowerPC
- ▶ 200KB footprint

Execution time vs. Competitor RTJVM



RTEMS demo

- fVM runs on RTEMS 4.9.2
 - ▶ Java threads run side-by-side with RTEMS C, C++, Ada threads
- Repeat every 10 ms
 - ▶ Allocate Integer[1000] array, fill with Integer instances
 - ▶ Allocate 1000 more Integer instances
- Run code as an RTEMS interrupt handler
 - ▶ *fVM's Java runtime is robust enough to allow pure Java code to run in an interrupt context while using all of Java's features*


```
class Demo {  
    static Integer[] arr; static int iter, iWGC; static long mDWoGC, mDWGC;  
    public static void main(String[] v) {  
        final Timer t=new Timer();  
        t.fireAfter(10,new Runnable(){  
            public void run() {  
                long before=HardRT.readCPUTimestamp();  
                iter++; if (GC.inProgress()) iWGC++;  
                if (arr==null) {  
                    arr=new Integer[100000];  
                    for (int i=0;i<arr.length;++i) arr[i] = new Integer(i);  
                } else  
                    for (int i=0;i<arr.length;++i)  
                        if (!arr[i].equals(new Integer(i))) throw new Error("failed "+i);  
                t.fireAfter(10,this);  
                long diff = before-HardRT.readCPUTimestamp();  
                if (GC.inProgress()){  
                    if (diff>mDWGC) mDWGC = diff;  
                } else if (diff > mDWoGC) mDWoGC = diff;  
            } );  
        for (;;) {  
            String res = "Number of timer interrupts: "+iter +  
                         "\nNumber of timer interrupts when GC running: "+iWGC +  
                         "\nMax interrupt exec time with GC: "+ mDWGC);  
            System.out.println(res);  
            Thread.sleep(1000);  
        } } }
```

```
class Demo {  
  
    static void main(String[] v) {  
        final Timer t = new Timer();  
        t.fireAfter(10, new Runnable(){  
            public void run() {  
                long before = HardRT.getCPUTimestamp();  
                if (GC.inProgress()) iterationsWGC++;  
                arr = new Integer[1000];  
                for (int i=0;i<arr.length;+i)  
                    arr[i] = new Integer(i);  
                t.fireAfter(10, this);  
                ...  
            } } );  
        ...  
    }  
}
```

```
t = new Timer();  
t.fireAfter(10,  
new Runnable(){ void run(){  
    long before=getCPUTimestamp();  
    if (GC.inProgress()) iWGC++;  
    arr = new Integer[1000];  
    for (int i=0;i<arr.length;+i)  
        arr[i] = new Integer(i);  
    t.fireAfter(10, this);  
    ...  
} } );
```

References and acknowledgements

- Team
 - ▶ F. Pizlo, L. Ziarek, T. Kalibera, D. Tang, L. Zhao
- Funding: NSF, Fiji Systems LLC
- Availability: to be GPLed for research
- Paper trail
 - Real-time Java in Space: Potential Benefits and Open Challenges. DASIA, 2009
 - A Technology Compatibility Toolkit for Safety Critical Java. 2009





Conclusion

- Realtime Specification for Java:
 - ▶ <http://www.rtsj.org>
- Safety Critical Java:
 - ▶ JSR-302 <http://jcp.org>
- Fiji VM:
 - ▶ <http://www.fiji-systems.com>
- Ovm:
 - ▶ <http://www.cs.purdue.edu/homes/jv>