

High-level Programming Models for Real-time

Jan Vitek

Real-time embedded systems

- **Large and complex** — from a few hundred lines of assembly to 20 mio lines of Ada for the Space Station Freedom
- **Concurrent control of separate components** — devices operate in parallel in the real-world; model this by concurrent entities
- **Facilities to interact with special purpose hardware** — need to be able to program devices in a reliable and abstract way
- **Extreme reliability and safe** — embedded systems control their environment; failure can result in loss of life, or economic loss
- **Guaranteed response times** — must predict with confidence the worst case; efficiency important but predictability is essential

A new software crisis?

- Development time, code & certification are increasingly criteria
- For instance in the automotive industry:
 - ▶ 90% of innovation driven by electronics and software — *Volkswagen*
 - ▶ 80% of car electronics in the future will be software-based — *BMW*
 - ▶ 80% of our development time is spent on software — *JPL*
- Worst, software is often the source of missed project deadlines.

A new software crisis?

- **Typical productivity**

- ▶ *5 Line of Code / person / day*
- ▶ From requirements to testing: *1 kloc / person / year*

- **Typical avionics “box”**

- ▶ *100 kloc ⇒ 100 person years of effort*
- ▶ Costs of modern aircraft is ~\$500M

A new software crisis?

- The important metrics are thus
 - ▶ Reusability
 - ▶ Software quality
 - ▶ Development time
- The challenges are
 - ▶ Sheer number and size of systems
 - ▶ Poor programmer productivity
- The solutions are
 - ▶ Better processes (software engineering)
 - ▶ Better tools (verification, static analysis, program generation)
 - ▶ Better languages and programming models


What programming models?

- The **programming model** for most real-time systems is ‘defined’ as a function of the hardware, operating system, and libraries.
 - ▶ Consequently real-time systems **are not portable** across platforms

What programming model?

- “Real-time systems require fine grained control over resources and thus the language of choice is C or assembly”
- ...entails the software engineering drawbacks of low-level code
- Consider the following list of defects that have to be eradicated (c.f. “Diagnosing Medical Device Software Defects” Medical DeviceLink, May 2009):
 - ▶ Buffer overflow and underflow (does not occur in a HLL)
 - ▶ Null object dereference (checked exception in a HLL)
 - ▶ Uninitialized variable (does not occur in a HLL)
 - ▶ Inappropriate cast (all casts are checked in a HLL)
 - ▶ Division by zero (checked exception in a HLL)
 - ▶ Memory leaks (garbage collection in a HLL)

What programming model?



- Some of the guarantees can be retrofitted on legacy C programs.
- [Implementation of the Memory-safe Full ANSI-C Compiler, PLDI 2009]

1




Java?

- Object-oriented programming helps **software reuse**
- Mature **development environment** and **libraries**
- **Garbage collected & Memory-safe** high-level language
- **Portable**, little implementation-specific behavior
- **Concurrency** built-in, support for SMP memory model
- **Popular** amongst educators and programmers

Java?

- ▶ Predictable?



- ▶ Java Collision Detector running at 20Hz.
 - Bartlett's Mostly Copying Collector. Ovm. Pentium IV 1600 MHz, 512 MB RAM, Linux 2.6.14, GCC 3.4.4
- ▶ GC pauses cause the collision detector to miss up to three deadlines...this is not a particularly hard should support KHz periods

The Real-time Specification for Java (RTSJ)

- Java-like programming model:
 - ▶ Shared-memory, lock-based synchronization, first class threads.
- Main real-time additions:
 - ▶ **Physical memory access** (memory mapped I/O, devices, ...)
 - ▶ **Real-time threads** (heap and no-heap)
 - ▶ **Synchronization, Resource sharing** (priority inversion avoidance)
 - ▶ **Memory Management** (region allocation + real-time GC)
 - ▶ **High resolution Time values and Clocks**
 - ▶ **Asynchronous Event Handling and Timers**
 - ▶ **Asynchronous Transfer of Control**

Ovm
The Real-time Java
experience



Ovm

- Started on Real-time Java in 2001, in a DARPA funded project. At the time, no real RTSJ implementation.
- Developed the Ovm virtual machine framework, a clean-room, open source RT Java virtual machine.
- Fall 2005, first flight test with Java on a plane.



Duke's Choice
Award




Case Study: ScanEagle



ScanEagle



ScanEagle



- Flight Software:


- ▶ 953 Java classes, 6616 methods.
- Multiple Priority Processing:
 - High (20Hz) - Communicate with Flight Controls
 - Medium (5 Hz) - Computation of navigation data
 - Low (1 Hz) - Performance Computation
- ▶ Embedded Planet 300 Mhz PPC, 256MB memory, Embedded Linux

- Java performed better than C++



References and acknowledgements

- Team
 - ▶ J. Baker, T. Cunei, C. Flack, D. Holmes, C. Grothoff, K. Palacz, F. Pizlo, M. Prochazka and also J. Thomas, K. Grothoff, E. Pla, H. Yamauchi, P. McGachey, J. Manson, A. Madan, B. Titzer
- Funding: DARPA, NSF, Lockheed Martin, Boeing
- Availability: open source, <http://www.cs.purdue.edu>
- Paper trail
 - *A Real-time Java Virtual Machine for Avionics.* RTAS, 2006
 - *Scoped Types and Aspects for Real-Time Systems.* ECOOP, 2006
 - *A New Approach to Real-time Checkpointing.* VEE, 2006
 - *Real-Time Java scoped memory: design patterns, semantics.* ISORC, 2004
 - *Subtype tests in real time.* ECOOP, 2003
 - *Engineering a customizable intermediate representation.* IVME, 2003





Fiji VM technology

- Proprietary ahead-of-time compiler

- ▶ Java bytecode to portable ANSI C
- ▶ high-performance, predictable execution
- ▶ Multi-core ready


- Proprietary real-time garbage collection

- ▶ easy-to-use, fully preemptible, small overhead
- ▶ zero pause times for RT tasks


- Current platforms

- ▶ OS X, Linux, RTEMS
- ▶ x86 and x64, SPARC, LEON2/3, ERC32, and PowerPC
- ▶ 200KB footprint


The Fiji VM Overview



The Runtime



Better view of Fiji CI



Performance/Predictability

local assignments,
simple arithmetic,
casts, conditionals



*same performance as
C/C++*

loops,
method invocation,
field/array access,
static initialization



*slightly slower than
C/C++*

allocation, locking,
exceptions



faster than C/C++

condition variables,
threading, I/O



identical to C/C++

CDx Benchmark

- Representative Real-time benchmark
 - Aircraft detection based on simulated radar frames
- CDc - written in idiomatic C
- CDj - written in idiomatic Java
- Uses many arrays and is computationally intensive

CDx Benchmark

- The algorithm detects a collision whenever the distance between aircraft is smaller than a specified “proximity radius”
- Step 1: ← eliminates planes at large distances
 - split aircraft into clusters
- Step 2: ← closer examinations of potential collisions
 - for each cluster determine actual collisions


CDx Benchmark

- What if we run CDx on a real-time setup?
 - RTEMS 4.9.1 (hard RTOS microkernel: no processes or virtual memory)
 - 40MHz LEON3 with 64MB RAM (radiation-hardened SPARC)
- This is the platform used by ESA and NASA




CDx Configuration

- 6 airplanes in our airspace
- execute over 10,000 radar frames
 - runs take on average 45 minutes
 - slight modification to generate frames
- 300ms period for the collision detector task
 - between 145ms - 275ms
 - leaves less than 50% of the schedule for the GC



CDc Summary



CDj Summary

30% slower on
average

Iterations w/ potential
Collisions

Number of Iterations

1000
500
100
50
10
5
1

100


150

200

250

300

Time (ms)




Source of overheads

measured using RTBx data logger

- Expect to see larger Java overheads when potential collisions are detected
- Array bounds checks
- Type checks
- Null checks




Array Bounds Checks




Number of Checks correlated against execution time

Null Checks



Number of Checks correlated against execution time

Type Checks



Number of Checks correlated against execution time

Correlation Java vs C when running on RTEMS/LEON3

