

High-level Programming Models for Real-time

Jan Vitek



PURDUE
UNIVERSITY

Fiji

Programming Models for Concurrency and Real-time

Real-time embedded systems

- **Large and complex** — from a few hundred lines of assembly to 20 mio lines of Ada for the Space Station Freedom
- **Concurrent control of separate components** — devices operate in parallel in the real-world; model this by concurrent entities
- **Facilities to interact with special purpose hardware** — need to be able to program devices in a reliable and abstract way
- **Extreme reliability and safe** — embedded systems control their environment; failure can result in loss of life, or economic loss
- **Guaranteed response times** — must predict with confidence the worst case; efficiency important but predictability is essential

Programming Models for Concurrency and Real-time

A new software crisis?

- Development time, code & certification are increasingly criteria
- For instance in the automotive industry:
 - ▷ 90% of innovation driven by electronics and software — Volkswagen
 - ▷ 80% of car electronics in the future will be software-based — BMW
 - ▷ 80% of our development time is spent on software — JPL
- Worst, software is often the source of missed project deadlines.

A new software crisis?

- Typical productivity

- ▶ 5 Line of Code / person / day
- ▶ From requirements to testing: 1 kloc / person / year

- Typical avionics “box”

- ▶ 100 kloc \Rightarrow 100 person years of effort
- ▶ Costs of modern aircraft is $\sim \$500M$

A new software crisis?

- The important metrics are thus

- ▶ Reusability
- ▶ Software quality
- ▶ Development time

- The challenges are

- ▶ Sheer number and size of systems
- ▶ Poor programmer productivity

- The solutions are

- ▶ Better processes (software engineering)
- ▶ Better tools (verification, static analysis, program generation)
- ▶ Better languages and programming models

What programming models?

- The **programming model** for most real-time systems is 'defined' as a function of the hardware, operating system, and libraries.

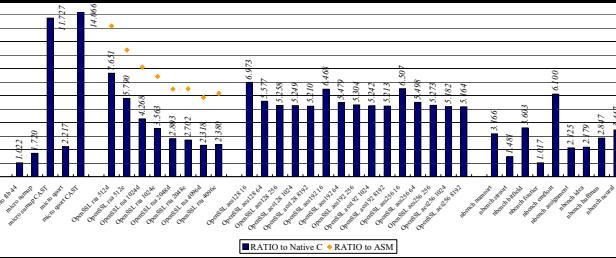
- ▶ Consequently real-time systems are **not portable** across platforms

What programming model?

- “Real-time systems require fine grained control over resources and thus the language of choice is C or assembly”
 - ...entails the software engineering drawbacks of low-level code
 - Consider the following list of defects that have to be eradicated
(§ 5 “Designing Medical Device Software Defects” Medical Devices Int., Mar. 2009)

- | | |
|---------------------------------|----------------------------------|
| ▶ Buffer overflow and underflow | (does not occur in a HLL) |
| ▶ Null object dereference | (checked exception in a HLL) |
| ▶ Uninitialized variable | (does not occur in a HLL) |
| ▶ Inappropriate cast | (all casts are checked in a HLL) |
| ▶ Division by zero | (checked exception in a HLL) |
| ▶ Memory leaks | (garbage collection in a HLL) |

What programming model?



- Some of the guarantees can be retrofitted on legacy C programs
 - [Implementation of the Memory-safe Full ANSI-C Compiler, PLDI 2009]

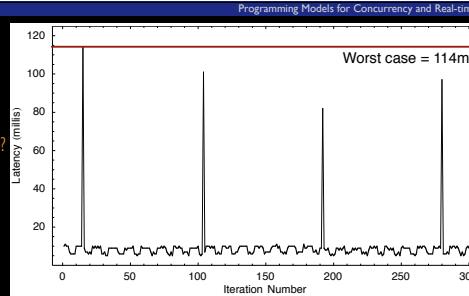


Java?

- Object-oriented programming helps **software reuse**
- Mature **development environment** and **libraries**
- **Garbage collected & Memory-safe** high-level language
- **Portable**, little implementation-specific behavior
- **Concurrency** built-in, support for SMP memory model
- **Popular** amongst educators and programmers

Java?

▷ **Predictable?**



▷ Java Collision Detector running at 20Hz.

- Bartlett's Mostly Copying Collector. Ovn. Pentium IV 1600 MHz, 512 MB RAM, Linux 2.6.14, GCC 3.4.4

▷ GC pauses cause the collision detector to miss up to three deadlines...this is not a particularly hard should support KHz periods

The Real-time Specification for Java (RTSJ)

- Java-like programming model:
 - ▷ Shared-memory, lock-based synchronization, first class threads.
- Main real-time additions:
 - ▷ **Physical memory access** (memory mapped I/O, devices, ...)
 - ▷ **Real-time threads** (heap and no-heap)
 - ▷ **Synchronization, Resource sharing** (priority inversion avoidance)
 - ▷ **Memory Management** (region allocation + real-time GC)
 - ▷ **High resolution Time values and Clocks**
 - ▷ **Asynchronous Event Handling and Timers**
 - ▷ **Asynchronous Transfer of Control**

Ovm

The Real-time Java
experience



Programming Models for Concurrency and Real-time

Ovm

- Started on Real-time Java in 2001, in a DARPA funded project. At the time, no real RTSJ implementation.
- Developed the Ovm virtual machine framework, a clean-room, open source RT Java virtual machine.
- Fall 2005, first flight test with Java on a plane.



Duke's Choice
Award



Programming Models for Concurrency and Real-time

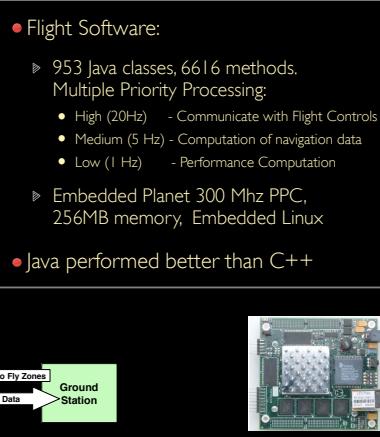
Case Study: ScanEagle



ScanEagle



ScanEagle



References and acknowledgements

Team

- J. Baker, T. Cunei, C. Flack, D. Holmes, C. Grothoff, K. Palacz, F. Pizlo, M. Prochazka and also J. Thomas, K. Grothoff, E. Pla, H. Yamauchi, P. McGachey, J. Mansor, A. Madan, B. Titzer

Funding: DARPA, NSF, Lockheed Martin, Boeing

Availability: open source, <http://www.cs.purdue.edu>

Paper trail

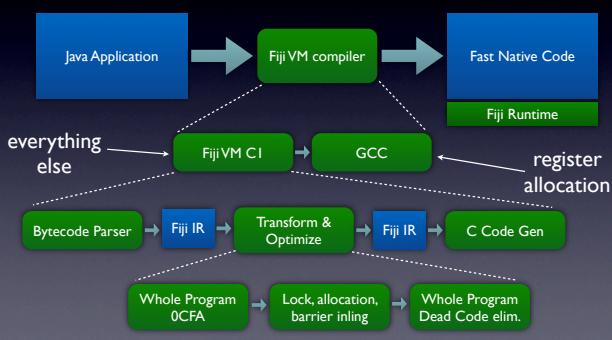
- A Real-time Java Virtual Machine for Avionics. RTAS, 2006
- Scoped Types and Aspects for Real-Time Systems. ECOOP, 2006
- A New Approach to Real-time Checkpointing. VEE, 2006
- Real-Time Java scoped memory: design patterns, semantics. ISORC, 2004
- Subtype tests in real time. ECOOP, 2003
- Engineering a customizable intermediate representation. IVME, 2003



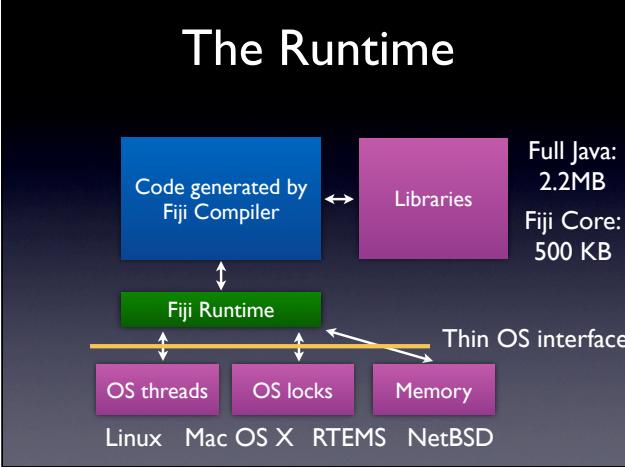
2



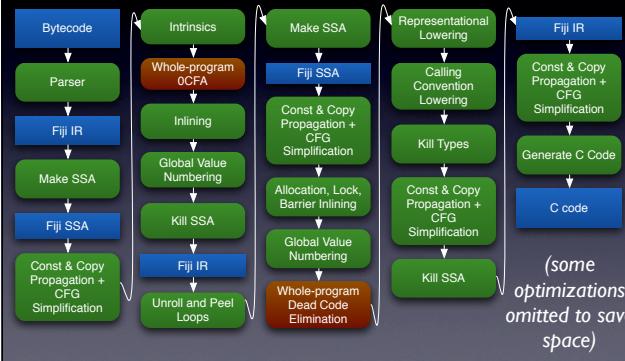
The Fiji VM Overview



The Runtime



Better view of Fiji CI



Performance/Predictability

| | |
|--|---------------------------------|
| local assignments, simple arithmetic, casts, conditionals | → same performance as C/C++ |
| loops, method invocation, field/array access, static initialization | → slightly slower than C/C++ |
| allocation, locking, exceptions | → faster than C/C++ |
| condition variables, threading, I/O | → identical to C/C++ |

CDx Benchmark

- Representative Real-time benchmark
- Aircraft detection based on simulated radar frames
- CDc - written in idiomatic C
- CDj - written in idiomatic Java
- Uses many arrays and is computationally intensive

CDx Benchmark

- The algorithm detects a collision whenever the distance between aircraft is smaller than a specified “proximity radius”
- Step 1: ← eliminates planes at large distances
 - split aircraft into clusters
- Step 2: ← closer examinations of potential collisions
 - for each cluster determine actual collisions

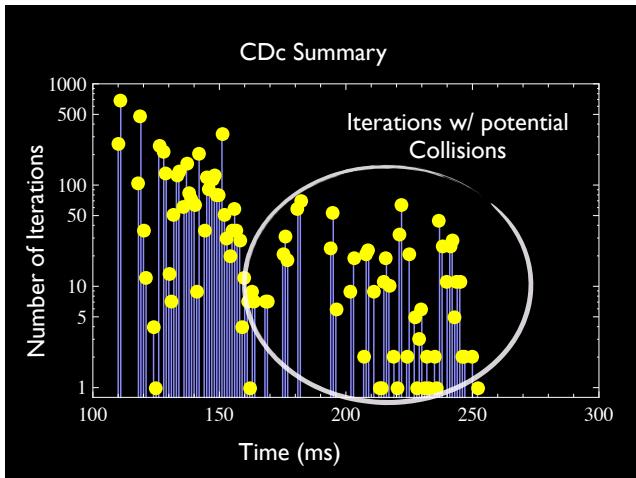
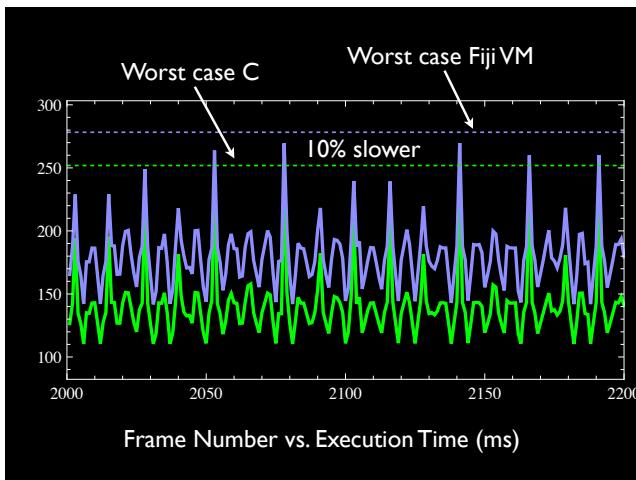
CDx Benchmark

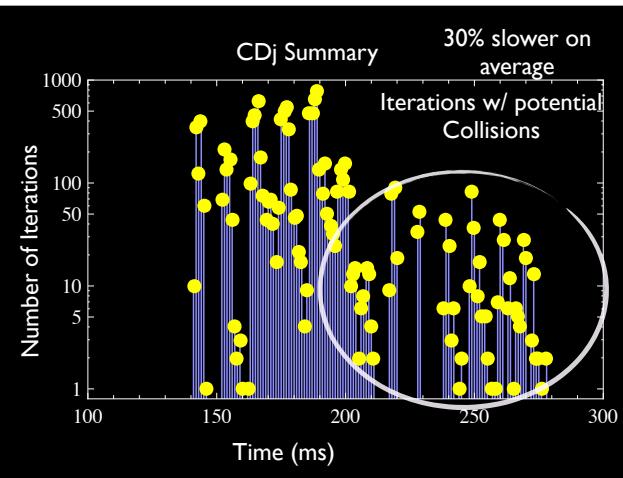
- What if we run CDx on a real-time setup?
- RTEMS 4.9.1 (hard RTOS microkernel: no processes or virtual memory)
- 40MHz LEON3 with 64MB RAM (radiation-hardened SPARC)
- This is the platform used by ESA and NASA



CDx Configuration

- 6 airplanes in our airspace
- execute over 10,000 radar frames
 - runs take on average 45 minutes
 - slight modification to generate frames
- 300ms period for the collision detector task
 - between 145ms - 275ms
- leaves less than 50% of the schedule for the GC

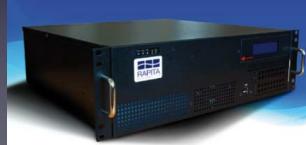




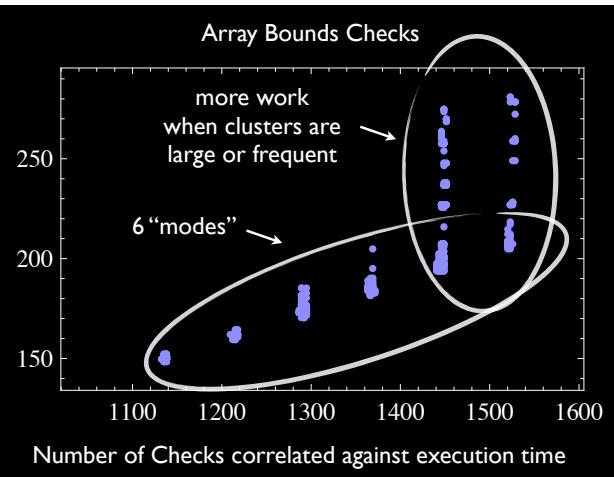
Source of overheads

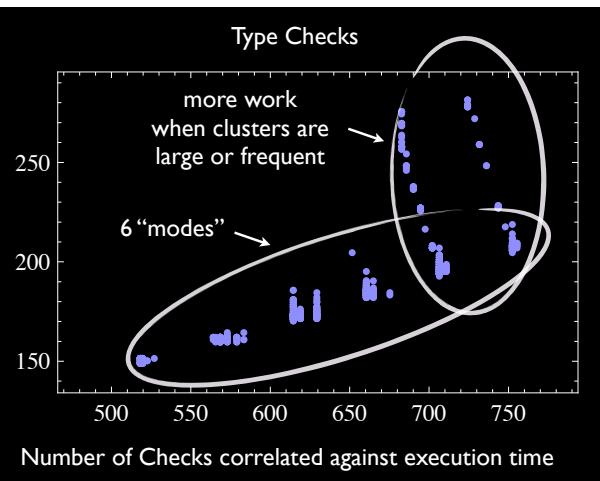
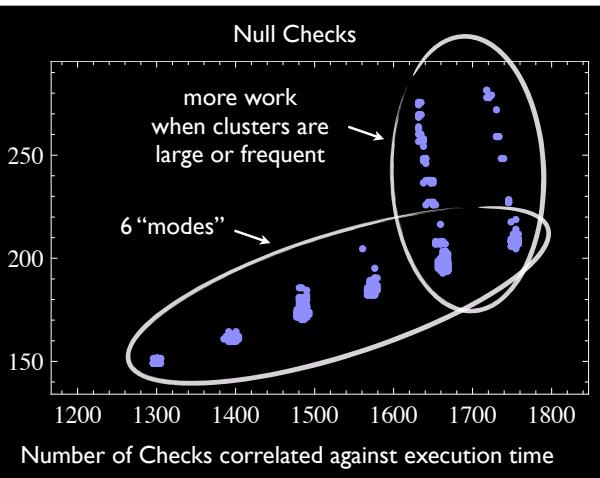
measured using RTBx data logger

- Expect to see larger Java overheads when potential collisions are detected
- Array bounds checks
- Type checks
- Null checks



www.rapitasystems.com





Correlation Java vs C when running on RTEMS/LEON3

