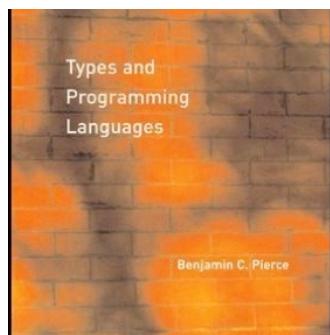


Virtual and Interface Methods



Method Invocation

- Method invocation is one of the key features of OOPs
- The process of determining which code to run is called “method dispatch” and depends on the runtime type of the receiver (late binding).
- We will take Java as a practical example and study how compile-time resolution and runtime dispatch interact.

Methods

- Methods implement the behavior associated with classes and objects

```
[public|private|protected|abstract|final|synchronized|native]*  
[<Type>]  <MethodName> ( [[final] <Type> <Name>]+ )  
[throws <Type>+]  
[ {<Block>} | ; ]
```

- `abstract` methods have ";" as body. All other methods must have a `{<block>}` body
- `final` methods can not be overridden
- constructors have no return types
- methods must not fall off the end of the body without returning the proper type. But the following is allowed

```
int m() throw E { throw new E(); }
```

Method signatures

- The signature of a method consists of the name of the method and types of formal parameters to the method.

- Example:

```
public int foo( char c, HashTbl b)  
has signature  
foo(C;HashTbl)
```

- A class may not declare two methods with the same signature.

Overriding & Hiding

[override] A method **A.m** overrides a method **B.n** if

1. $A <:_{s} B$,
2. $\text{signature}(m) = \text{signature}(n)$
3. $\text{return-type}(m) = \text{return-type}(n)$
4. $\text{throws}(m) \leq \text{throws}(n)$
5. if **n** is **public** then **m** is **public**, if **n** is **protected** then **m** is (**protected|public**), if **n** is default-access then **m** is not **private**
6. **n** is not **static**

[hide] A static method **A.m** hides a static method **B.n** if

1. $A <:_{s} B$,
2. $\text{signature}(m) = \text{signature}(n)$
3. $\text{return-type}(m) = \text{return-type}(n)$
4. $\text{throws}(m) \leq \text{throws}(n)$
5. if **n** is **public** then **m** is **public**, if **n** is **protected** then **m** is (**protected|public**), if **n** is default-access then **m** is not **private**

Overloading

- If two methods of a class have the same name but different signatures they are said to be overloaded.

Example:

```
class C {  
    int foo( C bar ) {}  
    int foo( B bore) {}  
    void foo( int i) throws Stuff {}  
    char foo( B bore) throws IllegalValue {} // error
```

Method invocation

- We consider the process of binding a source code method invocation expression to an actual method
 - partly done at compile time
 - partly done at run time
- Syntax:

```
<MethodName> ( <args> ) |  
<Primary> . Identifier ( <args> ) |  
super . Identifier ( <args> )
```

Examples:

```
foo(arg)  
obj.foo(arg)  
HashTbl.foo(arg)  
obj.foo(arg).bar()
```

Determine Unit and Name (step 1)

- Determine the name of the method and the unit (class or interface) in which the method is located. Consider the cases:

1. `<MethodName> (<args>)`

`if` `MethodName` `is`

- ... an identifier then the unit is the one containing the method invocation and `MethodName` is the name of the method
- ... a qualified name `TypeName.Identifier` then the unit is `TypeName` and `Identifier` is the method's name
- ... a qualified name `FieldName.Identifier` then the unit is type-of (`FieldName`) and `Identifier` is the method's name

2. `<Primary>.Identifier(<args>)`

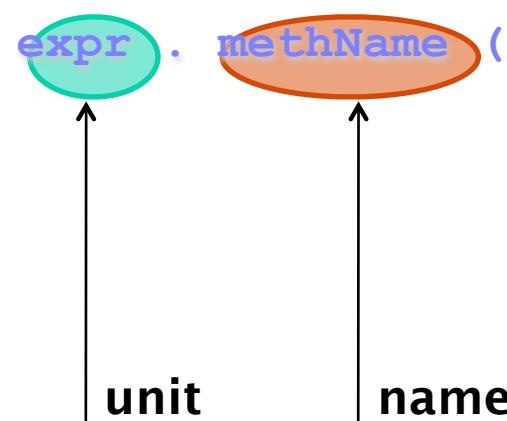
unit is type of `Primary` expr, `Identifier` is method's name

3. `super.Identifier(<args>)`

unit is the superclass, `Identifier` is method's name

Determine Unit and Name (step 1)

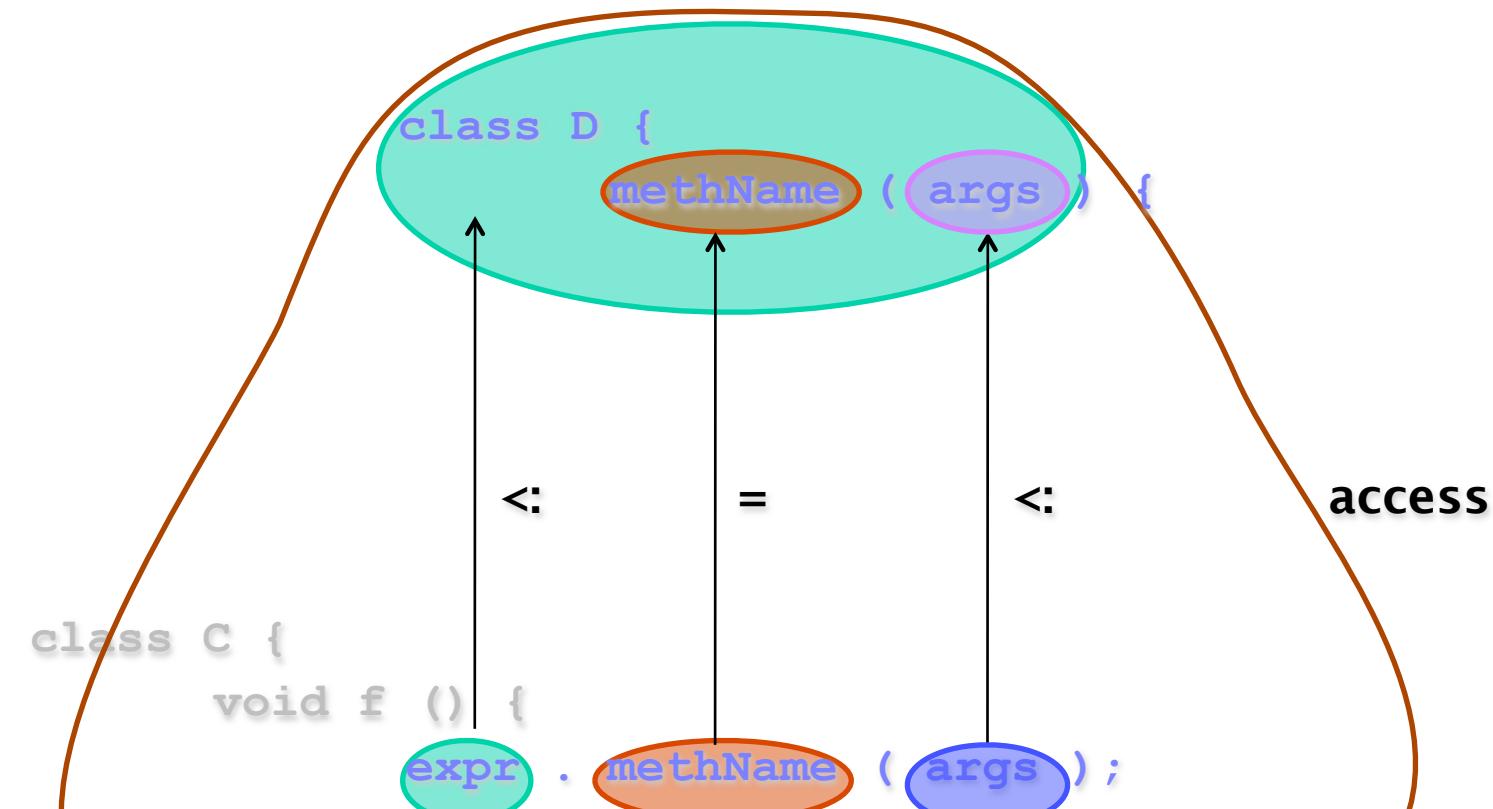
```
class C {  
    void f () {  
        expr . methName ( args );
```



Determine Method Signature (step 2)

- ➊ Search the unit for method that are **applicable** and **accessible**.
 1. A method is **applicable** if
 - number of parameters = number of arguments
 - type of each argument can be converted by method invocation conversion to the corresponding parameter
 - the name of the method is the one determined in step 1All inherited methods are included in this search.
 2. Whether a method is **accessible** to an invocation depends on the access modifier in the declaration and on where the method appears.

Determine Method Signature (step 2)



Determine Method Signature (step 3)

- Then choose the most specific method.

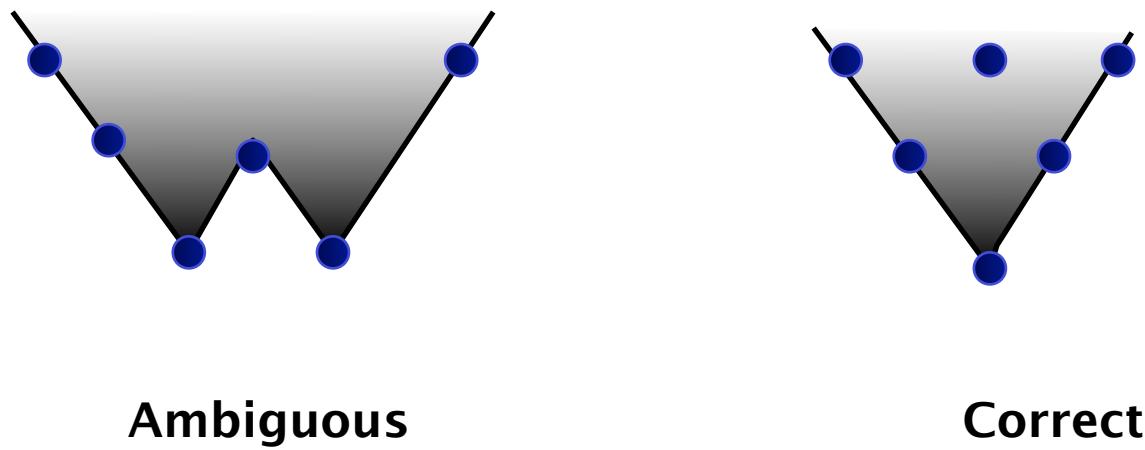
Given methods $T\ m(T_1, \dots, T_n)$ and $U\ m'(U_1, \dots, U_n)$, m is **more specific** than m' if

- T can be converted to U by method invocation conversion
- for all i from 1 to n , T_i can be converted to U_i by method invocation conversion

A method is said to be **maximally specific** for an invocation if it is applicable and accessible and there is no method more specific

A method is said to be **the most specific** if it is the only maximally specific method for an invocation.

Determine Method Signature (step 2)



Determine Appropriateness (step 3)

- If there is a most specific declaration for a method invocation, it is called the **compile-time declaration**.

Check:

1. `<MethodName>(<args>)`
if the invocation occurs within a static method or an initializer then the compile-time declaration must be static.
2. `<TypeName>. <MethodName>(<args>)`
then the compile-time declaration must be static.
3. If the return type is `void`, then the invocation must be a top-level expression.

- Output:

- method name,
- class or interface that contains the compile-time declaration
- number and type of parameters, result type,
- invocation mode: static, nonvirtual (private), super, interface, virtual

Method Invocation (step 4)

- At run-time, method invocation requires five steps:

- Compute the target reference

static methods have no target, otherwise it is either this or the value of the primary

- Evaluate arguments in order from left to right.

- Check accessibility of type and method

Throw `NoSuchMethodError` if method not found, `IncompatibleClassChangeError` if the class does not implement the interface in which the method was found, `IllegalAccessException` if the class or method is not accessible due to access modifier changes.

- Locate method to invoke

Throw `NullPointerException` if target is null.

- Create frame, synchronize and transfer control

Implementing Method Invocation



- Four types of invocation bytecodes:
 - invokespecial -- direct function call (compile/load-time binding)
 - invokestatic -- direct function call (compile/load-time binding)
 - invokevirtual -- indirect call (runtime binding)
 - invokeinterface -- indirect call (runtime binding)

The difference between invokevirtual and invokeinterface boils down to the difference between single and multiple inheritance. (A class may have a single parent, but can implement multiple interfaces)

Invokevirtual



- Given a call:

```
x.foo(c, "hello")
```

- At compile-time the type of `x` is some class `T`, the signature of the method is `foo(C, String);`
- At run-time, type preservation ensures that `x` will refer to an object of type `T`. If `T` is a class, the object referred to from `x` can be either a `T` or a subclass of `T`.
- Implementation:
 - Every method signature defined in class `T` is given a unique offset which must be larger than any offset of inherited methods
 - A table of code pointers is associated to every object.
 - The compiler generates a call to `x->vtable [FOO_C_STRING]`

Invokeinterface



- Given a call:

```
x.foo(c, "hello")
```

- At compile-time the type of `x` is some interface `T`, the signature of the method is `foo(C, String);`

- Fun fact one:

the JVM verifier does not check that `x` has a type that implements `T`!

- Fun fact two:

`T` needs only be loaded when the call is executed.

- How do we implement this efficiently?

Class Object Search



- Naive strategy: every class maintains a list of directly implemented interfaces.
- Implement `invokeinterface` by searching the hierarchy and performing all dynamic checks required by the specification.

Searched ITables



- Every class has a list of interface tables.
- Every interface has unique IID
- Every method defined by an interface has unique MID.
- Implementation of invokeinterface given an oref, IID, MID:
 - get the interface table from the oref
 - scan it for the IID
 - if found
 - move-to-front the IID in the table
 - use the MID to index in the itable and jump
 - if not found, do a dynamic type check,
 - load the interface and add the itable or throw an exception

Directly Index ITables



- Every class has an array of interface tables.
- Every interface has an IID (not necessarily unique)
- Every method defined by an interface has unique MID.
- Implementation of invokeinterface given an oref, IID, MID:
 - do a subtype test to ensure that oref implements the interface
 - get the interface table from the oref and index it at IID,
 - use the MID to index in the itable and jump
- As an optimization: initialize the ITable with thunks that do the subtype test and load the interfaces

Jikes IMTs [OOPSLA01]



- Every method defined in an interface is assigned an offset (not unique)
- Every class has an Interface Method Table (IMT)
- Invokeinterface uses the offset to index the table and when there is a conflict, compiles a stub for resolving which method to call.

Jikes IMTs [OOPSLA01]



```
// virtual invocation sequence
// t0 contains a reference to the receiver object
L s2, tibOffset(t0) // s2 := TIB of the receiver
L s2, vmtOffset(s2) // s2 := VMT entry for method being dispatched
MTLR s2             // move target address to the link register
BLRL                 // branch to it (setting LR to return address)

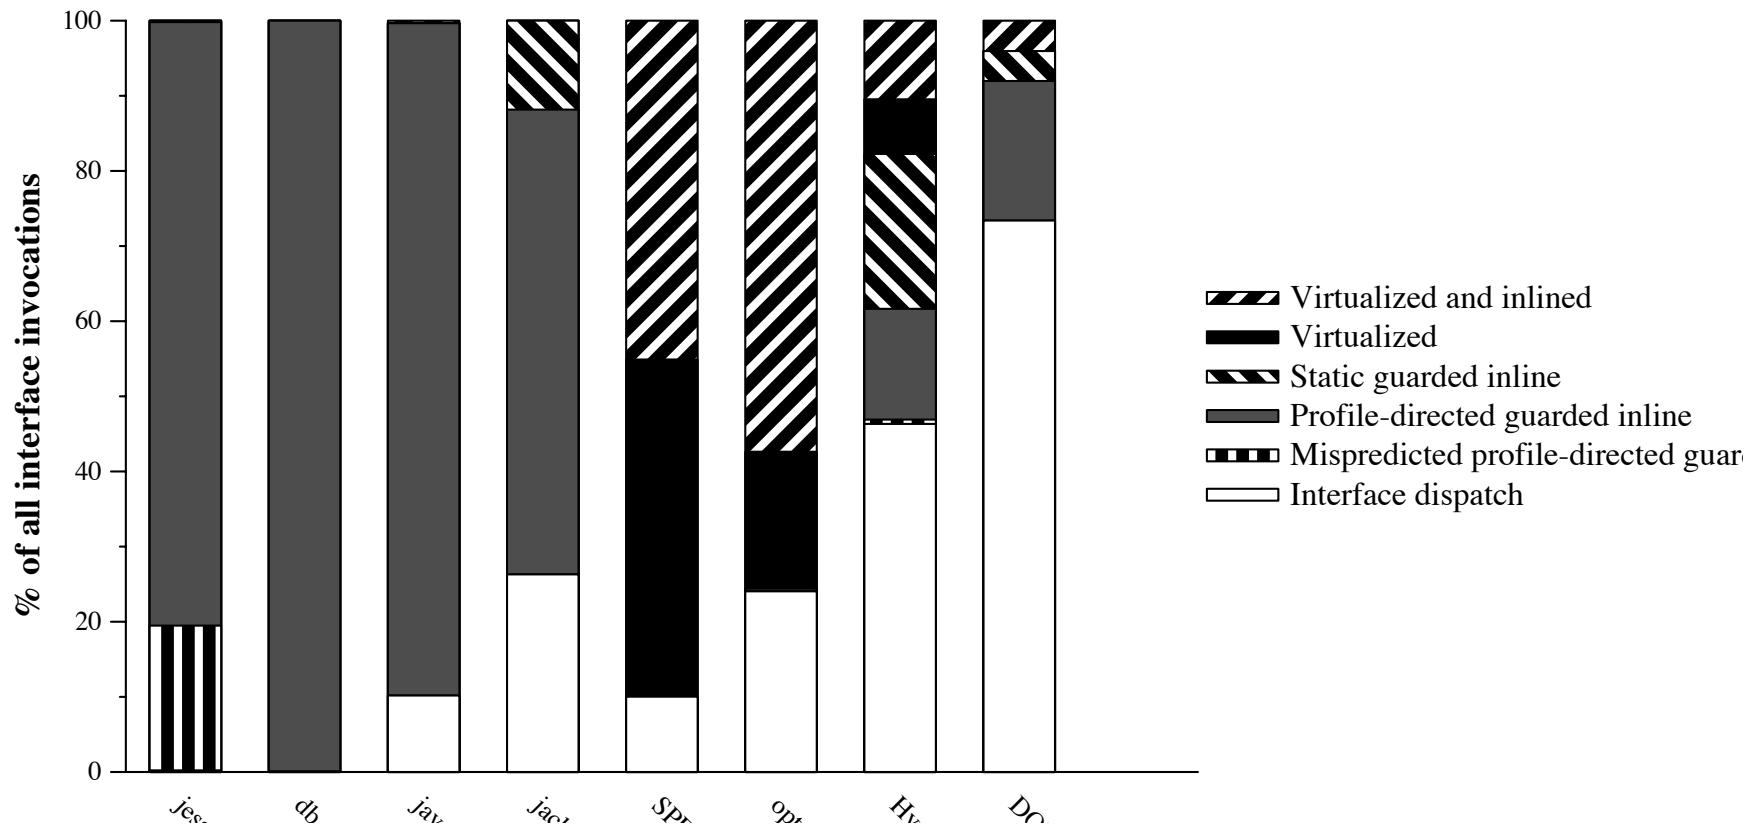
// IMT-based interface invocation sequence
// t0 contains a reference to the receiver object
L s2, tibOffset(t0) // s2 := TIB of the receiver
L s2, imtOffset(s2) // s2 := IMT entry for signature being dispatched
L s1, signatureId   // put signature id in hidden parameter register
MTLR s2             // move target address to the link register
BLRL                 // branch to it (setting LR to return address)
```

Benchmarks

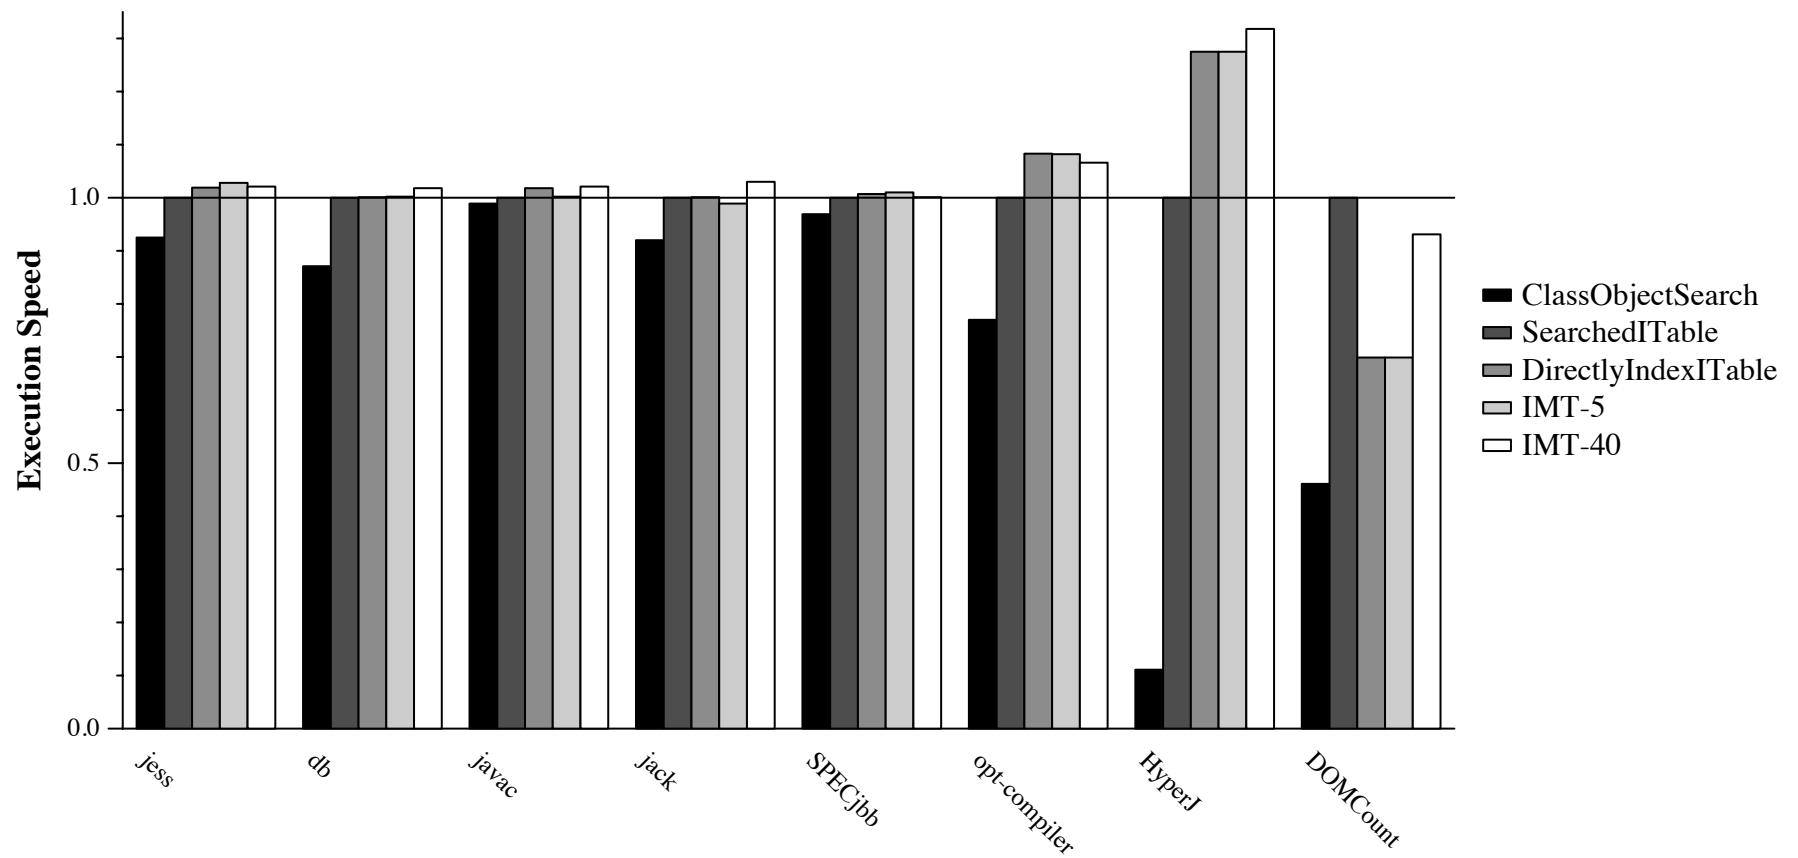


Dispatching Mechanism	Number of Methods in Interface	Trivial Callee
virtual	Not applicable	8.13
Embedded IMT-5	interface with 1 method (no IMT conflict)	8.23
Embedded IMT-5	interface with 100 methods (20 element stub)	20.25
Embedded IMT-40	interface with 1 method (no IMT conflict)	8.23
Embedded IMT-40	interface with 100 methods (2 or 3 element stub)	14.23
IMT-5	interface with 1 method (no IMT conflict)	10.18
IMT-5	interface with 100 methods (20 element stub)	22.20
IMT-40	interface with 1 method (no IMT conflict)	10.18
IMT-40	interface with 100 methods (2 or 3 element stub)	18.20
Directly Indexed ITables	interface with 1 method	12.18
Directly Indexed ITables	interface with 100 methods	12.18
Searched ITables	interface with 1 method	77.45
Searched ITables	interface with 100 methods	77.45
Class Object Search	interface with 1 method	352.55
Class Object Search	interface with 100 methods	1,743.13

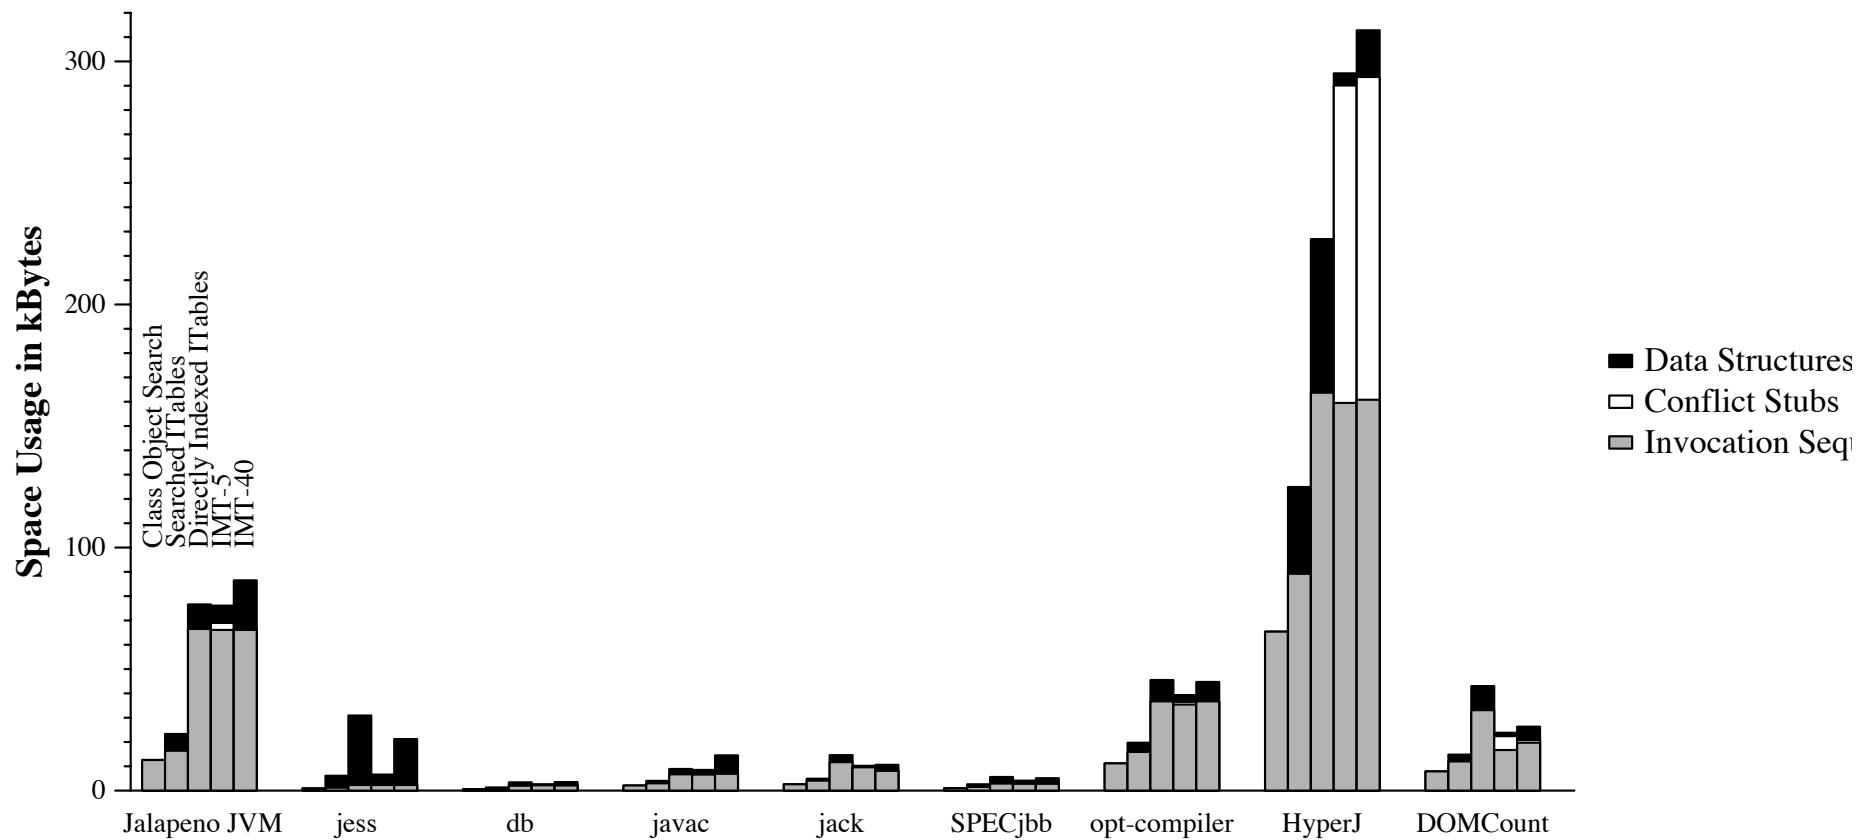
Benchmarks



Benchmarks

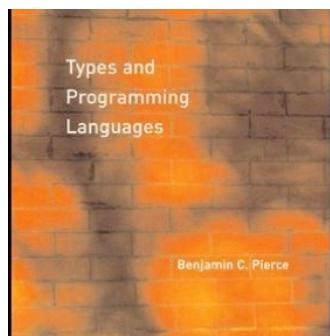


Benchmarks



Subtype tests

[Palacz,Vitek ECOOP'03]



Subtype tests

- Subtype tests are queries over a poset of types, (T, \lessdot) . Write $A \lessdot B$ if B belongs to transitive closure of parents of A .
- Dynamic subtype tests are needed to ensure run time type safety.
- Subtype tests have unpredictable time and space requirements.

Subtype tests in Java

- Dynamic subtype tests are frequent operations in many Java programs, consider:

`x instanceof A`

`(A) x`

`A[0] = x`

`try{...}catch(E e){...}`

- Furthermore, the Generic Java compiler (gj) works by translating genericity away and inserting subtype tests to obtain verifying code.

Goals of this work

- **Predictability:** an algorithm for subtype tests in constant time and (almost) constant space
- **Incrementality:** Java allows dynamic class loading, any subtype test implementation must thus support incremental type hierarchy updates

Definitions

$A <: B$

holds if A is a subtype of B

isSubtype(A,B)

returns true if $A <: B$

checkCast(A,B)

throws an exception if not $A <: B$

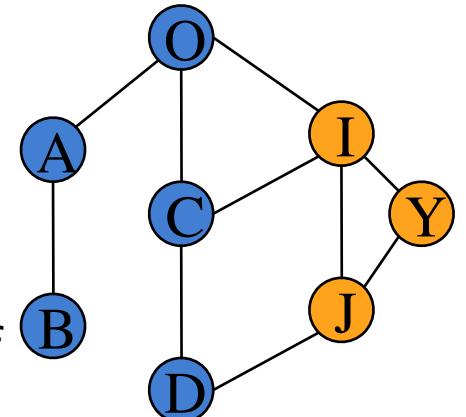
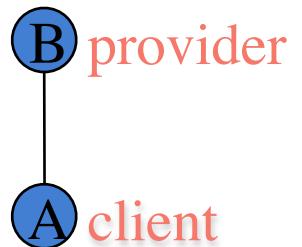
extends(A,B)

iff $A <: B$ and B is a class

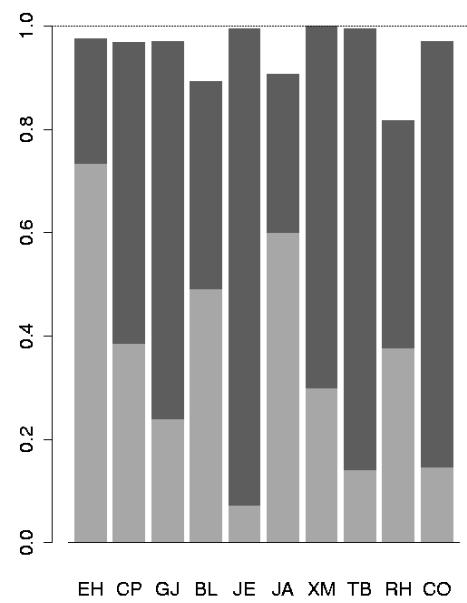
implements(A, B) iff $A <: B$ and B is an interface

a subtype test is **successful** if **isSubtype** returns true or if **checkCast** does not throw an exception

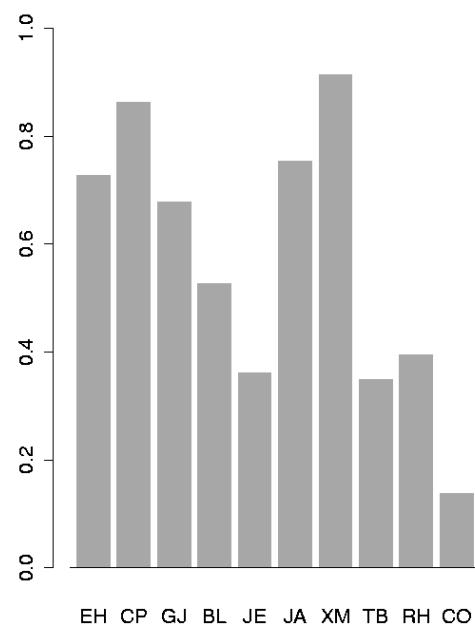
$A <: A$ is called a selftest (always true)



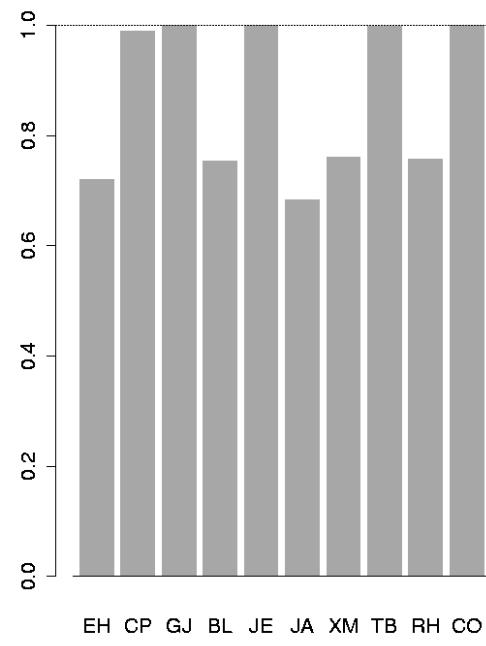
Dynamic characteristics of ST



Success rate



Checkcast rate
as % of dynamic tests



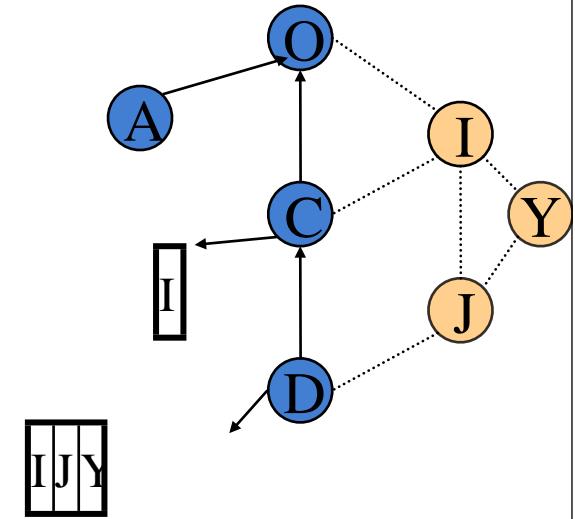
Extends rate

Observations

- Tests are mostly **successful**
 - ◆ A failed checkcast is often the last thing you do
 - ◆ Of course some programming styles invalidate this
- Selftests excepted, most types never in **provider** position
- Tests mostly **extends** tests
- **checkast** are highly dependent on programming style

Hierarchy traversal

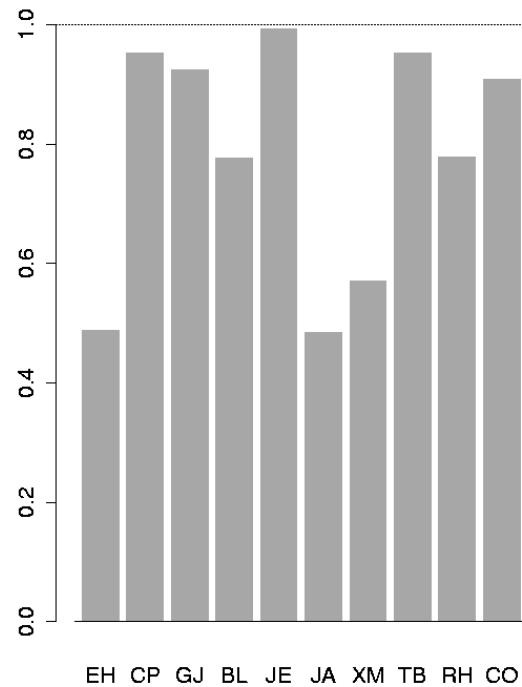
```
bool isSubtype(type cl, type pr) {  
    if (cl.cache==pr) return true;  
    if (isInterface(pr)) return extends(cl, pr);  
    else return implements(cl, pr); }  
  
bool implements(type cl, type pr) {  
    for (int i=0; i < pr.interfaces.length; i++)  
        if (cl==pr.interfaces[i]) {cl.cache=pr;return true;}  
    return false; }  
  
bool extends(type cl, type pr) {  
    for (type sp=cl.parent; sp!=null; sp=sp.parent)  
        if (sp==pr) {cl.cache=pr; return true;}  
    return false; }
```



Truth in advertising

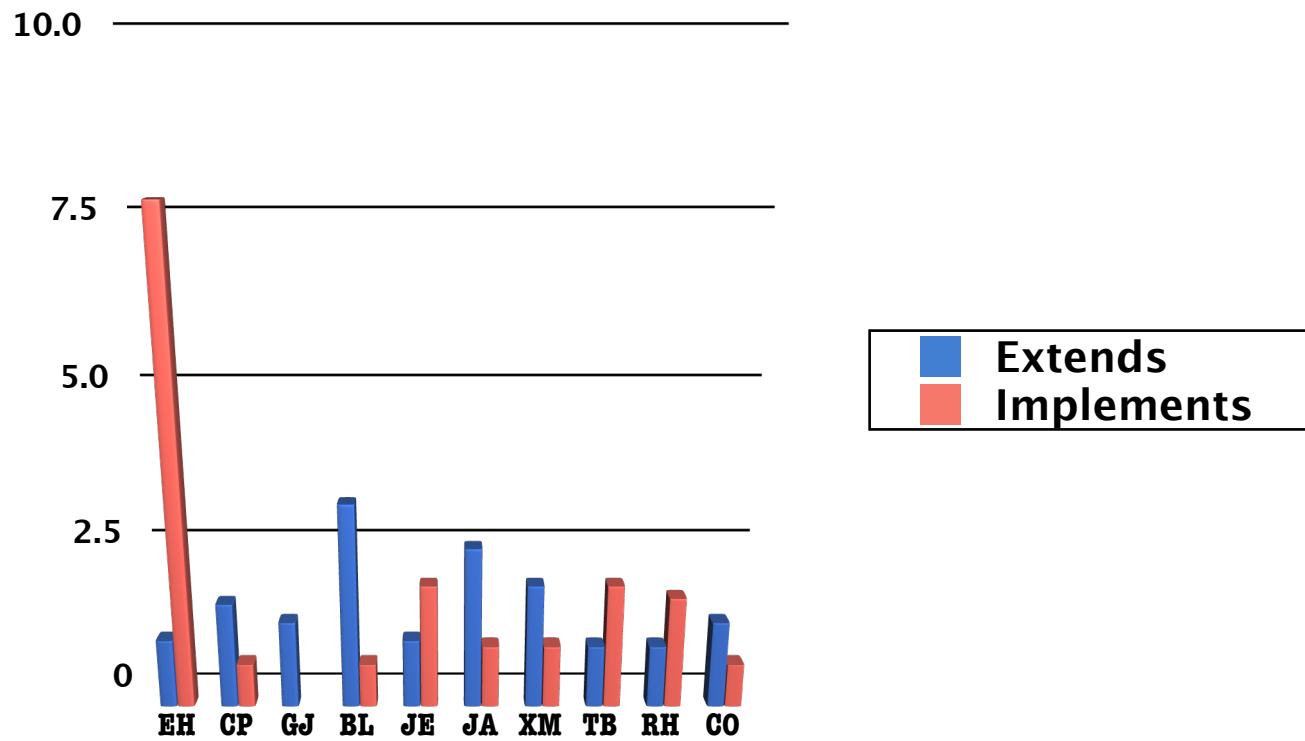
- The real subtype test procedure a bit more complex:
 - ◆ Cache test is JITed at the call site
 - ◆ Some interfaces may not be loaded
 - ◆ Arrays require comparison of element types
 - ◆ Some interfaces are treated specially (Cloneable and Serializable)
 - ◆ EVM uses two cache slots, one for array stores and one for other tests
 - ◆ `null instanceof A` throws an exception while `(A)null` is successful, go figure...

Evaluating the cache



Cache hit rates
as % of dynamic tests

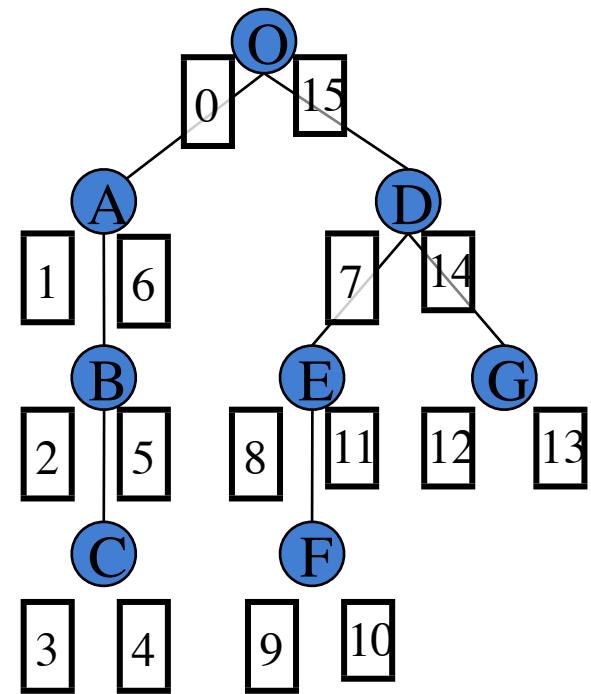
Hierarchy traversal cost



Avg number of iterations of extends and implements
loop without a cache.

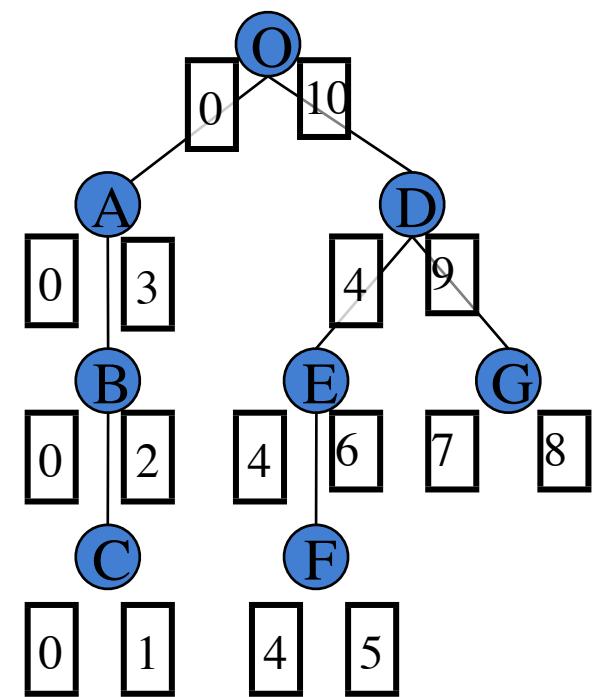
Constant time/space Extends

- Relative numbering
 - ◆ $A.\text{low} < B.\text{low} \&\& B.\text{high} < A.\text{high}$



Extends

```
bool extends(type cl, type pr) {  
  
    if (pr.low <= cl.low) {  
        if (cl.low < pr.high)  
            return true;  
    }  
    return false;  
}
```

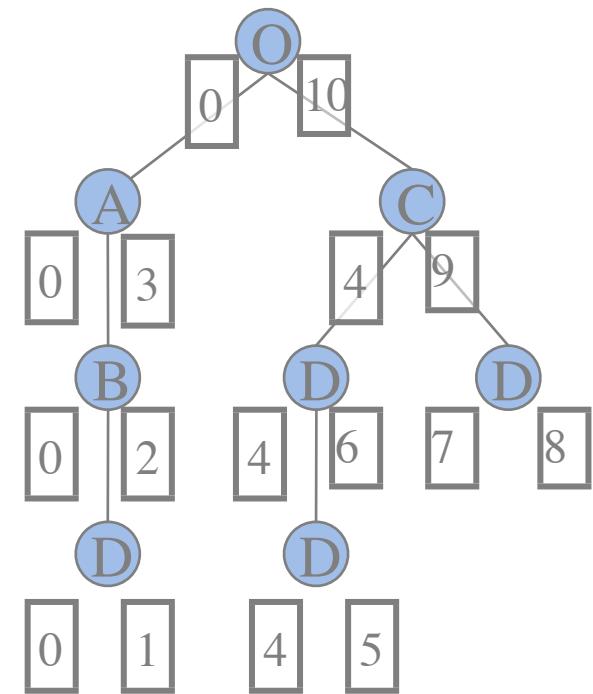


Incremental computation

- Recall:
 - ◆ Most types are not in provider position
 - ◆ Success is the common case

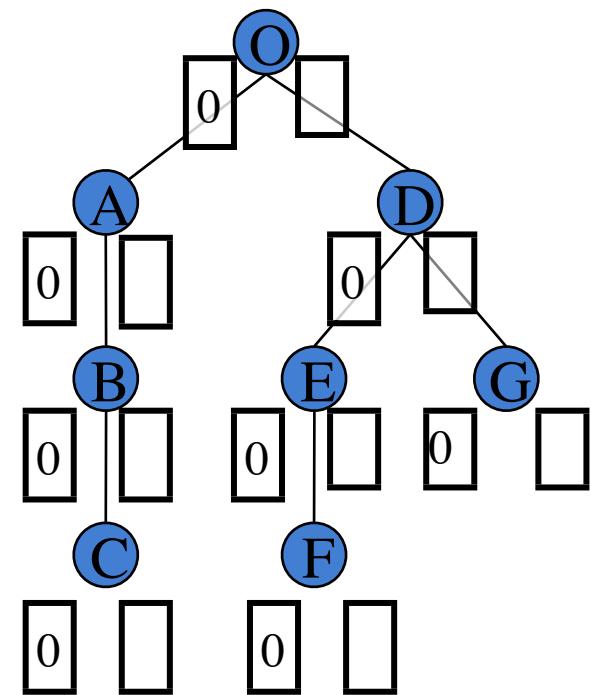
- Define:

```
bool isValid(type cl) {
    return (cl.high == -1);
}
```



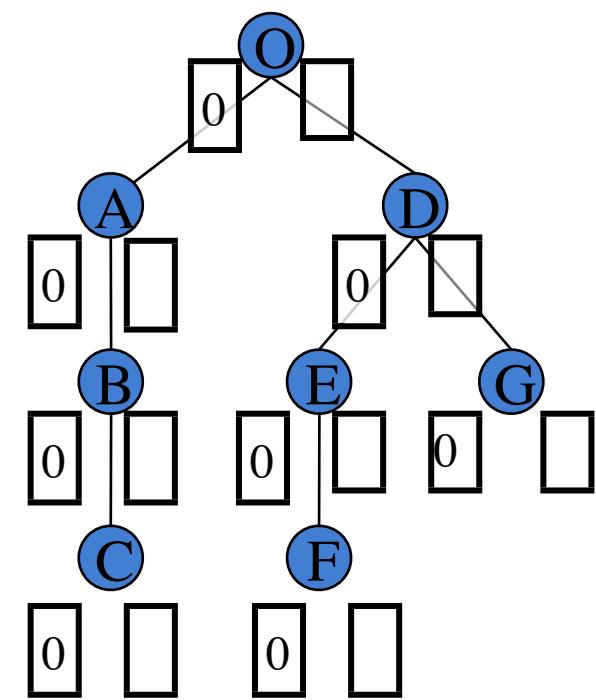
Incremental computation

```
bool extends(type cl, type pr) {  
    if (pr == cl) return true;  
    if (pr.low <= cl.low) {  
        if (cl.low < pr.high)  
            return true;  
        if (isValid(pr)) {  
            recompute(pr);  
            return extends(cl, pr);  
        }  
    }  
    return false;  
}
```



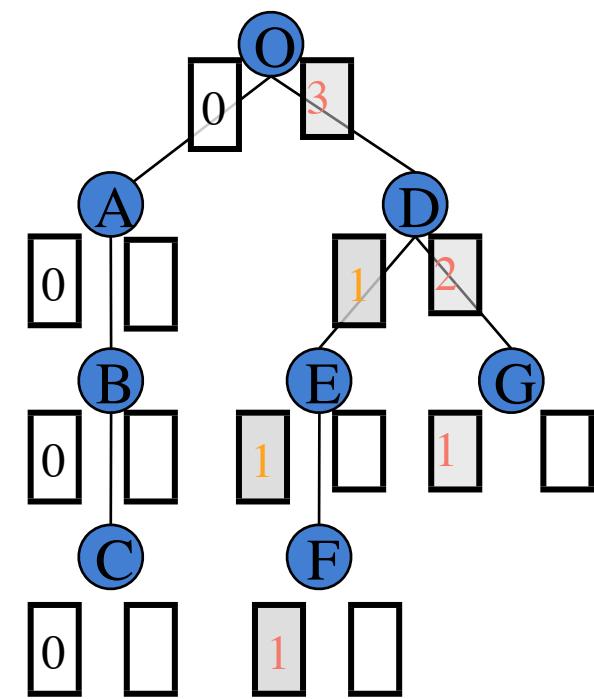
Incremental computation

extends (F ,D) ?



Incremental computation

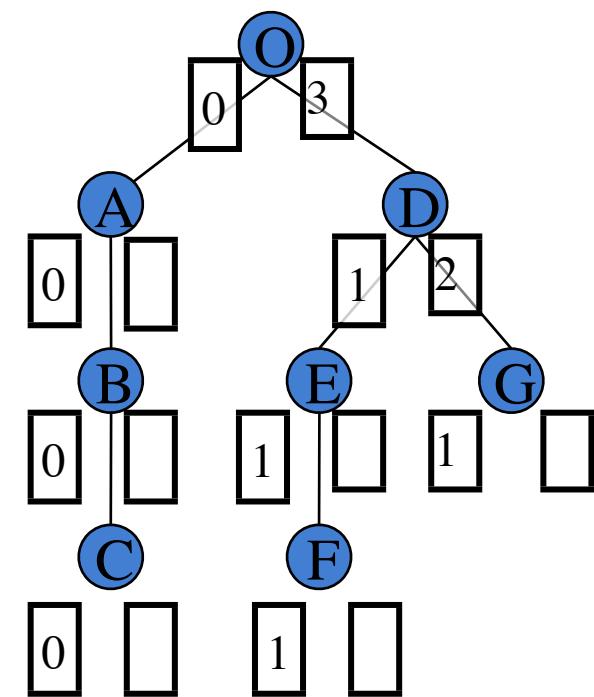
extends (F,D) ?



Incremental computation

extends (F , D)

extends (G , D) ?

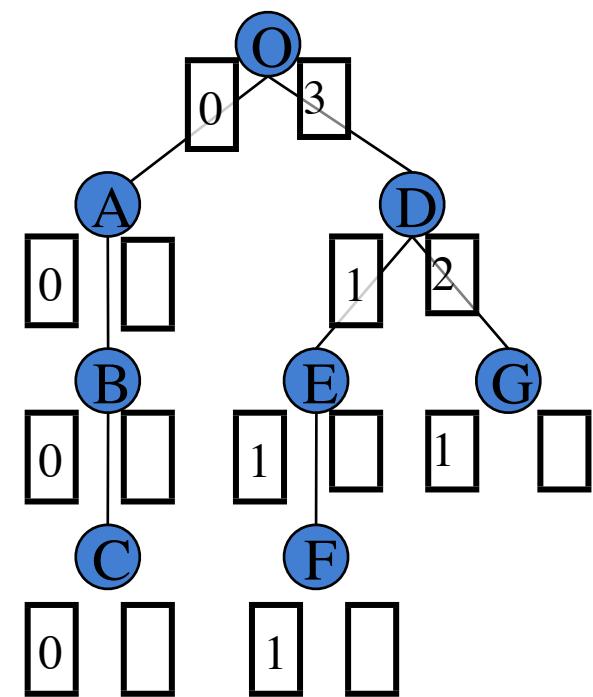


Incremental computation

extends (F ,D)

extends (G ,D)

extends (F ,E) ?



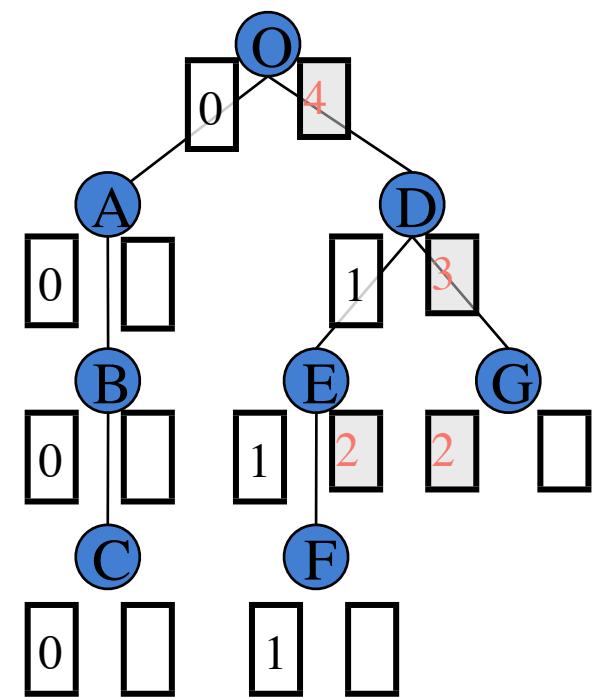
Incremental computation

extends (F ,D)

extends (G ,D)

extends (F ,E)

extends (C ,C) ?



Incremental computation

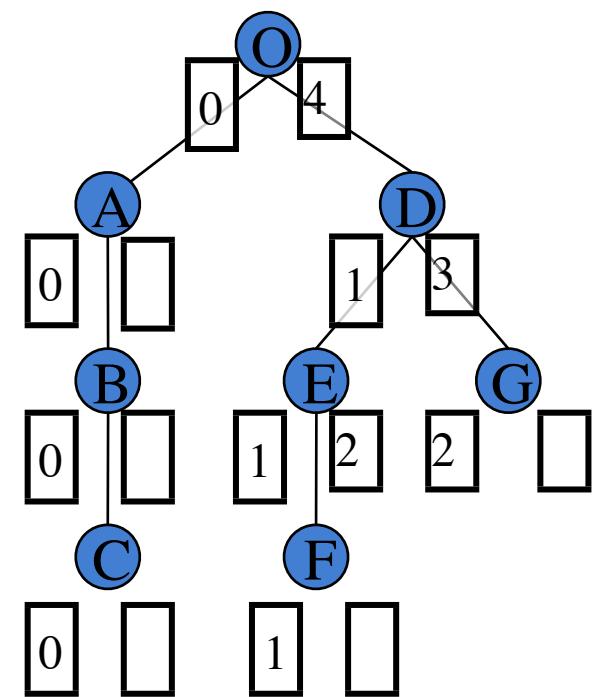
extends (F ,D)

extends (G ,D)

extends (F ,E)

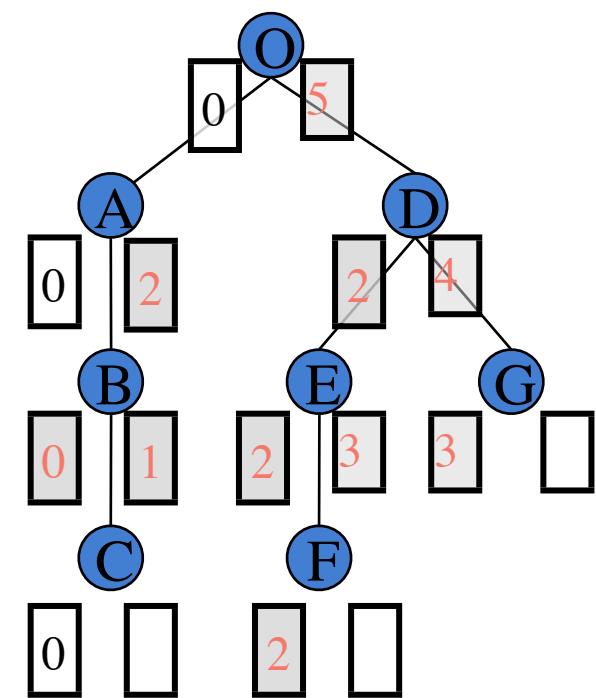
extends (C ,C)

extends (C ,B) ?



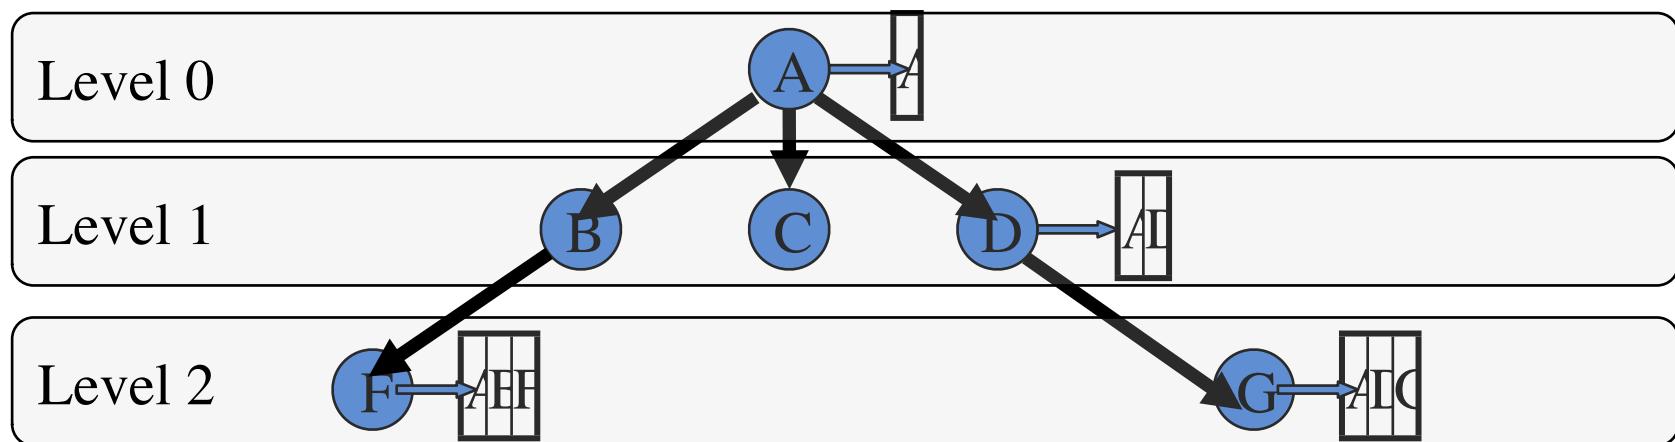
Incremental computation

extends (F ,D)
extends (G ,D)
extends (F ,E)
extends (C ,C)
extends (C ,B)



Implements

Cohen's displays for single inheritance...



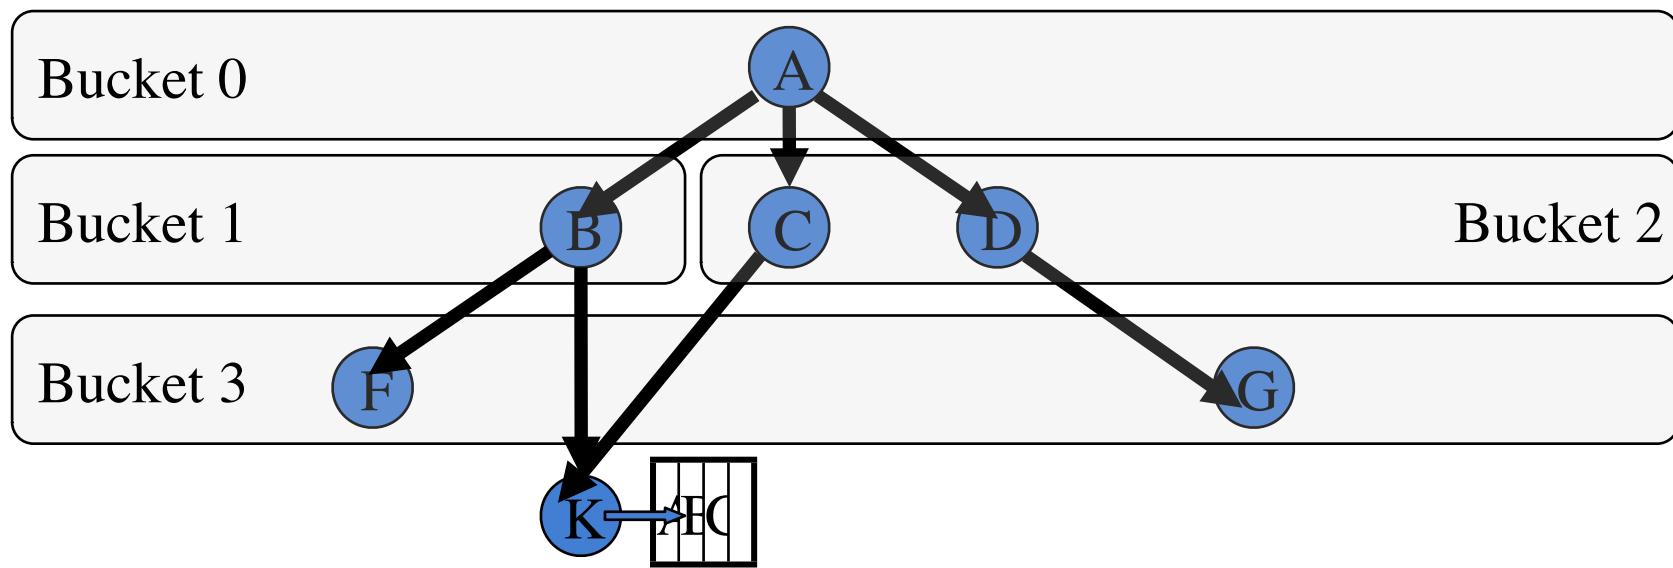
$x <: y$

iff

`x.display.length < y.level && x.display[y.level] == y.id`

Multiple subtyping

Interfaces in the same bucket have no common descendants



$x <: y$

iff

`x.display[y.bucket_num] == y.id`

Display Updates

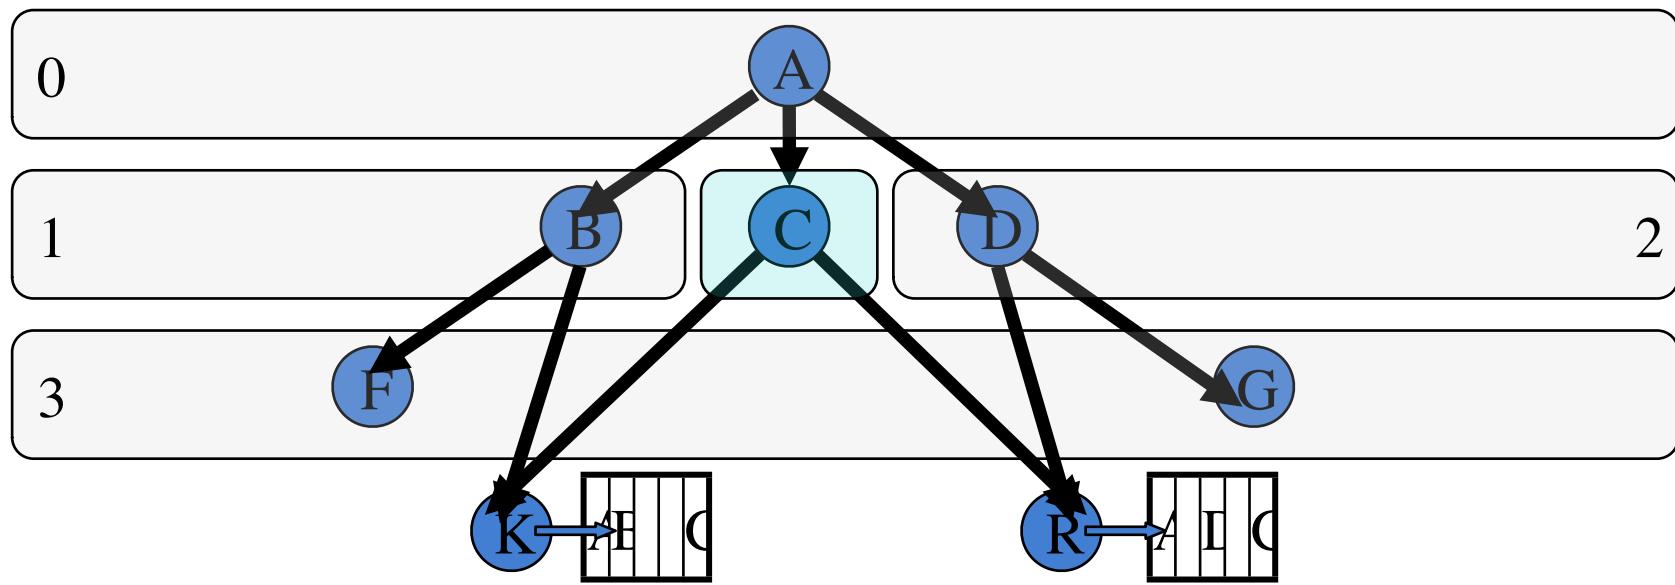
- Maintain invariant:

$$\begin{aligned} I.\text{bucket_num} == J.\text{bucket_num} \\ \Leftrightarrow \\ \text{no } C \text{ such that } C < : I \text{ \&\& } C < : J \end{aligned}$$

- When loading class **A**

```
Foreach I: if A < : I do  
    mark[I.bucket_num]++  
  
Foreach b: if mark[b_num]==k && k<1 do  
    split b into k buckets of same size,  
    assign A's interfaces to different buckets  
recompute displays
```

Class Loading



Implements, ct'd

- When loading interfaces: which bucket to add to?
 - ◆ choose smallest bucket among m most recently created buckets (for us $m=5$)
 - ◆ heuristics: avoids putting system interfaces together with application-defined interfaces while attempting to balance bucket sizes (likely to conflict)
- Simple but good enough
(display size only twice maximum number of implemented interfaces)

Summary

- We have presented two simple techniques for dealing with extends and implements tests
- Some heuristics are particularly tuned to the characteristics of Java
- We have implemented these techniques in the SUN Research VM (aka EVM)

Results

