

Garbage Collection



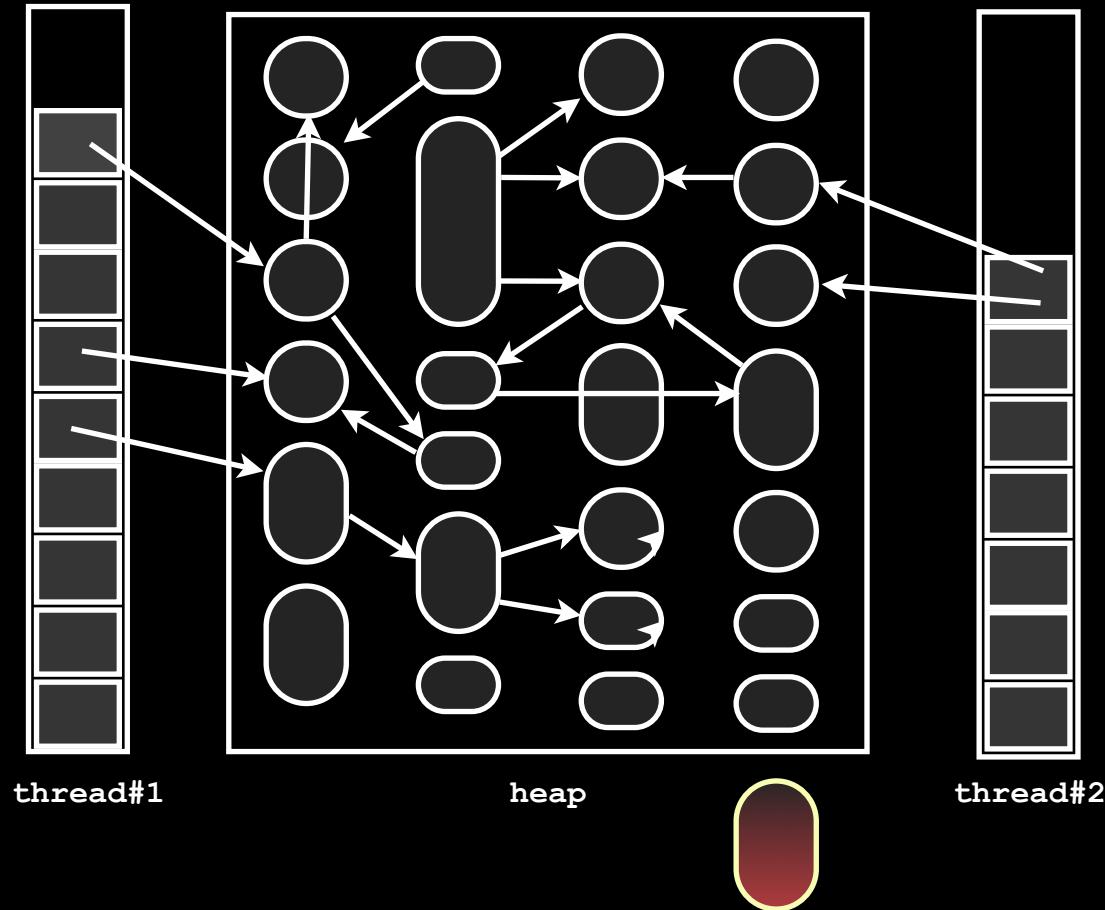
Garbage Collection

- A Garbage Collector (GC) is an algorithm that automatically finds unused objects in the memory of an application and prepares them for reuse
- GC frees programmers from worrying about the exact lifetime of objects and ensures that the heap will not be corrupted by access to previously freed data
- ... *but introduces pauses that may be $O(\text{heap})$ and can increase the memory required.*
Moreover, pauses occur at unpredictable times, especially in concurrent programs

Garbage Collection

Phases

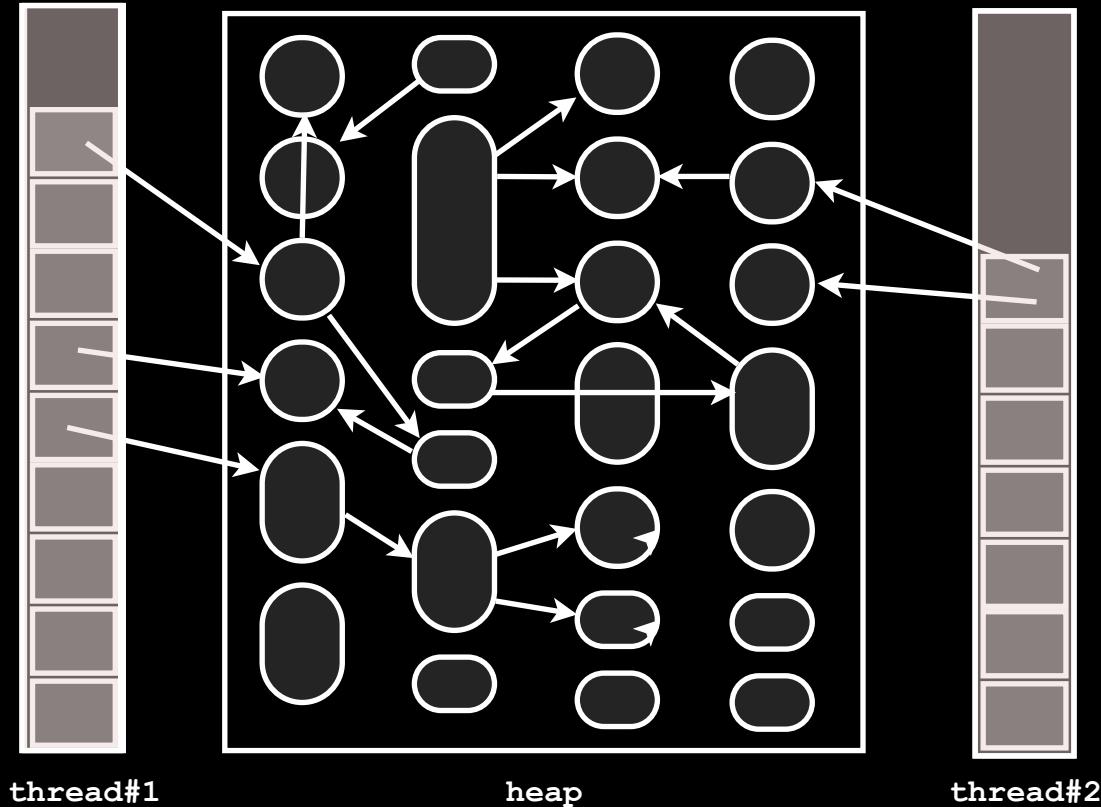
- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction



Garbage Collection

Phases

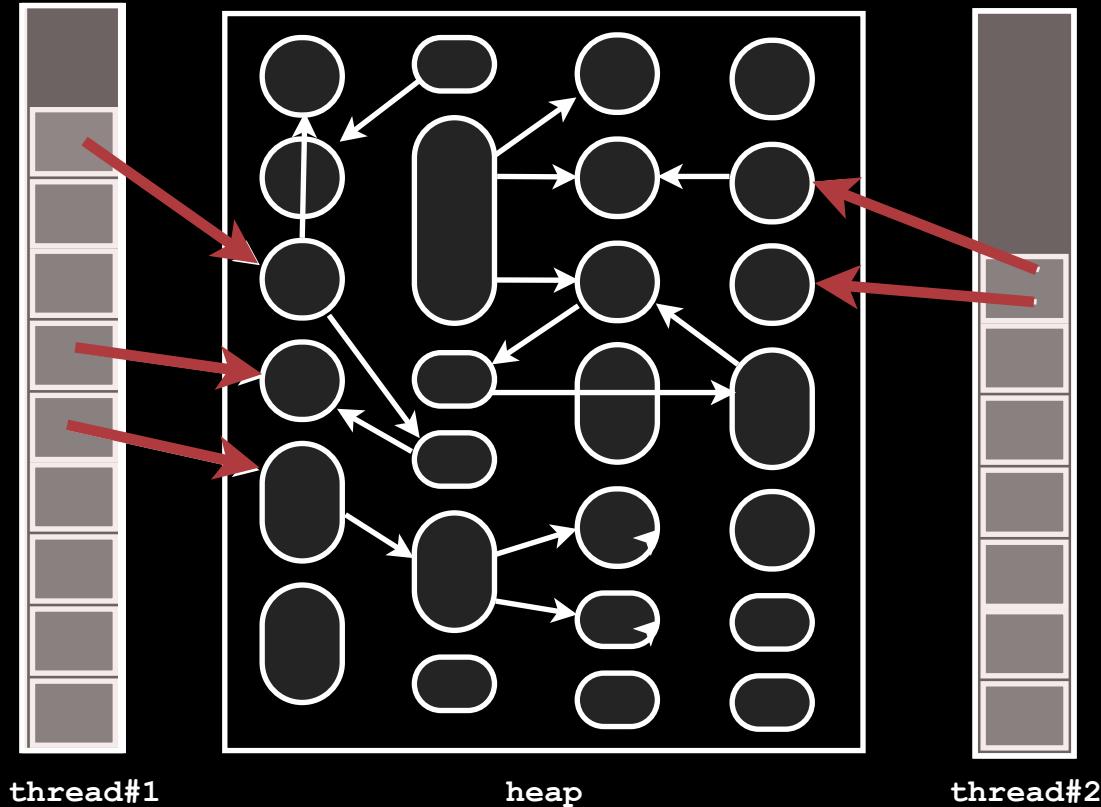
- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction



Garbage Collection

Phases

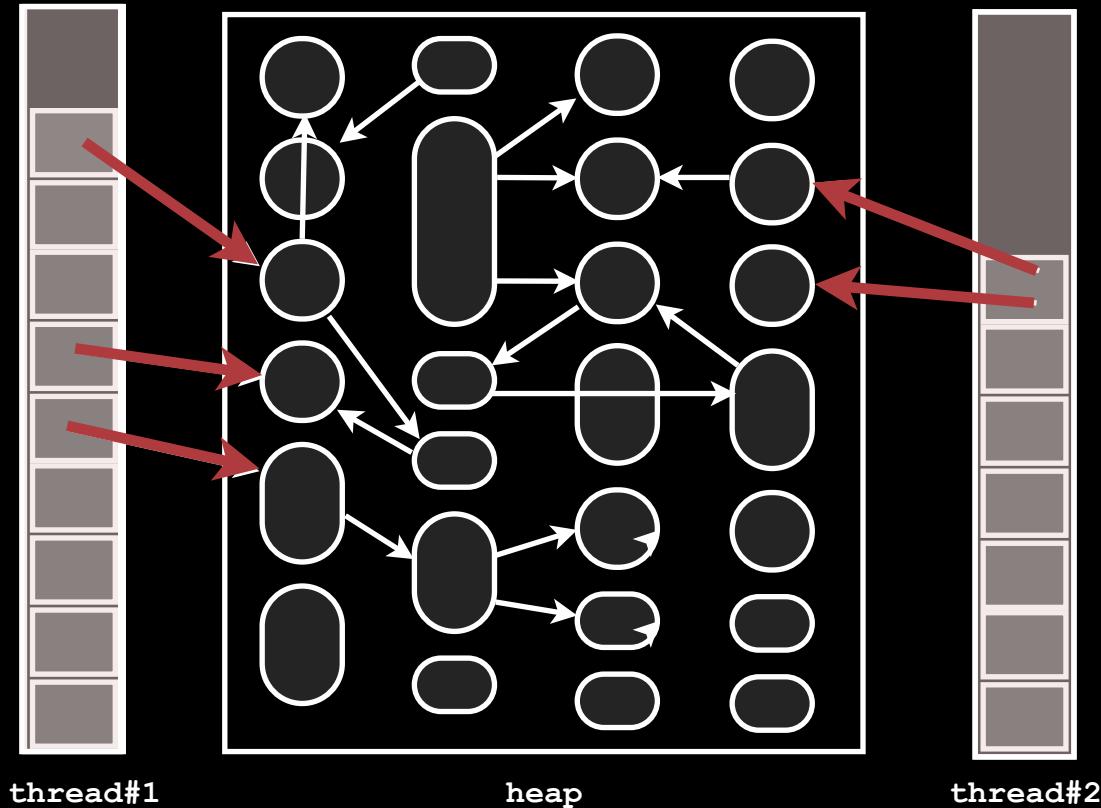
- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction



Garbage Collection

Phases

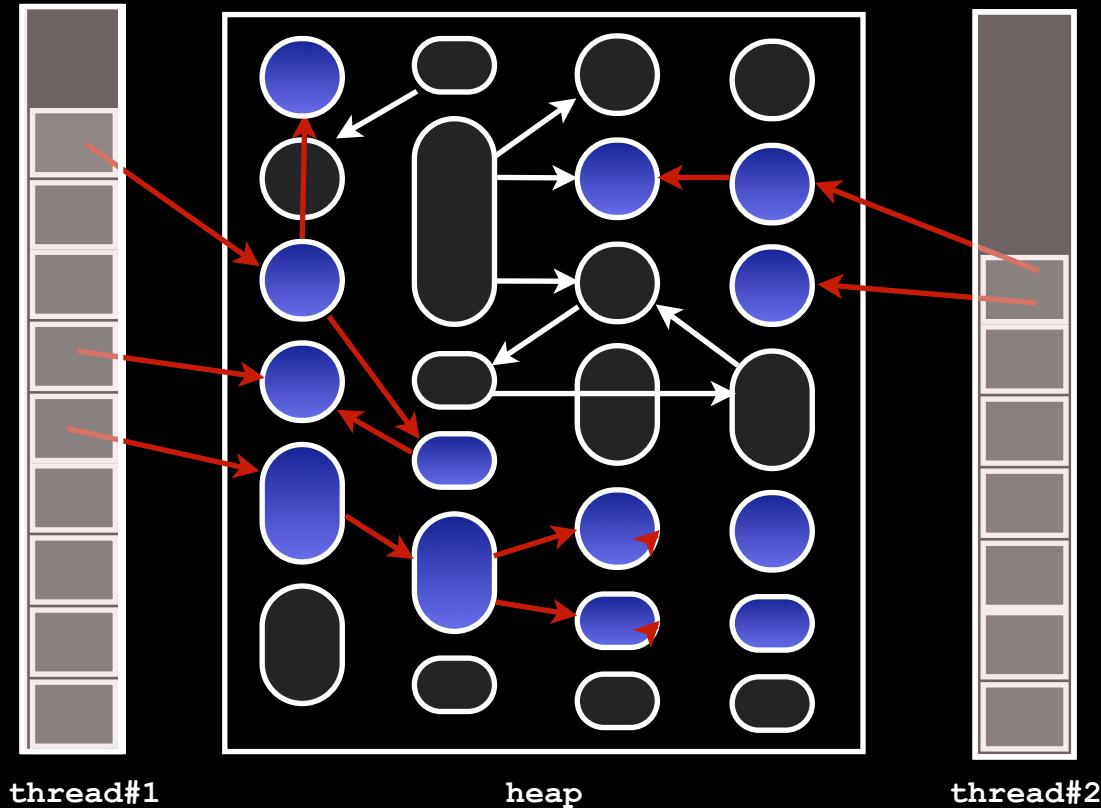
- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction



Garbage Collection

Phases

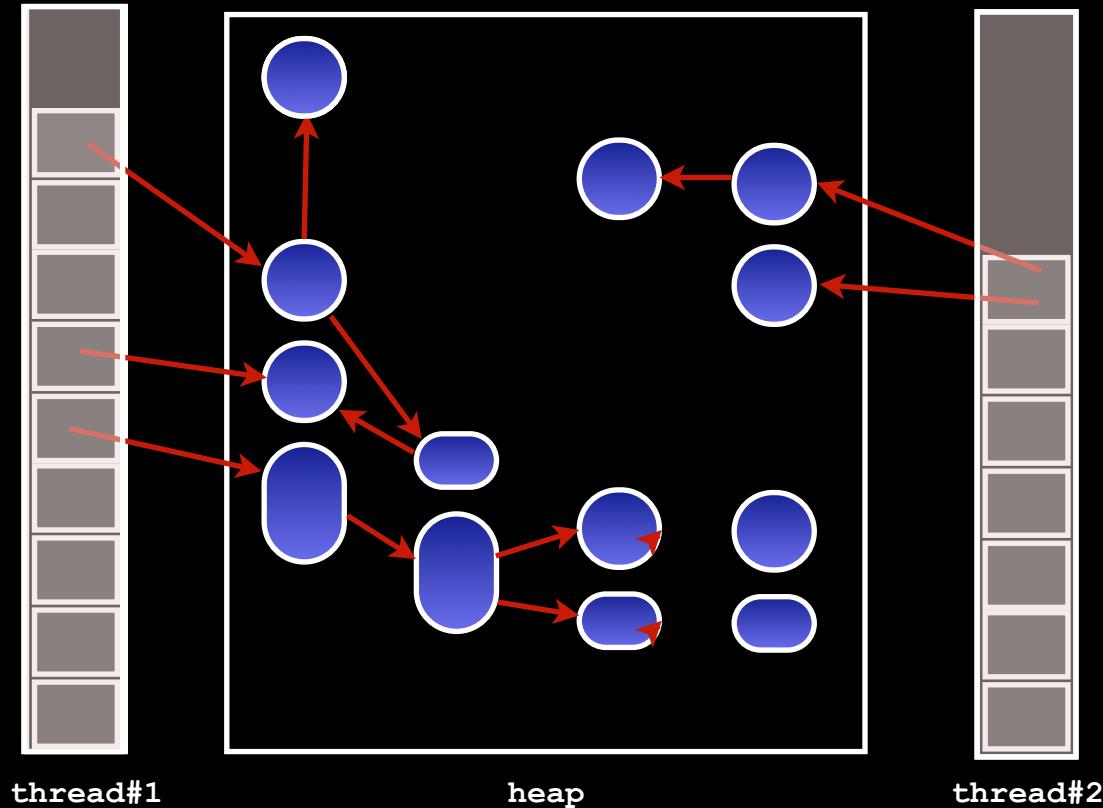
- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction



Garbage Collection

Phases

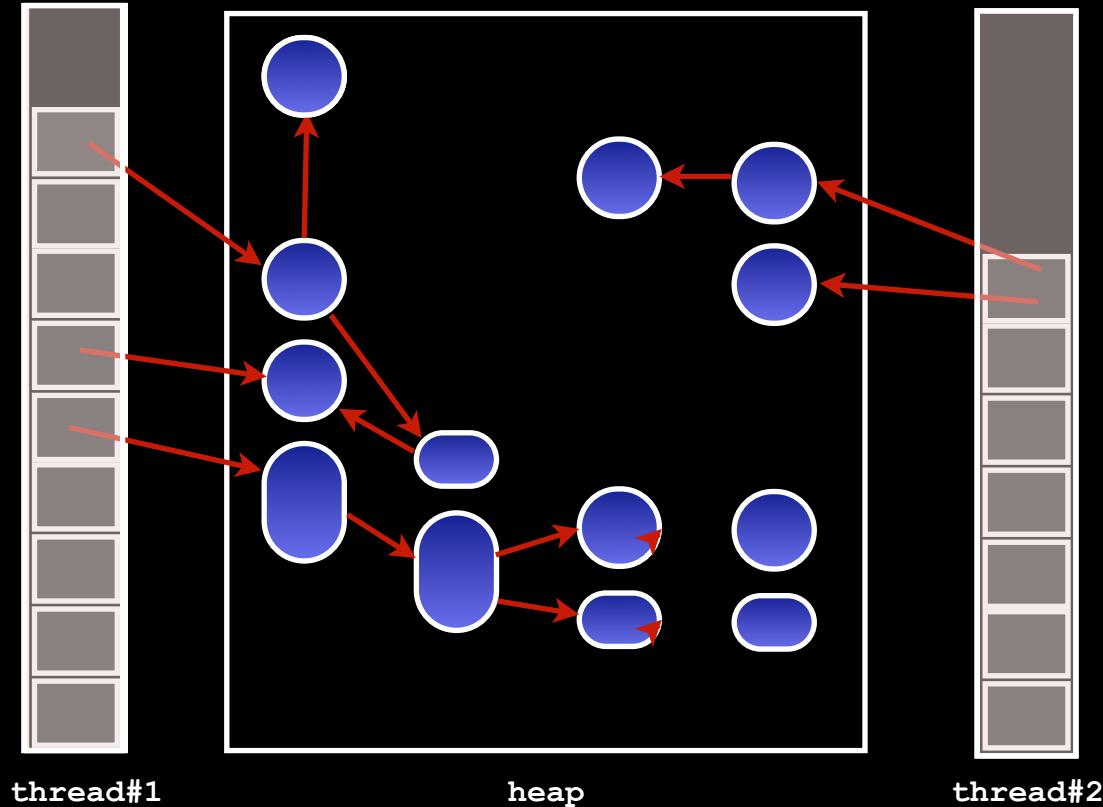
- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction



Garbage Collection

Phases

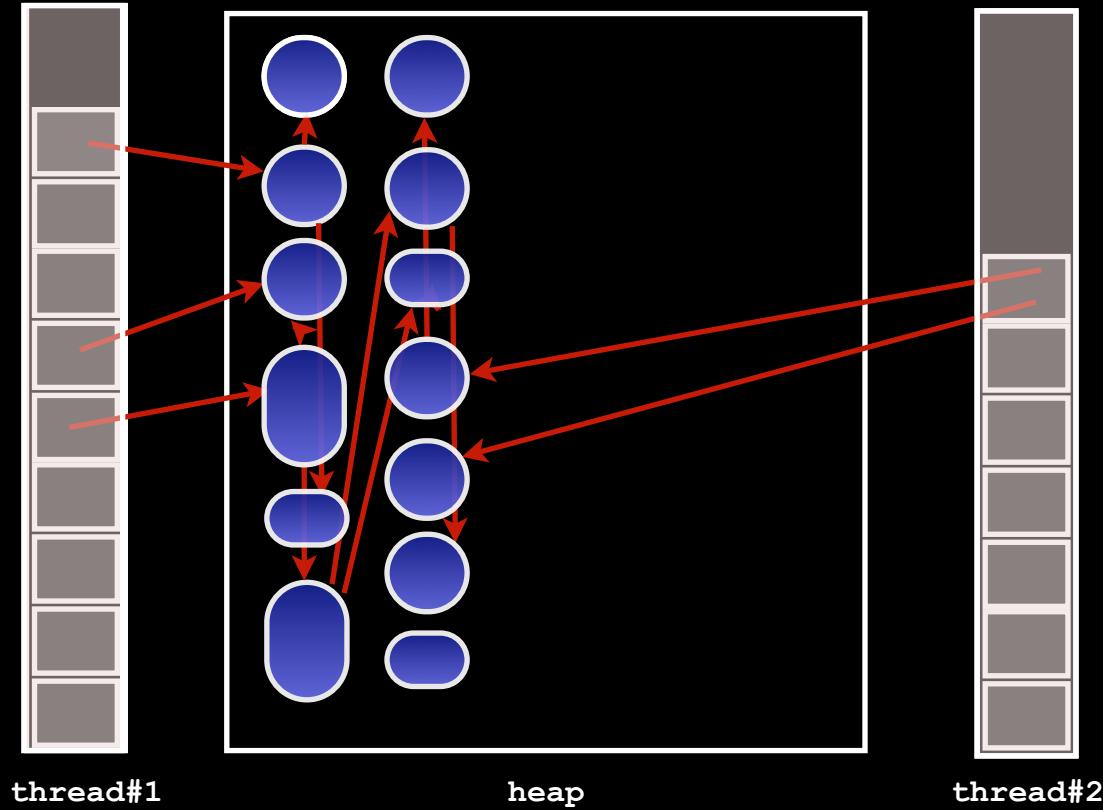
- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction



Garbage Collection

Phases

- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction



GC is Easy

- If responsiveness is not an issue,
the GC can complete in one long pause under the
assumption that there is no interleaved application
activity
- Marking is easy
if the graph does not change while you are searching it.
- Copying/compacting objects and fixing up the heap is
easy
if the application is prevented from accessing the heap

Real-time Garbage Collection



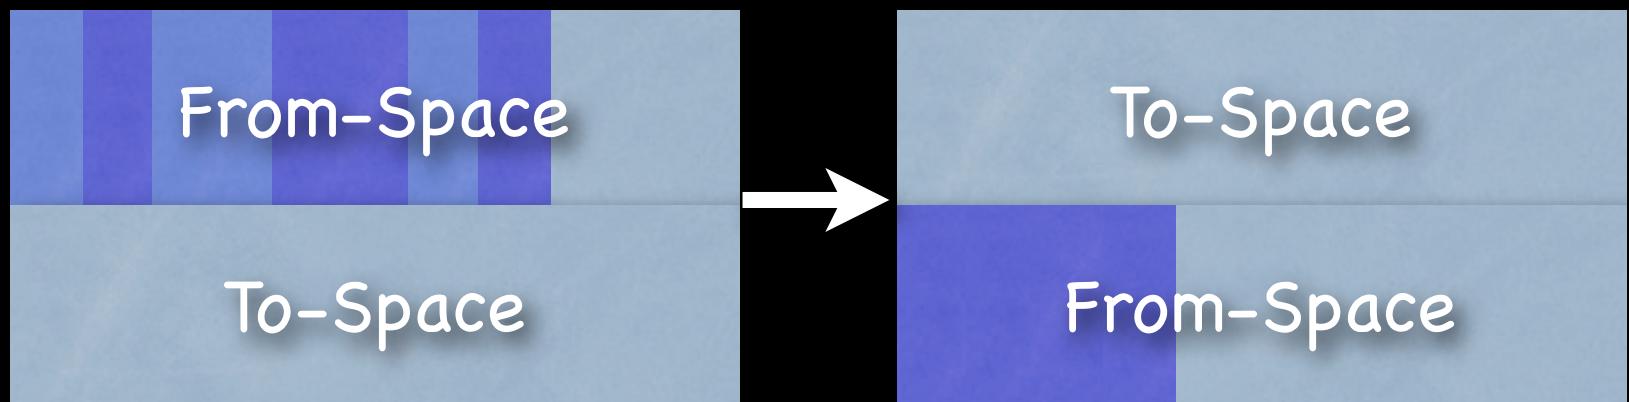
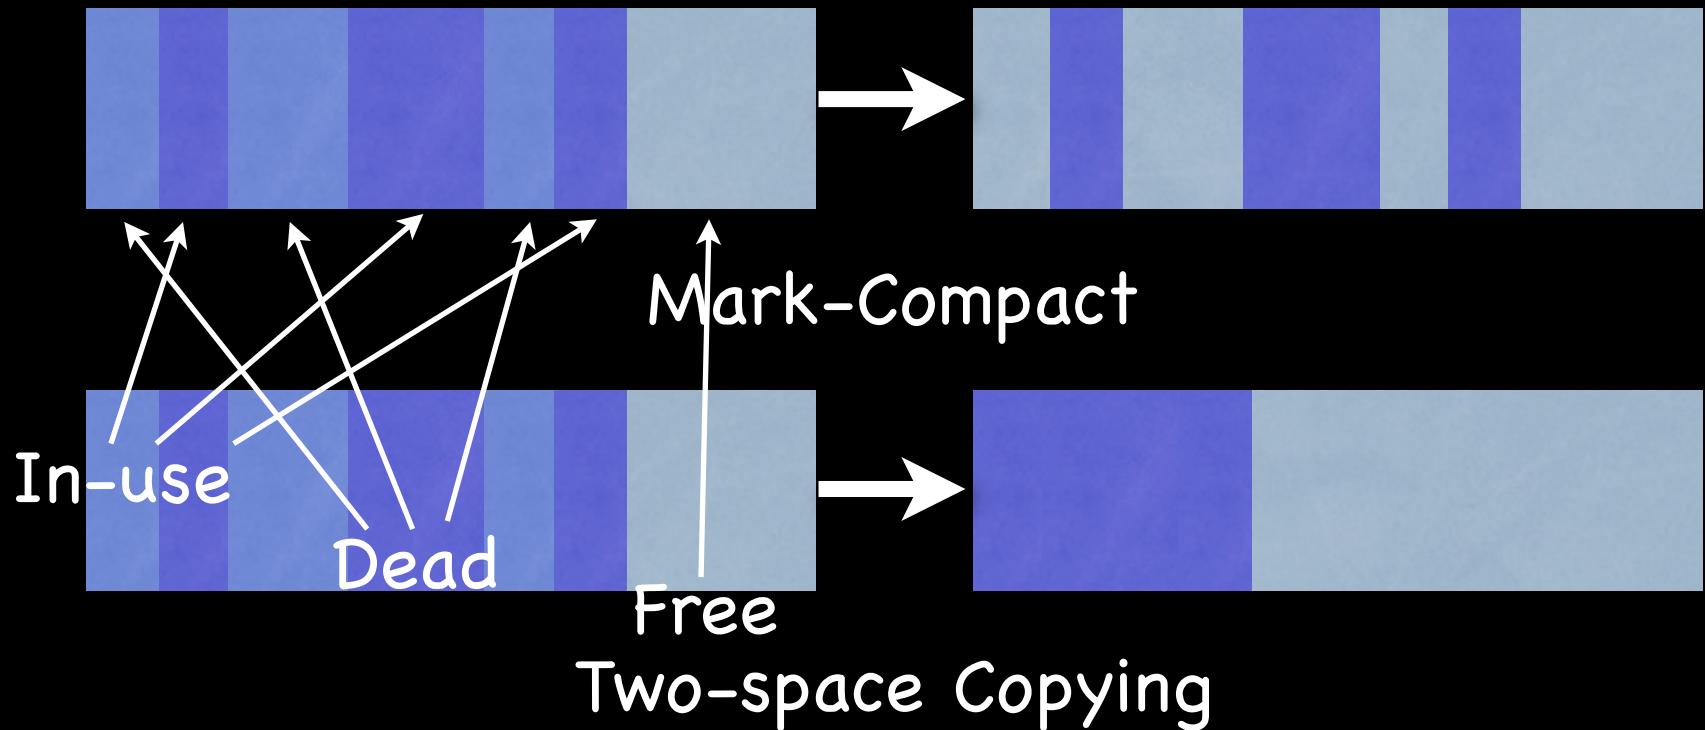
Real-time GC

- A Real-time GC must provide time and space predictability
 - ▶ provide a performance model that can be used to guarantee that programs do not run out of memory or experience pauses that violate their timing constraints
- A Real-time-GC must support defragmentation of the heap if it is to be used with long-lived applications
- Multi-processor support is unavoidable
- Throughput should not degrade overly

Main Collector Types

- Reference counting
 - ▶ keep count of incoming references. Requires auxiliary GC to reclaim cycles.
- Mark-Sweep
 - ▶ one phase to mark used objects, another to sweep unused space. No copying is performed.
- Mark-Compact
 - ▶ one phase to mark used objects, another to copy all used objects to one side of the heap.
- Copying
 - ▶ one phase to simultaneously mark and evacuate all used objects; the space that previously held all objects is then completely free.

Mark-Sweep



RTGC: state of the art

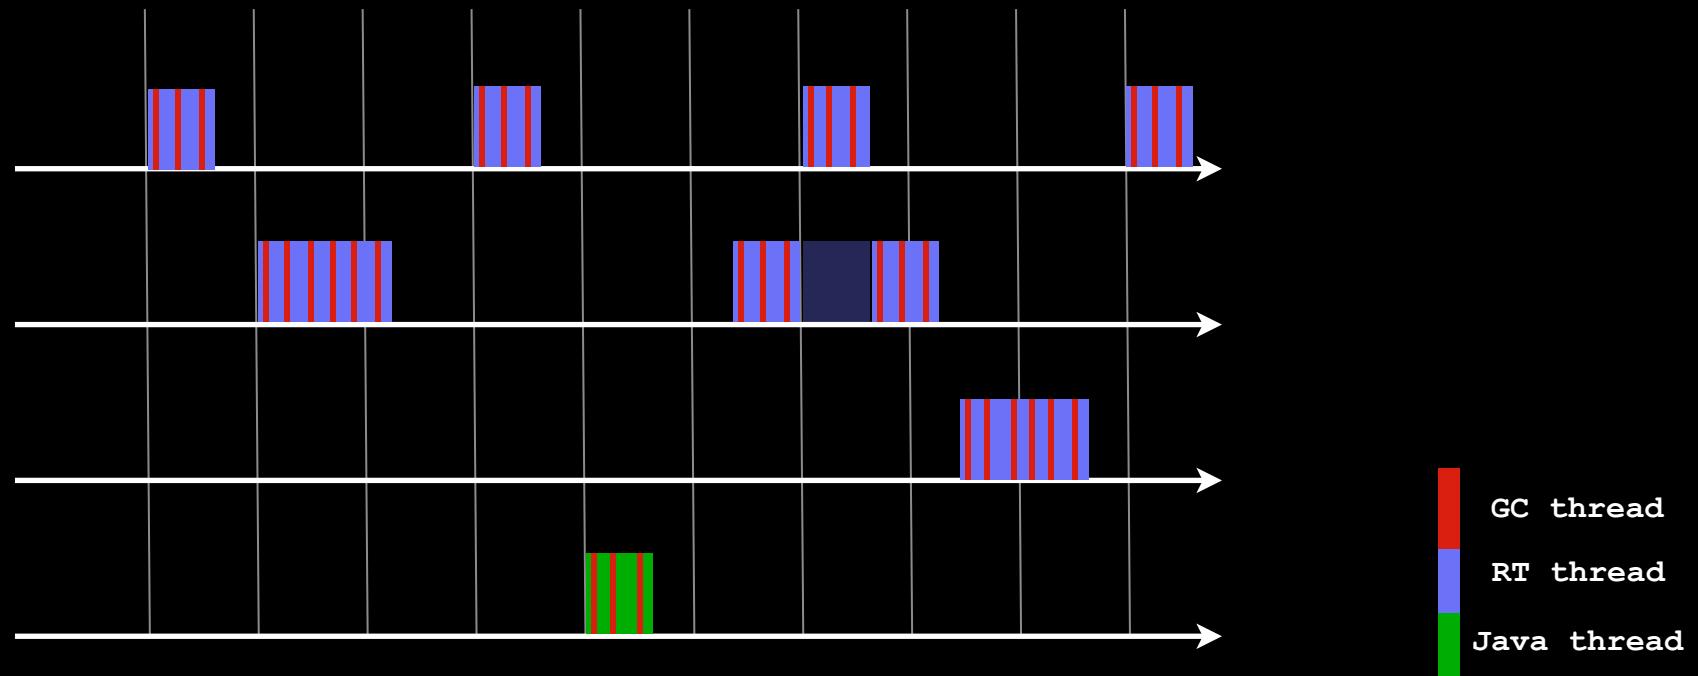
- Sort-of Real-time GCs:
 - ▶ Too many to list, Baker (1978) was first
- Really Real-time GCs:
 - ▶ JamaicaVM (1999) - *work-based*
 - ▶ Henriksson (1998) - *slack-based*
 - ▶ Metronome (2003) - *time-based*

Work-based RTGC

Baker'78, Siebert'99



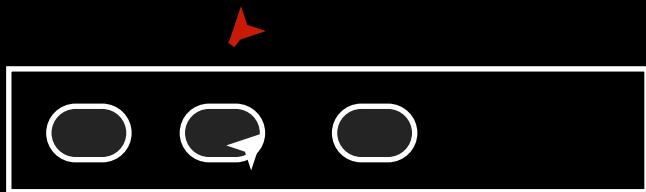
Work-based scheduling



Baker

- First attempt at RTGC by Baker '78
- For LISP where all objects are the same, small, size
- Copying collector that instruments the mutator to maintain:
 - ▶ To-space invariant:
 - mutator uses to-space, never sees from-space,
objects are copied on first read
 - ▶ Steady state:
 - collector keeps up with mutator allocations,
collection work performed on each allocation

Incrementalizing marking



Collector marks object



Application updates
reference field



Compiler inserted
write barrier marks object



Dissecting Baker

- The chief problem with Baker is the to-space invariant
 - ▶ Reads are frequent, a heavy read barrier leads to poor performance
 - ▶ Variable sized objects in Java make costs harder to bound
- What RT programmers fear:
 - ▶ a burst of copying reads, causing substantial slow-downs

Siebert

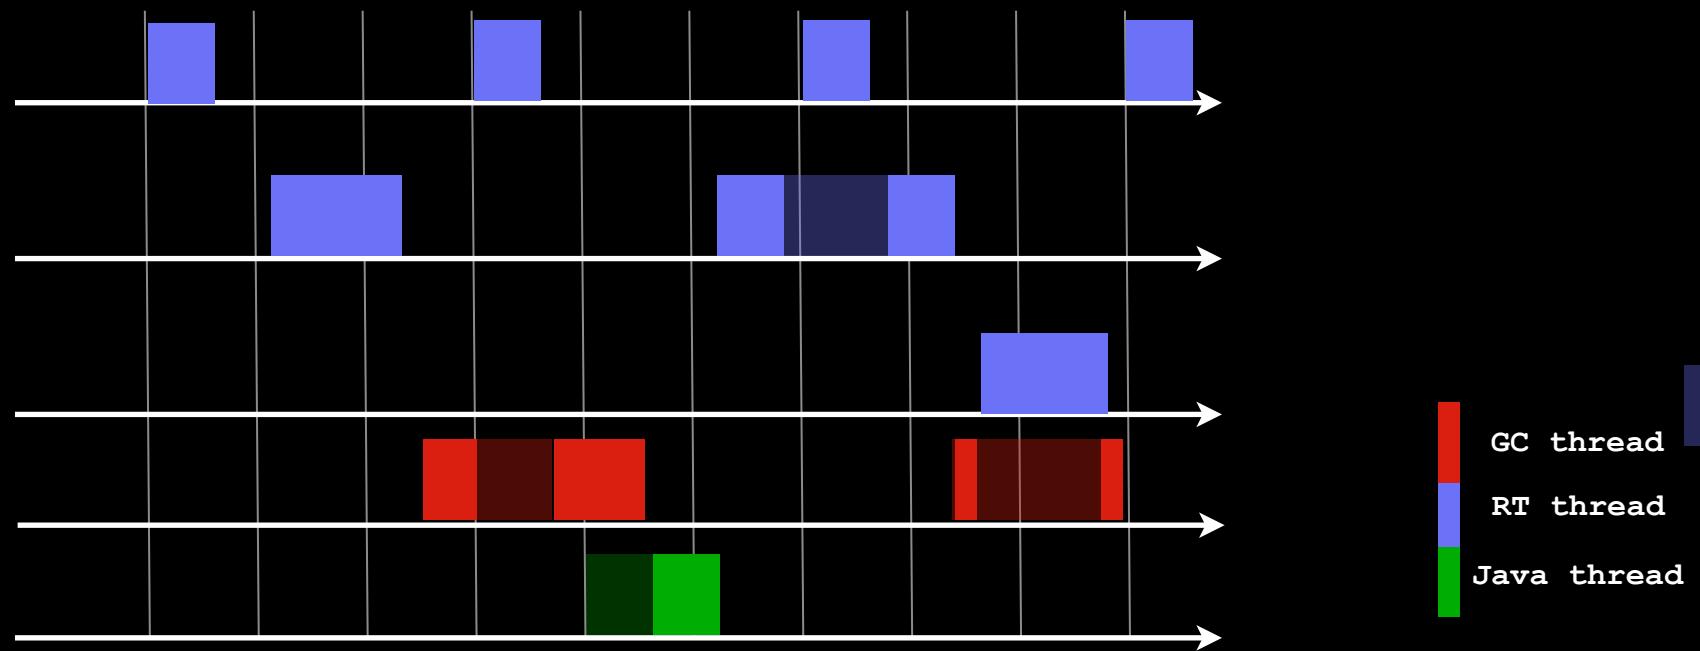
- Modern work-based collector in the JamaicaVM:
 - ▶ Mark-Sweep, no copying
 - ▶ All objects are same size, 32 bytes
(large objects \Rightarrow lists, arrays \Rightarrow tries) Stack is logically an object
 - ▶ Root scanning is fully incremental
 - ▶ Write barrier to mark objects (including the stack)
 - ▶ Allocation triggers a bounded amount of collector work

Slack-based RTGC

Henriksson'98



Slack-based GC Scheduling

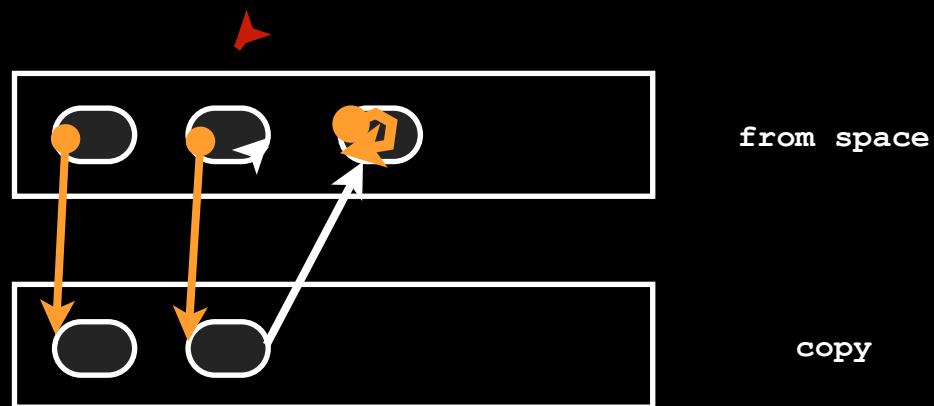


Slack-based GC

- Henriksson takes a different approach to real-time guarantees:
 - ▶ tries to guarantee real-time tasks never experience GC interference
- How is this done?
 - ▶ RT tasks can always preempt GC
 - ▶ Objects have normal representation.
 - ▶ Read & write barriers are used to ensure marking/copying soundness.
 - ▶ Copying collector; with roll backs when the mutator preempts GC

Incrementalizing copying/compaction

- Forwarding pointers are used to refer to the current version of the object.
- Every access must start with a dereference



Dissecting Henriksson

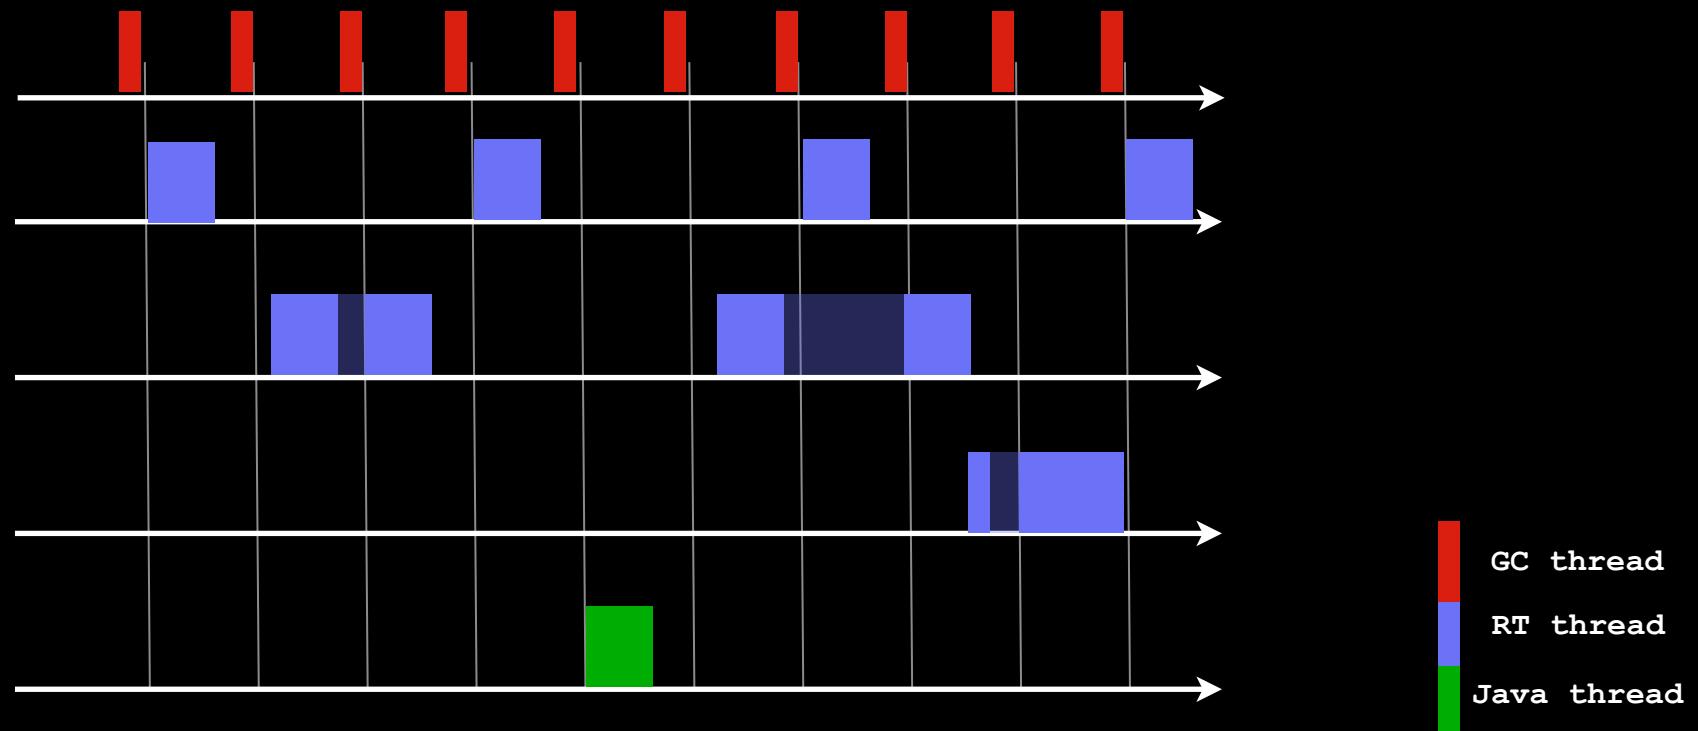
- Pros:
 - ▶ Gives the programmer full control over the GC schedule
 - ▶ SUN uses a variant of this collector in their RTSJVM
- Cons:
 - ▶ Read barrier are expensive
 - ▶ Mutator induced copy aborts make collector progress hard to ensure
 - ▶ Programmer must know allocation rate, and to budget for GC work
 - ▶ Priority inversion can lead to GC interference

Time-based RTGC

Metronome'03



Time-based GC Scheduling



Time-based GCs

- Metronome runs periodically for a bounded amount of time
 - ▶ Amount of collector interference is easy to understand
 - ▶ Progress is easier to guarantee than in Henriksson
- A Mark-Sweep-and-sometimes-Copy collector:
 - ▶ Write barrier to ensure marking soundness
 - ▶ Light read barrier to ensure copying soundness
 - ▶ Almost-empty pages evacuated in response to fragmentation
 - ▶ Arrays are split into tries with page-sized leaves
 - ▶ Collector increments lower-bounded by time to copy a page-size objects

Suming up



Siebert

Local ref. assignment	$\text{mark}(b)$ $a = b$
Primitive heap load	$a = b \rightarrow f$
Reference heap load	$a = b \rightarrow f$ $\text{mark}(a)$
Primitive heap store	$a \rightarrow f = b$
Reference heap store	$\text{mark}(b)$ $a \rightarrow f = b$

Henriksson

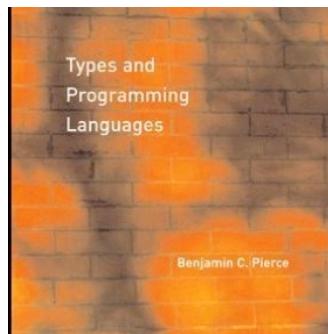
Local ref. assignment	$a = b$
Primitive heap load	$a = b \rightarrow \text{forward} \rightarrow f$
Reference heap load	$a = b \rightarrow \text{forward} \rightarrow f$ $\text{mark}(a)$
Primitive heap store	$a \rightarrow \text{forward} \rightarrow f = b$
Reference heap store	$\text{mark}(b)$ $a \rightarrow \text{forward} \rightarrow f = b \rightarrow \text{forward}$

Metronome

Local ref. assignment	$a = b$
Primitive heap load	$a = b \rightarrow \text{forward} \rightarrow f$
Reference heap load	$a = b \rightarrow \text{forward} \rightarrow f$
Primitive heap store	$a \rightarrow \text{forward} \rightarrow f = b$
Reference heap store	$\text{mark}(a \rightarrow \text{forward} \rightarrow f)$ $a \rightarrow \text{forward} \rightarrow f = b \rightarrow \text{forward}$

The Essence of JavaScript

by Arjun Guha, Claudiu Saftoiu, and
Shriram Krishnamurthi

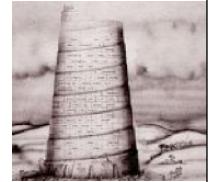


<http://www.cs.brown.edu/people/arjun/lambdajs>

A small step operational semantics

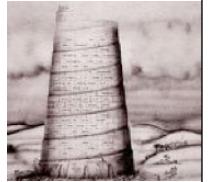


- λ_{JS} is a small-step, operational semantics for the core of JavaScript
- The full language is obtained by desugaring into λ_{JS}
- An implementation is available (in Haskell)
- Treatment of `with` is particularly interesting...
- Some security applications which we will not discuss here.



Functions and Objects

```
c = num | str | bool | undefined | null  
v = c | func(x...) { return e } | { str:v... }  
e = x | v | let (x = e...) e | e(e...) | e[e] | e[e] = e | delete e[e]  
E = • | let (x = v... x = E, x = e...) e | E(e...) | v(v... E, e...)  
| {str: v... str:E, str:e...} | E[e] | v[E] | E[e] = e | v[E] = e  
| v[v] = E | delete E[e] | delete v[E]
```



Functions and Objects

$$\text{let } (x = v \dots) e \hookrightarrow e[x/v] \dots \quad (\text{E-LET})$$

$$(\text{func}(x_1 \dots x_n) \{ \text{return } e \})(v_1 \dots v_n) \hookrightarrow e[x_1/v_1 \dots x_n/v_n] \quad (\text{E-APP})$$

$$\{ \dots \text{ str: } v \dots \}[\text{str}] \hookrightarrow v \quad (\text{E-GETFIELD})$$

$$\frac{\text{str}_x \notin (\text{str}_1 \dots \text{str}_n)}{\{ \text{str}_1: v_1 \dots \text{str}_n: v_n \} [\text{str}_x] \hookrightarrow \text{undefined}} \quad (\text{E-GETFIELD-NOTFOUND})$$

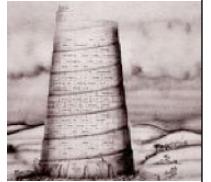
$$\begin{aligned} & \{ \text{str}_1: v_1 \dots \text{str}_i: v_i \dots \text{str}_n: v_n \} [\text{str}_i] = v \\ & \hookrightarrow \{ \text{str}_1: v_1 \dots \text{str}_i: v \dots \text{str}_n: v_n \} \end{aligned} \quad (\text{E-UPDATEFIELD})$$

$$\frac{\text{str}_x \notin (\text{str}_1 \dots)}{\{ \text{str}_1: v_1 \dots \} [\text{str}_x] = v_x \hookrightarrow \{ \text{str}_x: v_x, \text{str}_1: v_1 \dots \}} \quad (\text{E-CREATEFIELD})$$

$$\begin{aligned} & \text{delete } \{ \text{str}_1: v_1 \dots \text{str}_i: v_x \dots \text{str}_x: v_n \} [\text{str}_x] \\ & \hookrightarrow \{ \text{str}_1: v_1 \dots \text{str}_i: v \dots \text{str}_n: v_n \} \end{aligned} \quad (\text{E-DELETEFIELD})$$

$$\frac{\text{str}_x \notin (\text{str}_1 \dots)}{\text{delete } \{ \text{str}_1: v_1 \dots \} [\text{str}_x] \hookrightarrow \{ \text{str}_1: v_1 \dots \}} \quad (\text{E-DELETEFIELD-NOTFOUND})$$

Functions and Objects



- Dot syntax is desugared to string lookup
- ...

```
let (obj = { "x" : 500, "y" : 100 })
  let (select = func(name) { return obj[name] })
    select("x") + select("y")
→* 600
```

- A program that looks up a non-existent field does not result in an error; instead, JavaScript returns the value `undefined` (E-GETFIELD-NOTFOUND):

```
{ "x" : 7 }["y"] → undefined
```

- Field update in JavaScript is conventional (E-UPDATEFIELD)—

```
{ "x" : 0 }["x"] = 10 → { "x" : 10 }
```

—but the same syntax also creates new fields (E-CREATEFIELD):

```
{ "x" : 0 }["z"] = 20 → { "z" : 20, "x" : 10 }
```

State



$l = \dots$	Locations
$v = \dots l$	Values
$\sigma = \langle(l, v) \dots \rangle$	Stores
$e = \dots e = e \text{ref } e \text{deref } e$	Expressions
$E = \dots E = e v = E \text{ref } E \text{deref } E$	Evaluation Contexts



Example

```
function sum(arr) {  
    var r = 0;  
    for (var i = 0; i < arr["length"]; i = i + 1) {  
        r = r + arr[i] };  
    return r };  
  
sum([1,2,3]) // returns 6  
var a = [1,2,3,4];  
delete a["3"];  
sum(a) // returns NaN !
```

State



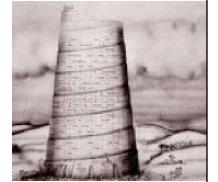
$$\frac{e_1 \hookrightarrow e_2}{\sigma E[e_1] \rightarrow \sigma E[e_2]}$$

$$\frac{l \notin \text{dom}(\sigma) \quad \sigma' = \sigma, (l, v)}{\sigma E[\text{ref } v] \rightarrow \sigma' E[l]} \quad (\text{E-REF})$$

$$\sigma E[\text{deref } l] \rightarrow \sigma E[\sigma(l)] \quad (\text{E-DEREF})$$

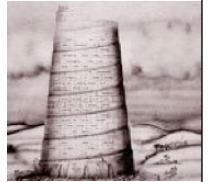
$$\sigma E[l = v] \rightarrow \sigma[l/v]E[l] \quad (\text{E-SETREF})$$

Prototypes



$$\frac{str_x \notin (str_1 \cdots str_n) \quad \text{"__proto__"} \notin (str_1 \cdots str_n)}{\{ str_1 : v_1 , \dots , str_n : v_n \} [str_x] \hookrightarrow \text{undefined}} \text{ (E-GETFIELD-NOTFOUND)}$$

$$\frac{str_x \notin (str_1 \cdots str_n)}{\{ str_1 : v_1 \cdots \text{"__proto__": } p \cdots str_n : v_n \} [str_x] \hookrightarrow (\text{deref } p)[str_x]} \text{ (E-GETFIELD-PROTO)}$$



Desugaring

```
desugar[{prop : e ...}] = ref {
  prop : desugar[e]...,
  "__proto__": (deref Object)["prototype"]
}
```

```
desugar[function(x...) { stmt... }] = ref {
  "code": func(this,x...) { return desugar[stmt...] },
  "prototype": ref { "__proto__":
    (deref Object)["prototype"]
  }
}
```

```
desugar[new ef(e...)] =
  let (constr = deref desugar[ef]
       obj = ref { "__proto__" : constr["prototype"] })
  constr["code"](obj, desugar[e]...);
  obj
```

Desugaring



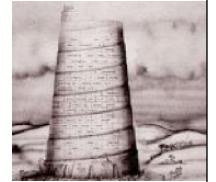
```
desugar[obj[field](e...)] =  
  let (obj = desugar[obj])  
    let (f = (deref obj)[field])  
      f ["code"](obj, desugar[e]...)
```

```
desugar[ef(e...)] =  
  let (obj = desugar[ef])  
    let (func = deref obj)  
      func ["code"](window, desugar[e]...)
```

Desugaring



```
desugar[obj instanceof constr] =  
  let (obj = ref (deref desugar[obj]),  
       constr = deref desugar[constr])  
  done: {  
    while (deref obj !== null) {  
      if ((deref obj).__proto__ === constr.prototype) {  
        break done true }  
      else { obj = (deref obj).__proto__ } };  
  false }  
  
desugar[this] = this (an ordinary identifier, bound by functions)
```



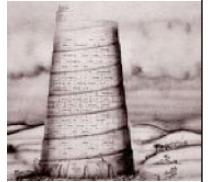
Local variables

- Local variables are automatically lifted to the top of the function:

```
function foo(){ if (true) { var x = 10; } return x }
foo() // returns 10
```

the `return` statement has access to the variable that appeared to be defined inside a branch of the `if`. This can result in unintuitive answers:

```
function bar() { x = 10; var x = x * x; return x }
bar() // returns 100
x // error: unbound identifier
```



Local variables

- All local variable declarations `var x = e` are desugared into simple assignments, `x = e`.
- Furthermore, locals are desugared to let-bindings at the top of functions:

```
let (x = ref undefined) ...
```

Global Variables

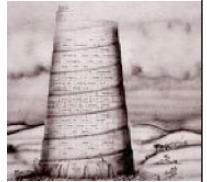


- Global variables are properties of the root scope object, called `window`. `window` has a field `window.window` that references itself:

```
window.window === window // returns true
```

A program can obtain a reference to the global scope object by simply writing `window`. As a consequence, globals seem to break lexical scope, since we can observe that they are properties of `window`:

```
var x = 0;  
window.x = 50;  
x // returns 50  
x = 100;  
window.x // returns 100
```

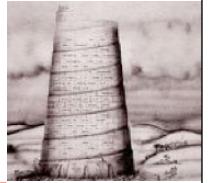


Global Variables

However, `window` is the only scope object that is directly accessible to JavaScript programs [6, Section 10.1.6]. We can maintain lexical scope by abandoning global variables. That is, we simply desugar the code above to the following:

```
window.x = 0;  
window.x = 50;  
window.x // returns 50  
window.x = 100;  
window.x // returns 100
```

with



- The `with` statement adds an arbitrary object to the front of the scope chain:

```
function(x, obj) {  
  with(obj) {  
    x = 50; //if obj.x exists then obj.x=50 else x=50  
    return y } } //return either obj.y, or window.y
```

We desugar `with` by turning comments above into code:

```
function(x, obj) {  
  if(obj.hasOwnProperty("x")) {obj.x=50} else {x=50}  
  if (obj.hasOwnProperty("y")) { ...
```