

Memory Management for Real-time Java

Jan Vitek

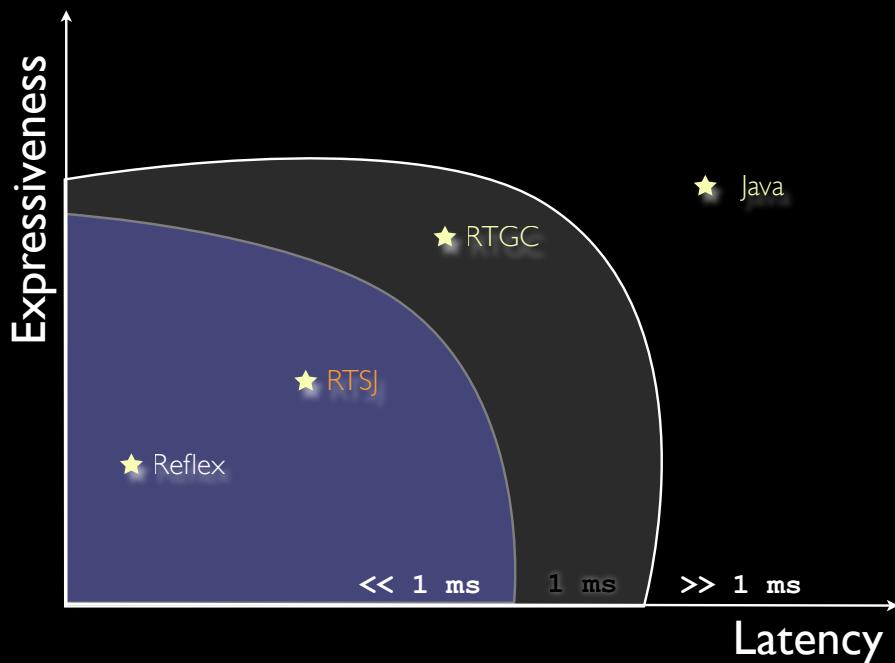


ECS 2009

Overview

- Memory management & Java
- Scoped Memory
- Real-time Collectors:
 - ▷ AICAS Jamaica, Sun McKinak, IBM Metronome
- Lock-free Collectors
- Conclusions

Design space



Manual Memory Management

- Memory management is the source of many errors (dangling pointers, memory leaks) in traditional software systems due to the non-local nature of the allocation/deallocation protocol
- One can sometimes prevent such errors by avoiding allocation, at the cost of over provisioning, but not in general.
- Programmers often use custom ‘allocators’ in conjunction with object pools
- ... but rich data structures and concurrency make manual management hard*
- ... requires synchronization on multicores*
- ... must deal with fragmentation*

Garbage Collection



ECS 2009

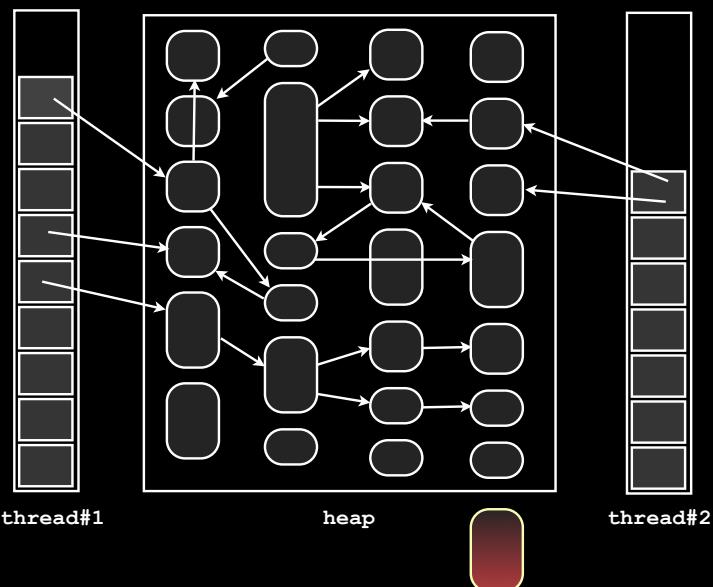
Garbage Collection

- A Garbage Collector (GC) is an algorithm that automatically finds unused objects in the memory of an application and prepares them for reuse
- GC frees programmers from worrying about the exact lifetime of objects and ensures that the heap will not be corrupted by access to previously freed data
- ... *but introduces pauses that may be $O(\text{heap})$ and can increase the memory required.*
Moreover, pauses occur at unpredictable times, especially in concurrent programs

Garbage Collection

Phases

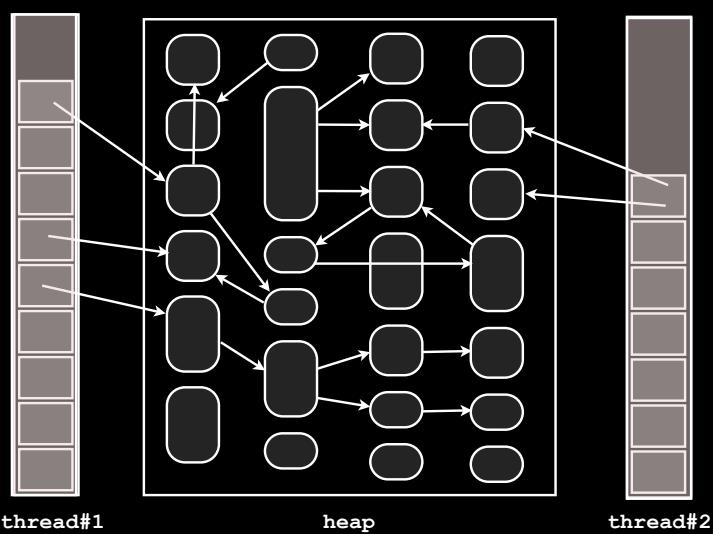
- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction



Garbage Collection

Phases

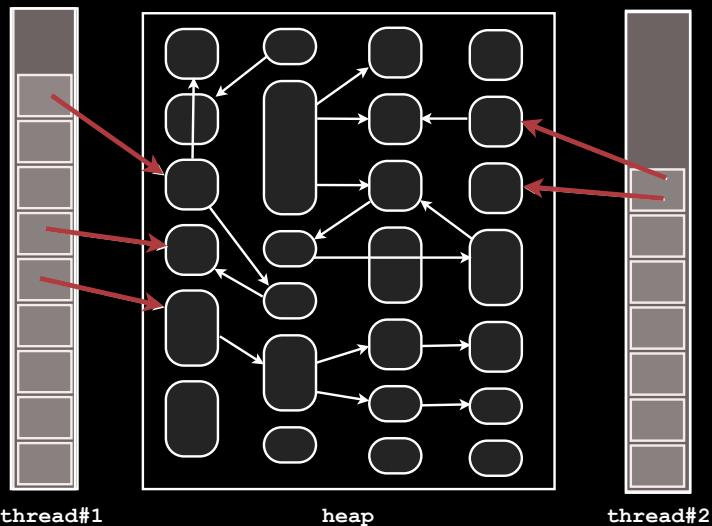
- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction



Garbage Collection

Phases

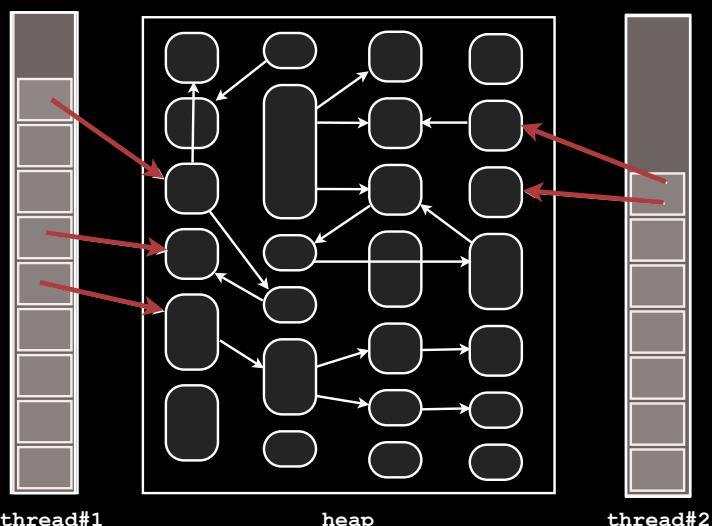
- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction



Garbage Collection

Phases

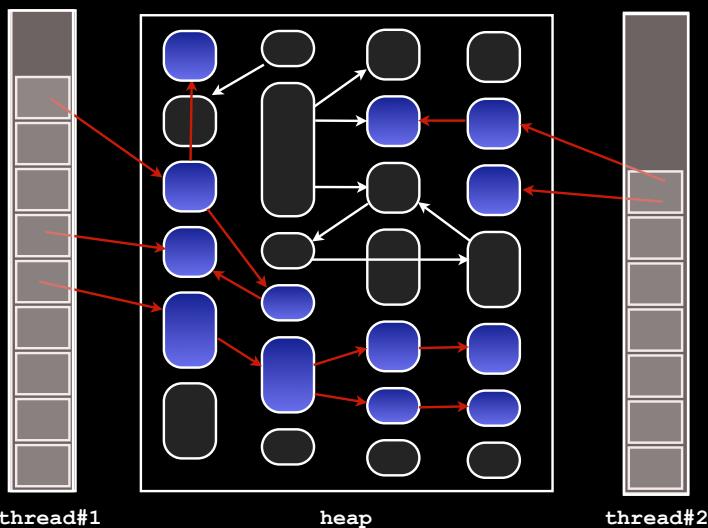
- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction



Garbage Collection

Phases

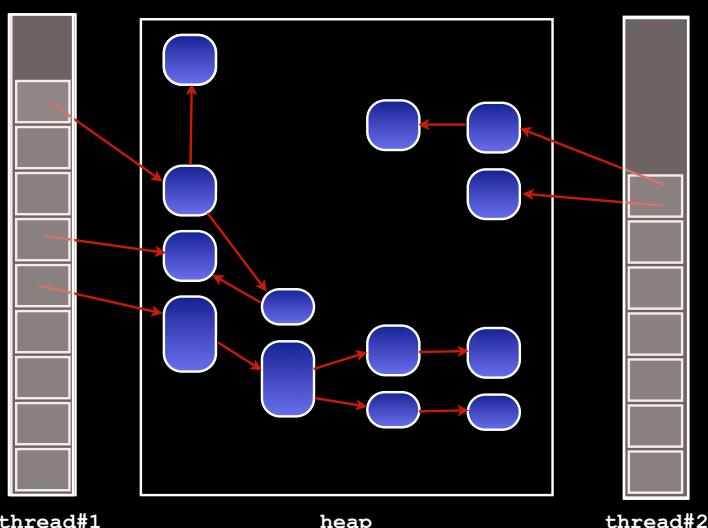
- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction



Garbage Collection

Phases

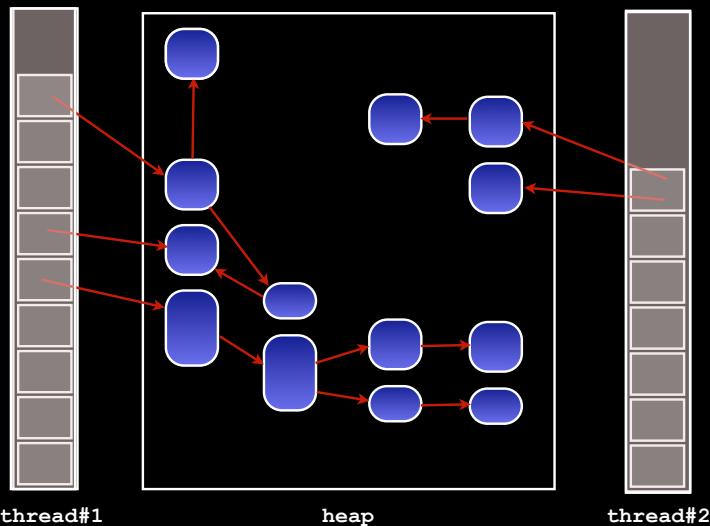
- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction



Garbage Collection

Phases

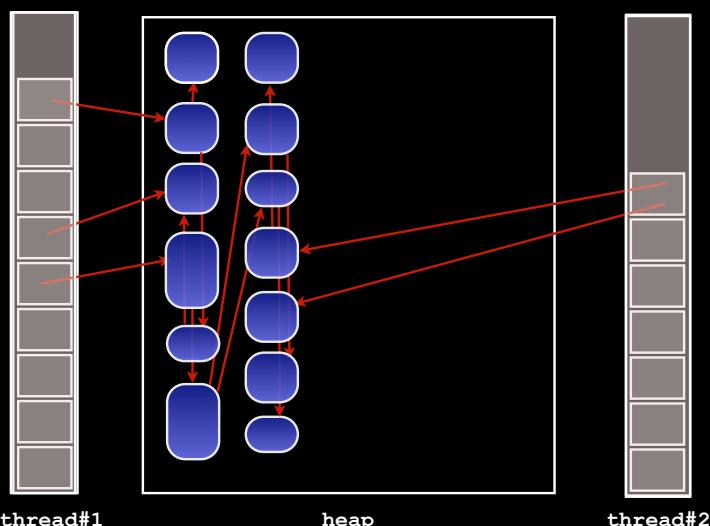
- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction



Garbage Collection

Phases

- Mutation
- Stop-the-world
- Root scanning
- Marking
- Sweeping
- Compaction



GC is Easy

- If responsiveness is not an issue,
the GC can complete in one long pause under the
assumption that there is no interleaved application activity
- Marking is easy
if the graph does not change while you are searching it.
- Copying/compacting objects and fixing up the heap is easy
if the application is prevented from accessing the heap

Real-time Garbage Collection

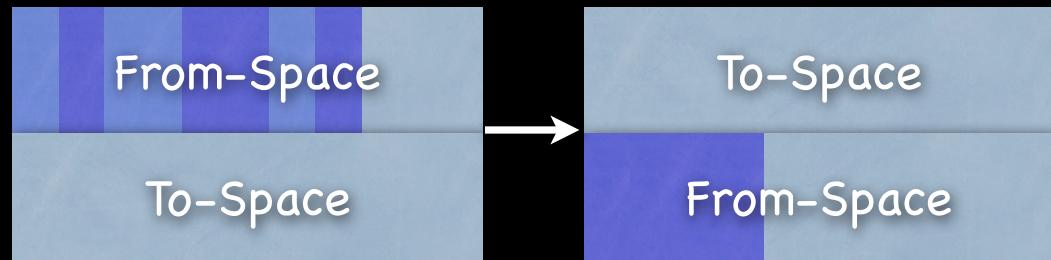
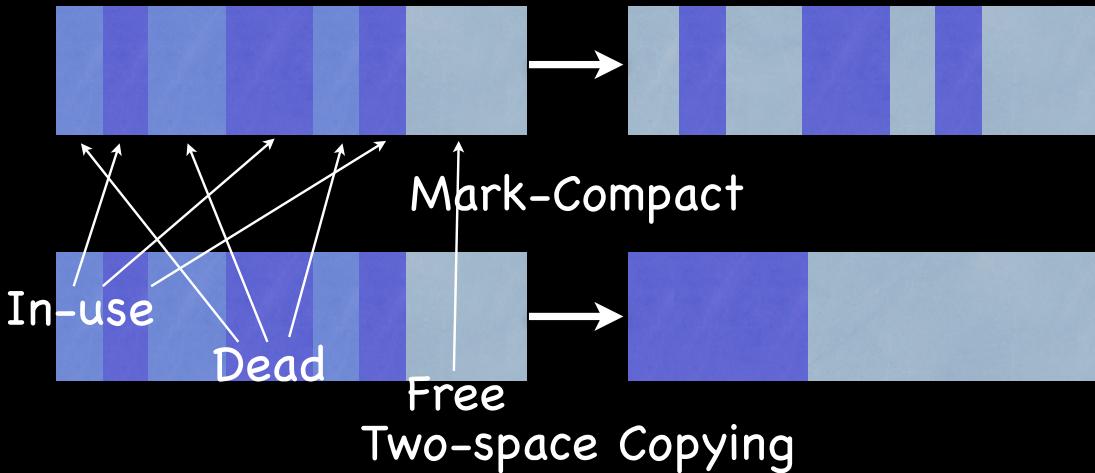
Real-time GC

- A Real-time GC must provide time and space predictability
 - ▶ provide a performance model that can be used to guarantee that programs do not run out of memory or experience pauses that violate their timing constraints
- A Real-time-GC must support defragmentation of the heap if it is to be used with long-lived applications
- Multi-processor support is unavoidable
- Throughput should not degrade overly

Main Collector Types

- Reference counting
 - ▶ keep count of incoming references. Requires auxiliary GC to reclaim cycles.
- Mark-Sweep
 - ▶ one phase to mark used objects, another to sweep unused space. No copying is performed.
- Mark-Compact
 - ▶ one phase to mark used objects, another to copy all used objects to one side of the heap.
- Copying
 - ▶ one phase to simultaneously mark and evacuate all used objects; the space that previously held all objects is then completely free.

Mark-Sweep



RTGC: state of the art

- Sort-of Real-time GCs:
 - ▷ Too many to list, Baker (1978) was first
- Really Real-time GCs:
 - ▷ JamaicaVM (1999) - work-based
 - ▷ Henriksson (1998) - slack-based
 - ▷ Metronome (2003) - time-based

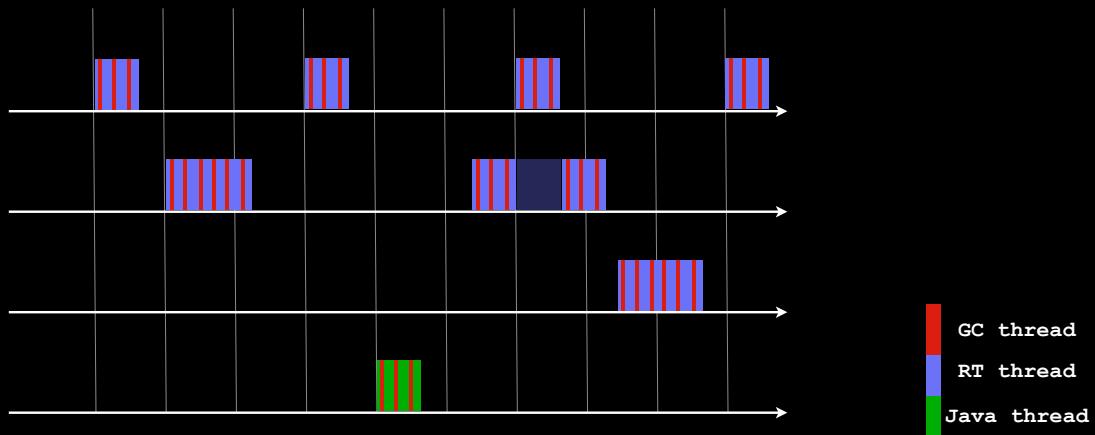
Work-based RTGC

Baker'78, Siebert'99



ECS 2009

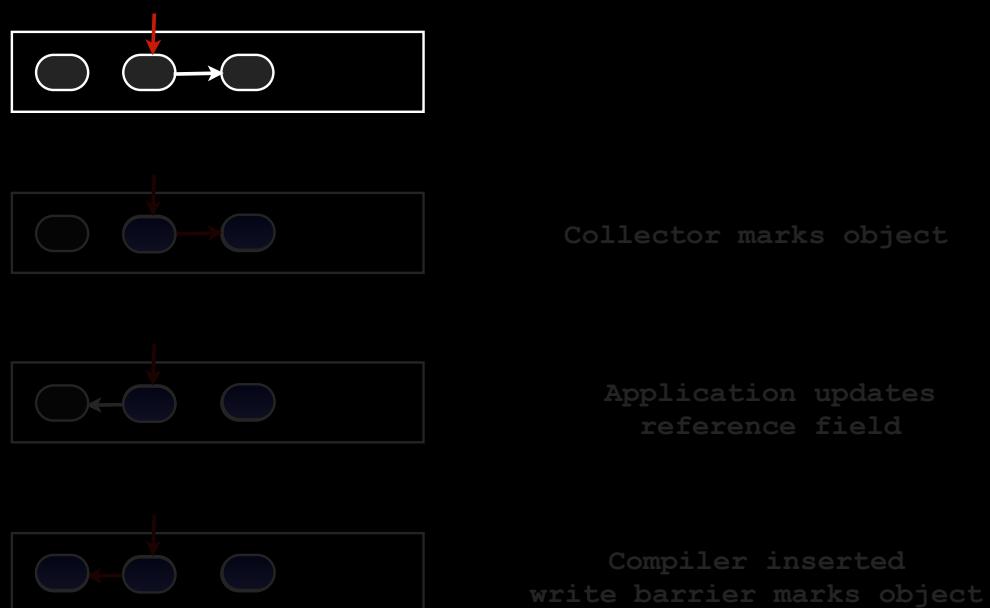
Work-based scheduling



Baker

- First attempt at RTGC by Baker '78
- For LISP where all objects are the same, small, size
- Copying collector that instruments the mutator to maintain:
 - ▶ To-space invariant:
 - mutator uses to-space, never sees from-space, objects are copied on first read
 - ▶ Steady state:
 - collector keeps up with mutator allocations, collection work performed on each allocation

Incrementalizing marking



Dissecting Baker

- The chief problem with Baker is the to-space invariant
 - ▷ Reads are frequent, a heavy read barrier leads to poor performance
 - ▷ Variable sized objects in Java make costs harder to bound
- What RT programmers fear:
 - ▷ a burst of copying reads, causing substantial slow-downs

Siebert

- Modern work-based collector in the JamaicaVM:
 - ▷ Mark-Sweep, no copying
 - ▷ All objects are same size, 32 bytes
(large objects⇒ lists, arrays⇒ tries) Stack is logically an object
 - ▷ Root scanning is fully incremental
 - ▷ Write barrier to mark objects (including the stack)
 - ▷ Allocation triggers a bounded amount of collector work

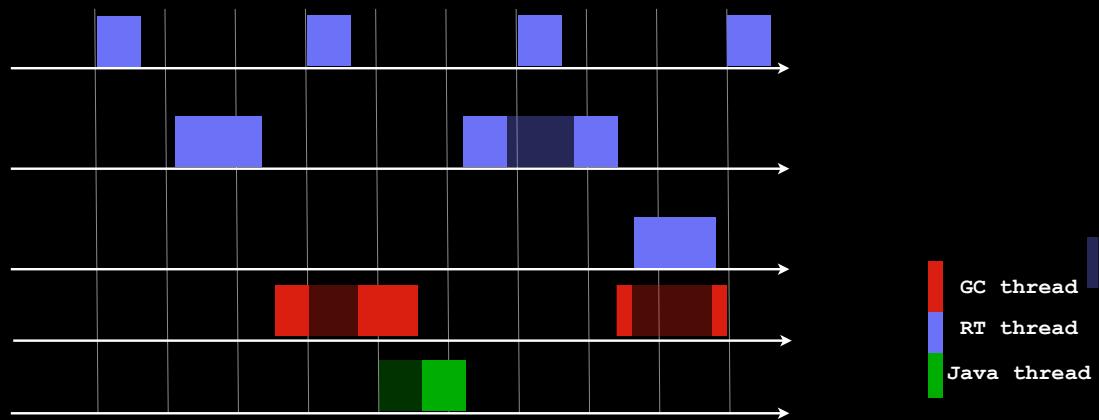
Slack-based RTGC

Henriksson'98



ECS 2009

Slack-based GC Scheduling

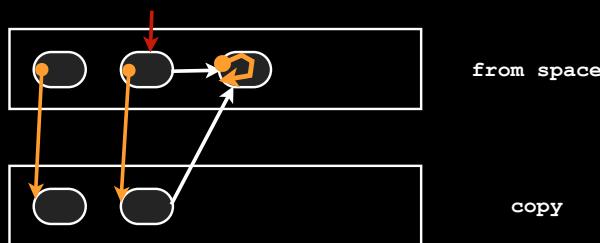


Slack-based GC

- Henriksson takes a different approach to real-time guarantees:
 - ▷ tries to guarantee real-time tasks never experience GC interference
- How is this done?
 - ▷ RT tasks can always preempt GC
 - ▷ Objects have normal representation.
 - ▷ Read & write barriers are used to ensure marking/copying soundness.
 - ▷ Copying collector, with roll backs when the mutator preempts GC

Incrementalizing copying/compaction

- Forwarding pointers are used to refer to the current version of the object.
- Every access must start with a dereference



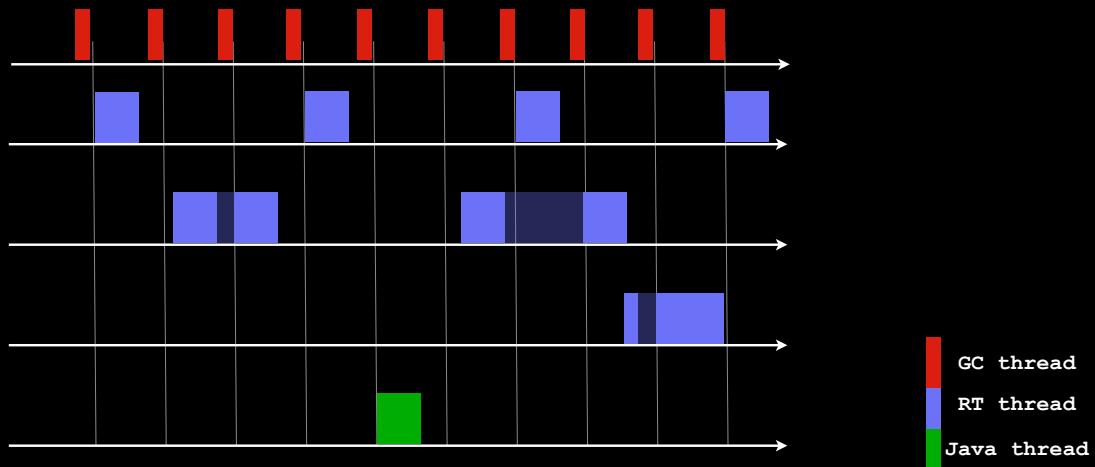
Dissecting Henriksson

- Pros:
 - ▶ Gives the programmer full control over the GC schedule
 - ▶ SUN uses a variant of this collector in their RTSJVM
- Cons:
 - ▶ Read barrier are expensive
 - ▶ Mutator induced copy aborts make collector progress hard to ensure
 - ▶ Programmer must know allocation rate, and to budget for GC work
 - ▶ Priority inversion can lead to GC interference

Time-based RTGC

Metronome'03

Time-based GC Scheduling



Time-based GCs

- Metronome runs periodically for a bounded amount of time
 - ▶ Amount of collector interference is easy to understand
 - ▶ Progress is easier to guarantee than in Henriksson
- A Mark-Sweep-and-sometimes-Copy collector:
 - ▶ Write barrier to ensure marking soundness
 - ▶ Light read barrier to ensure copying soundness
 - ▶ Almost-empty pages evacuated in response to fragmentation
 - ▶ Arrays are split into tries with page-sized leaves
 - ▶ Collector increments lower-bounded by time to copy a page-size objects

Suming up



ECS 2009

Siebert

Local ref. assignment	mark(b) $a = b$
Primitive heap load	$a = b \rightarrow f$
Reference heap load	$a = b \rightarrow f$ mark(a)
Primitive heap store	$a \rightarrow f = b$
Reference heap store	mark(b) $a \rightarrow f = b$

Henriksson

Local ref. assignment	$a = b$
Primitive heap load	$a = b \rightarrow \text{forward} \rightarrow f$
Reference heap load	$a = b \rightarrow \text{forward} \rightarrow f$ $\text{mark}(a)$
Primitive heap store	$a \rightarrow \text{forward} \rightarrow f = b$
Reference heap store	$\text{mark}(b)$ $a \rightarrow \text{forward} \rightarrow f = b \rightarrow \text{forward}$

Metronome

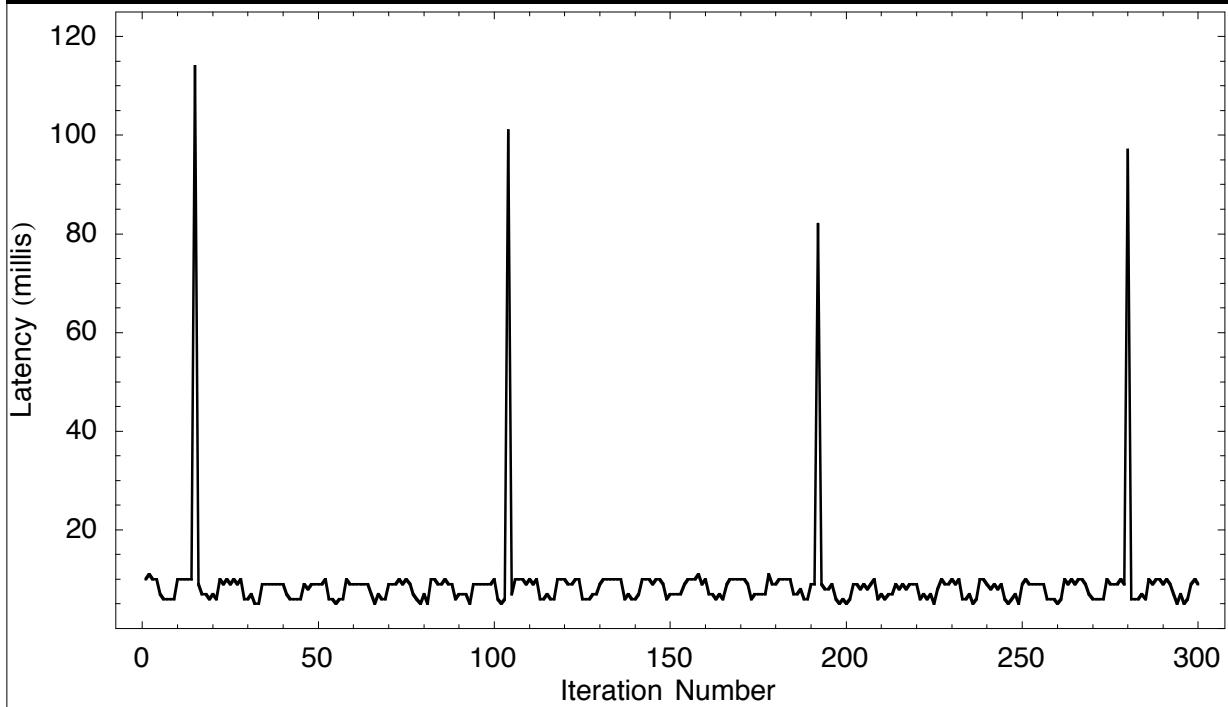
Local ref. assignment	$a = b$
Primitive heap load	$a = b \rightarrow \text{forward} \rightarrow f$
Reference heap load	$a = b \rightarrow \text{forward} \rightarrow f$
Primitive heap store	$a \rightarrow \text{forward} \rightarrow f = b$
Reference heap store	$\text{mark}(a \rightarrow \text{forward} \rightarrow f)$ $a \rightarrow \text{forward} \rightarrow f = b \rightarrow \text{forward}$

Collision Detector

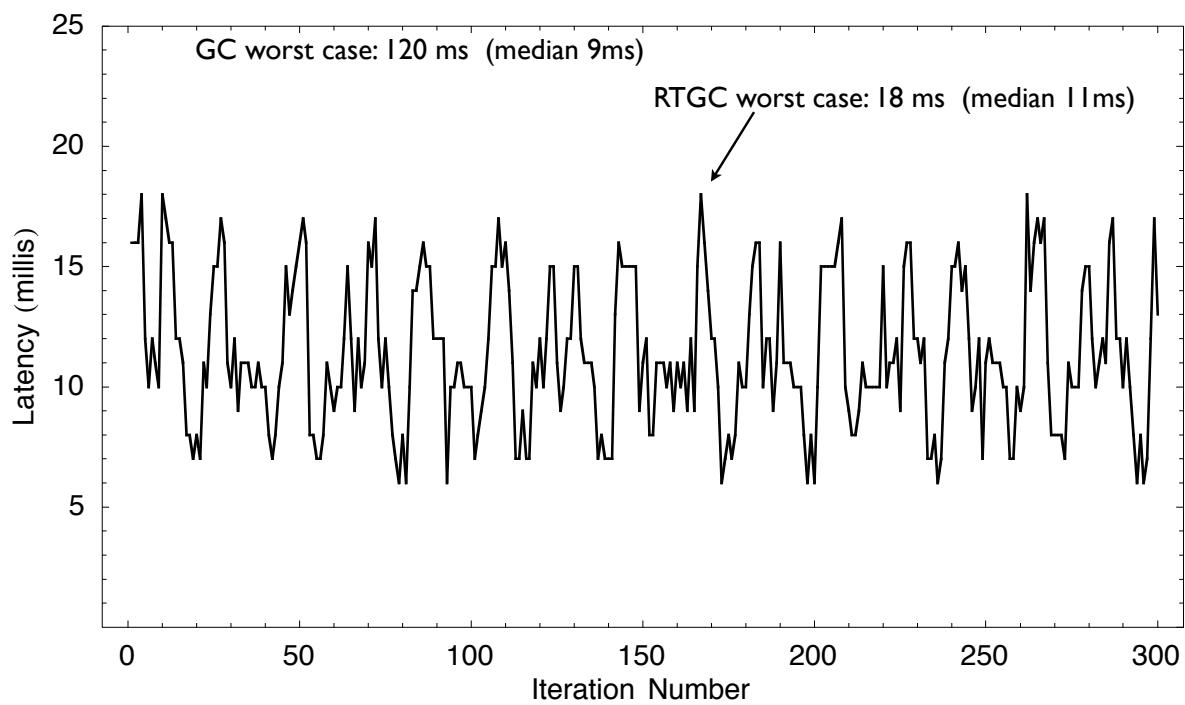
- Experiment:

- ▶ Pentium IV 1600 MHz, 512 MB RAM, Linux 2.6.14, GCC 3.4.4
- ▶ Application: Real-time Java collision detector (20Hz)
- ▶ Virtual machine: Ovm

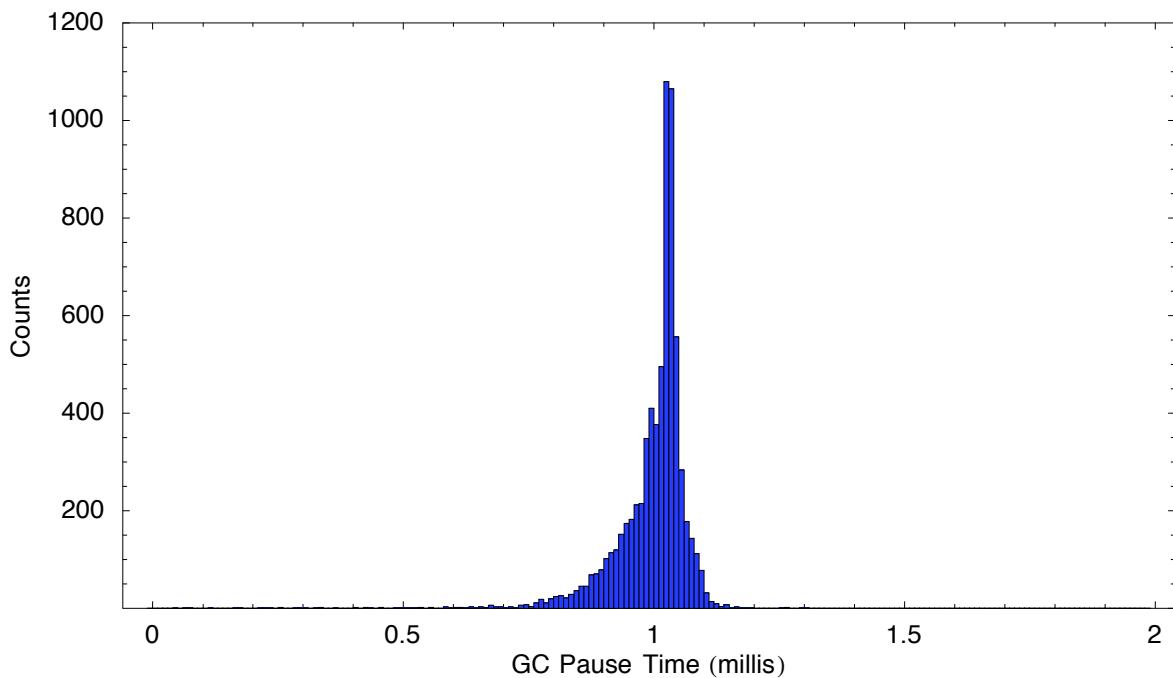
CD with GC



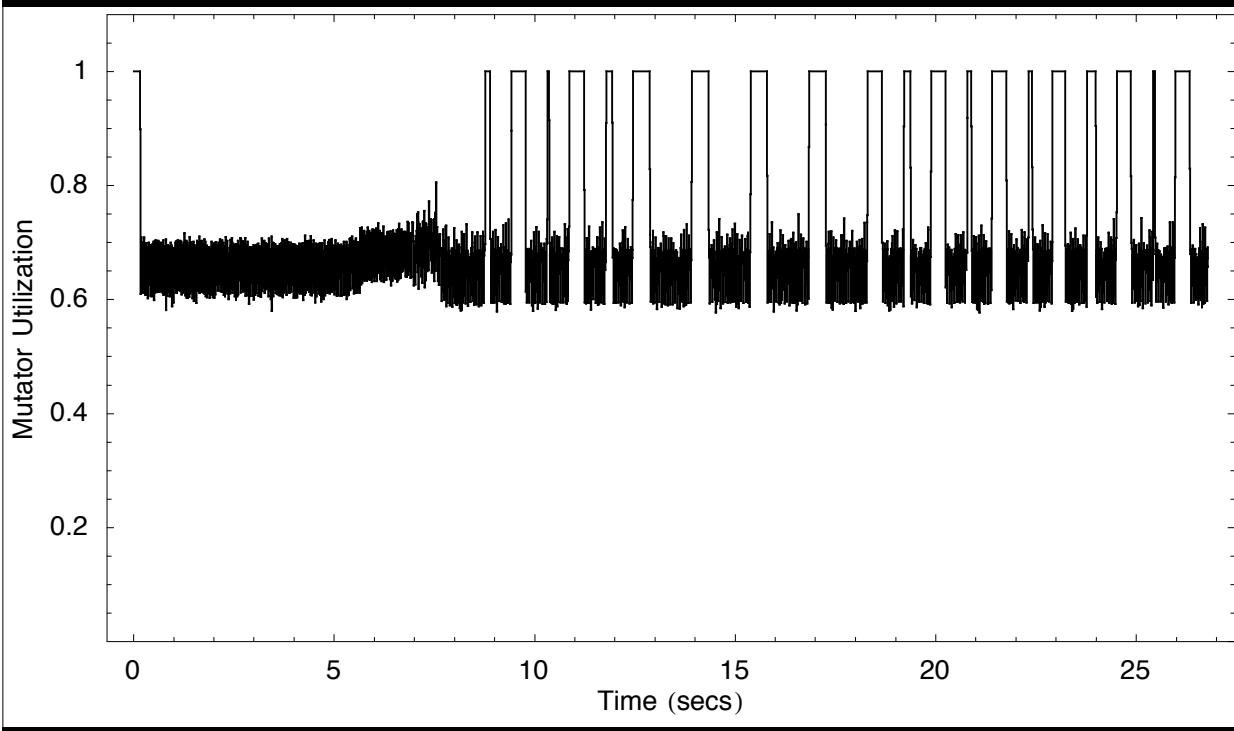
CD with RTGC



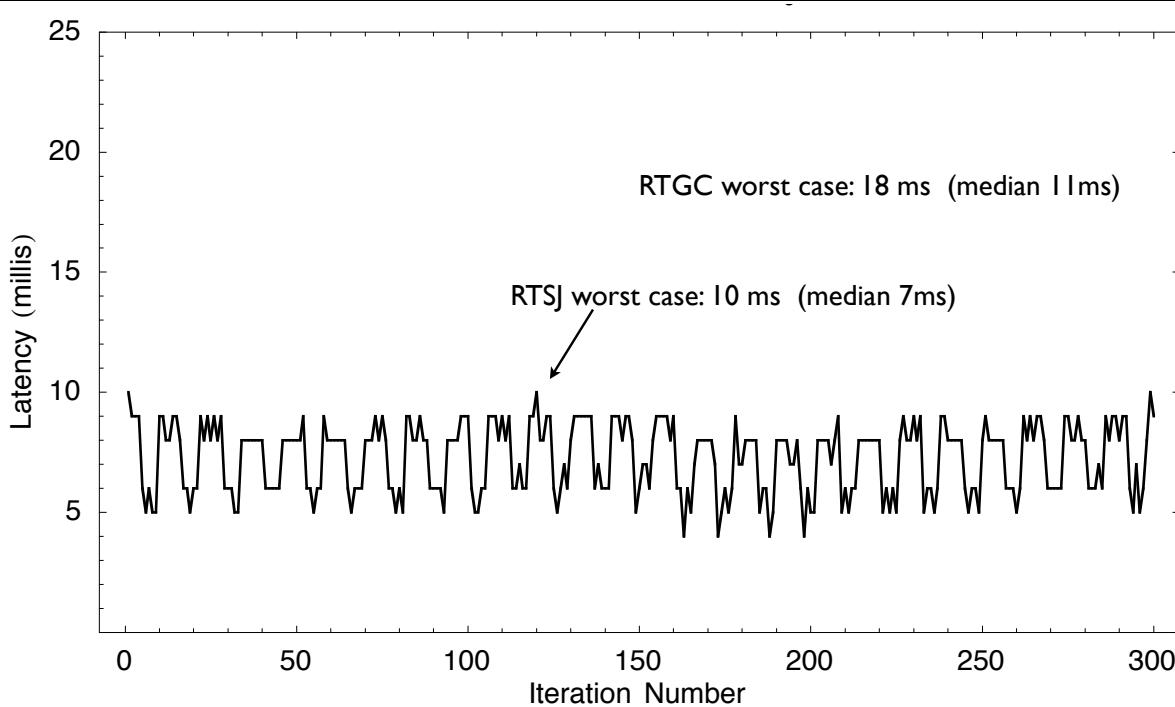
RTGC Pause time distribution



Utilization

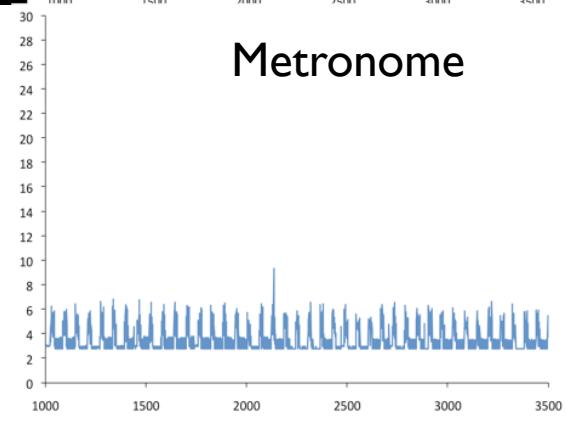
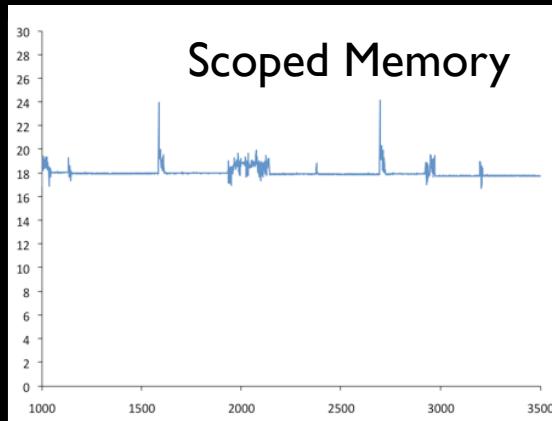
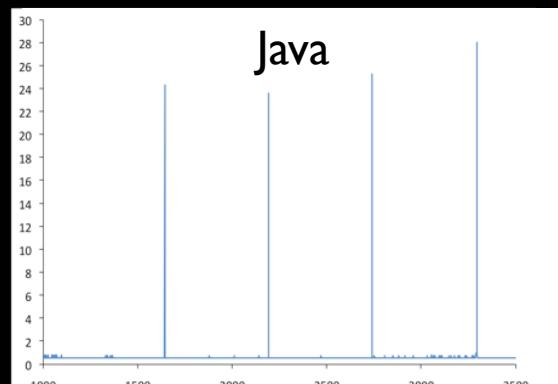


CD with RTSJ



CD on J9

- The implementation seems to favor GC over scopes



Lessons from the RTSJ

1/9/90 Rule

- The main reason of RTSJ's success is that it allows mixing freely real-time and timing-oblivious code in the same platform
- Rule of thumb: 1% hard real-time, 9% soft real-time, 90% non-RT
- RT Programming Model should:
 - ▶ Allow non-intrusive injection of real-time in a timing-oblivious system
 - ▶ Avoid paradigm switches---single semantic framework for RT/non-RT

Small Footprint

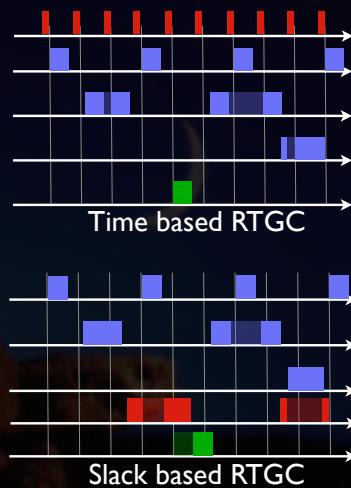
- The RTSJ requires sweeping changes in the Virtual Machine and has complex interaction with Java features.
 - ▶ The RTSJ changes the semantics of “normal” Java code:
 - any access to a reference variable may throw an exception,
 - meaning of “throws” changed to support asynchronous exception,
 - finalization interacts with memory management.
- Avoid complex APIs with ill-defined feature interactions
- Don’t change the semantics of non-RT code

Schedule Flexibly

- Support for real-time scheduling is crucial with a clean integration in the language concurrency model
 - ▷ Priority Preemptive Scheduling for periodic tasks
 - ▷ Support PIP and PCE locks for priority inversion avoidance
- Provide facilities for writing non-blocking algorithms
- Allow for user-defined schedulers.

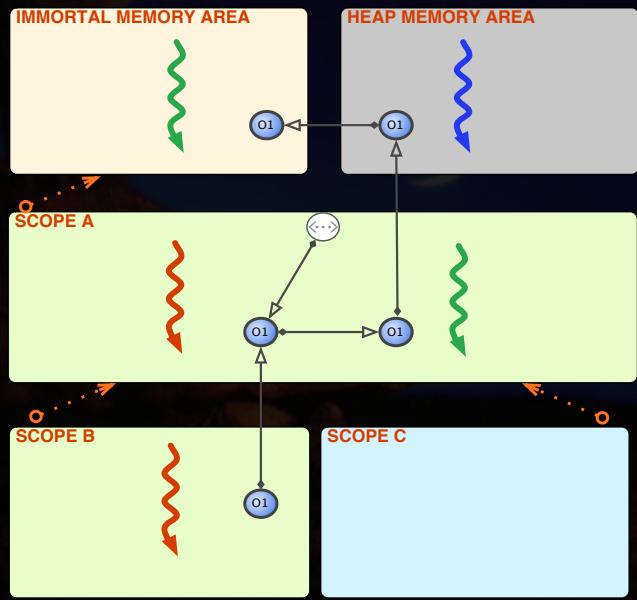
Don't Fear Garbage

- Real-time Garbage Collectors good enough for most RT apps
 - ▷ Time-based collectors (Metronome) can easily bound pauses to one millisecond.
 - ▷ Automatic defragmentation for long lived apps.
 - ▷ Slack-based collectors (SUN's) ensure RT code never preempted by GC.
 - ▷ Ensuring that the GC can keep up requires some understanding of whole-program allocation behavior.
- Avoid manual memory management unless absolutely necessary.



Type-check Region-based Allocation

- Regions are the worst mistake of the RTSJ
 - ▷ Dynamic safety checks slow down programs and cause runtime errors that are hard to prevent
 - ▷ Distinction between heap-using and non-heap thread is tricky and error prone.
- Use static type system for assurance and efficiency



Shared-memory Concurrency is Hard

- Real-time presents a number of concurrency related challenges.
 - ▷ Critical sections must be short --- avoid over-synchronization
 - ▷ Avoid priority inversion between real-time and non-RT code
 - ▷ Avoid blocking operations in RT code
- Limit/localize communication between RT and non-RT.
- Support flexible isolation to encourage determinism.